

Image Warping

Logistics

- We will have lecture next Tuesday (changed to later flight)

Date	Topic	References	Misc
1/16	Introduction		[First class]
1/18	Filters		Special Recitation on HW0; NSH 3305 at 5pm
1/23	Edges		HW0 - Python Intro (not graded)
1/25	Texture		
1/30	Linear algebra review		
2/1	Warping		HW1 - HOG Image Classification
2/6	Alignment	Deva in town!	
2/8	Frequency	Guest lecture	
2/13	Correspondence		
2/15	Descriptors		HW2 - LK Tracking
2/20	Image Formation		
2/22	Cameras		
2/27	Motion		
2/29	Flow		HW3 - Homographies
3/5	NO CLASS [Spring Break]		
3/7	NO CLASS [Spring Break]		
3/12	Two-view		
3/14	Stereo		
3/19	SFM		
3/21	Recognition		HW4 - Reconstruction
3/26	Classifiers	Guest Lecture	
3/28	MLPs	Guest Lecture	
4/2	Training Recipes		
4/4	CNNs		
4/9	Interpreting Networks		HW5 - Neural Networks
4/11	NO CLASS [Spring Carnival]		
4/16	Radiometry		
4/18	Color		
4/23	Recent Trends		
4/25	Catchup		HW6 - Photometric Stereo

Printing PDFs; nice YouTube posted to Piazza on multiple pathways to generate PDF
Suggestion; give yourself at least 30 min to iron this out for tonight's deadline

Linear algebra digression

https://www.cs.cmu.edu/~zkolter/course/linalg/linalg_notes.pdf

My own writeup (in lec gdrive folder): updated for today

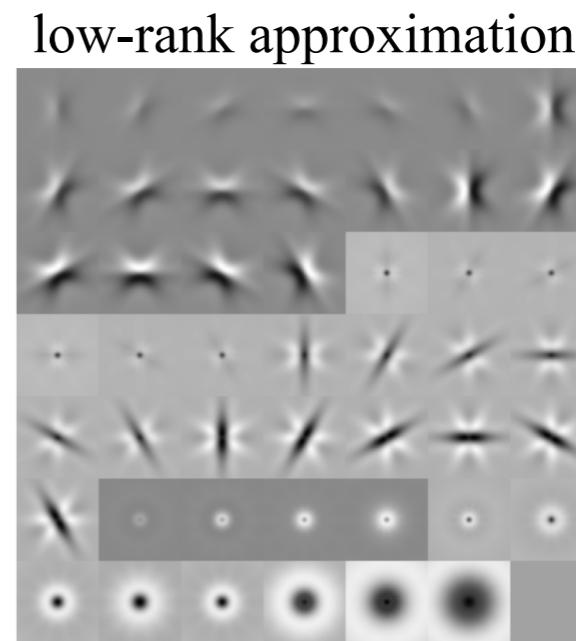
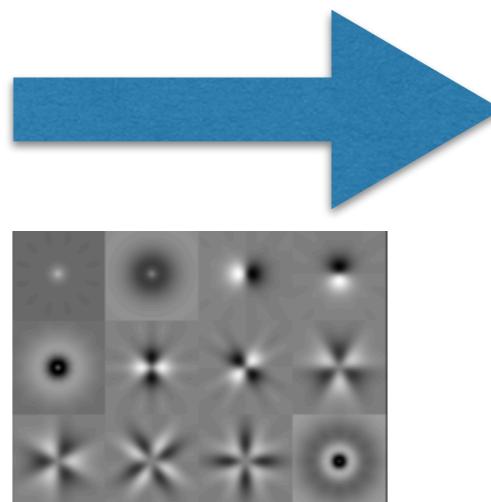
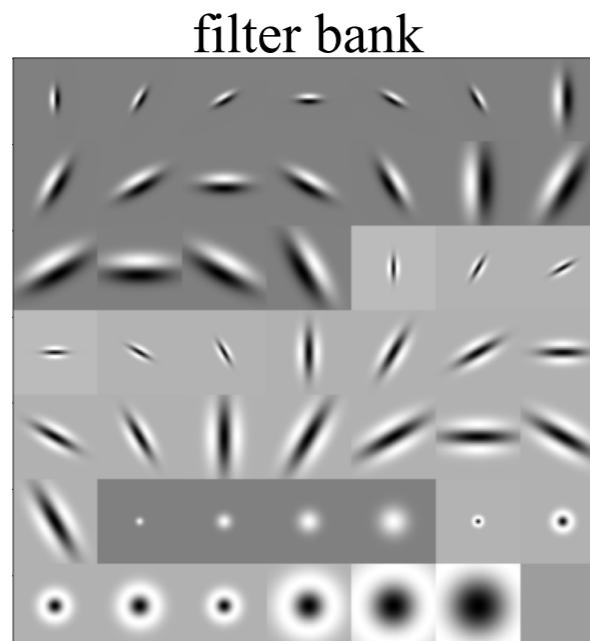
https://drive.google.com/file/d/1c-kjMqIQEKJCHyfMgLzdxRJNReOu6ssk/view?usp=drive_link

Least-squares method of steerability

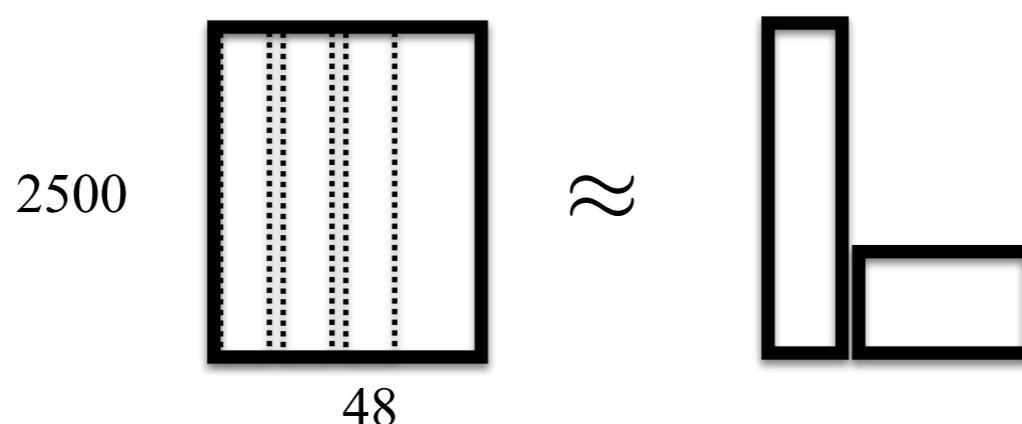
demo_faces.ipynb

demo_filter_bank.ipynb

Shy & Perona, CVPR94

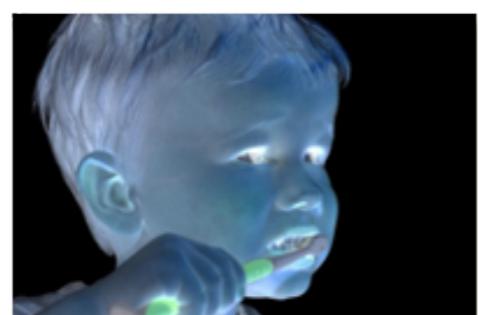


- Column-vectorize each of 48 50x50 filters and stack'em into a 2500x48 matrix
- Apply SVD to generate low-rank approximation of 48 filters as linear combination of 12 *basis* filters



What kind of image transformations can we perform?

Filtering

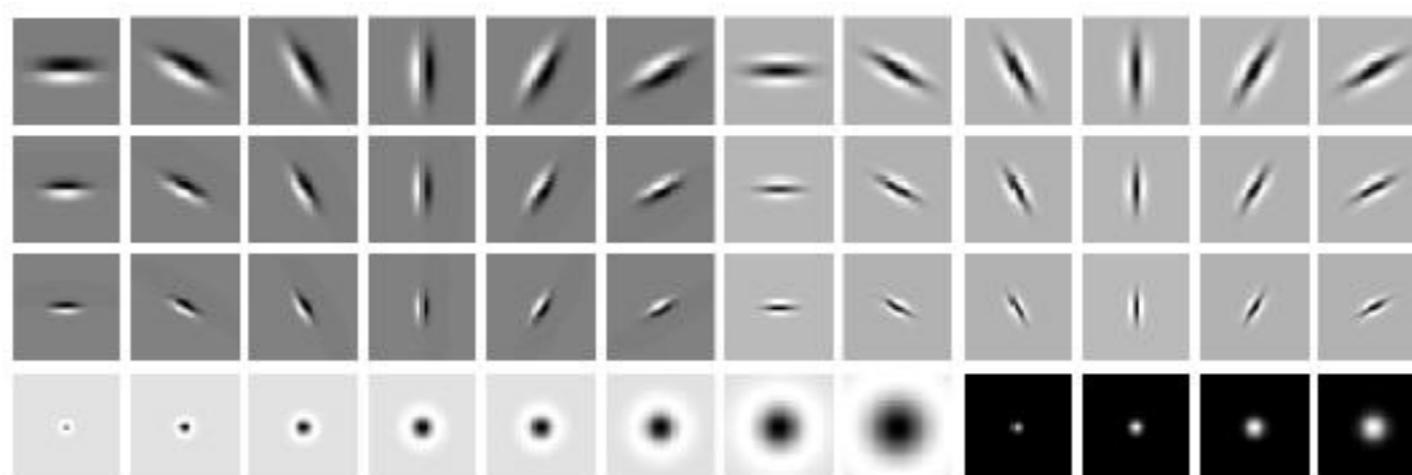


Changes pixel values $I[x,y]$

Warping



Changes pixel coordinates (x,y)



2D transformations



translation



rotation



aspect



affine

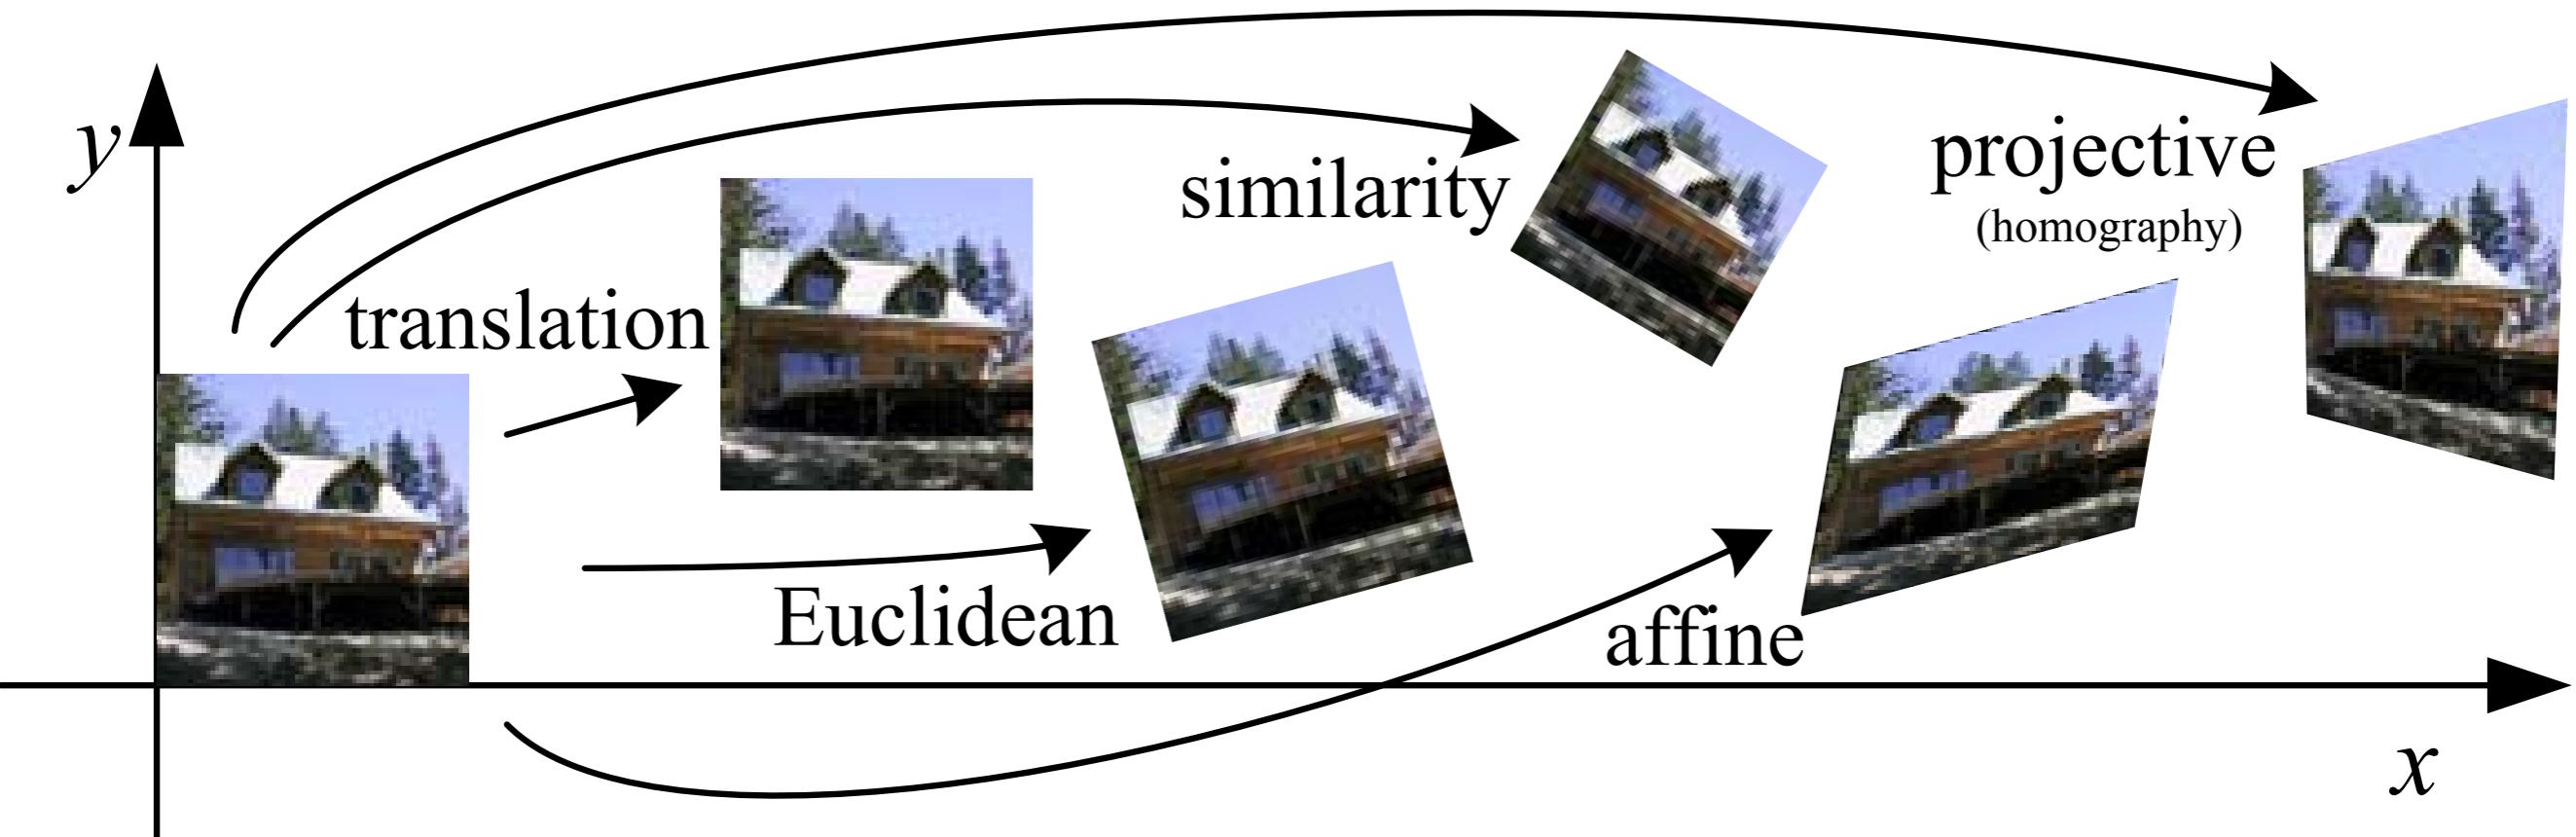


perspective



cylindrical

Family of “nice” 2D transformations (can be represented with a matrix)



Question: which of these transformations are linear-shift-invariant?

translation is; rotation breaks it

(rotating + then shifting image does not equal shifting then rotating)

Family of 2D warps

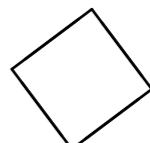
Transformation	Matrix	# DoF	Preserves	Icon
translation	$\left[\begin{array}{c c} \mathbf{I} & \mathbf{t} \end{array} \right]_{2 \times 3}$	2	orientation	

$x' = x + t_x$

$y' = y + t_y$

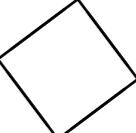
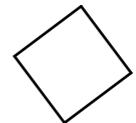
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Family of 2D warps

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	

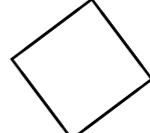
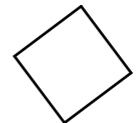
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & t_x \\ \sin \theta & \cos \theta & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Family of 2D warps

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	

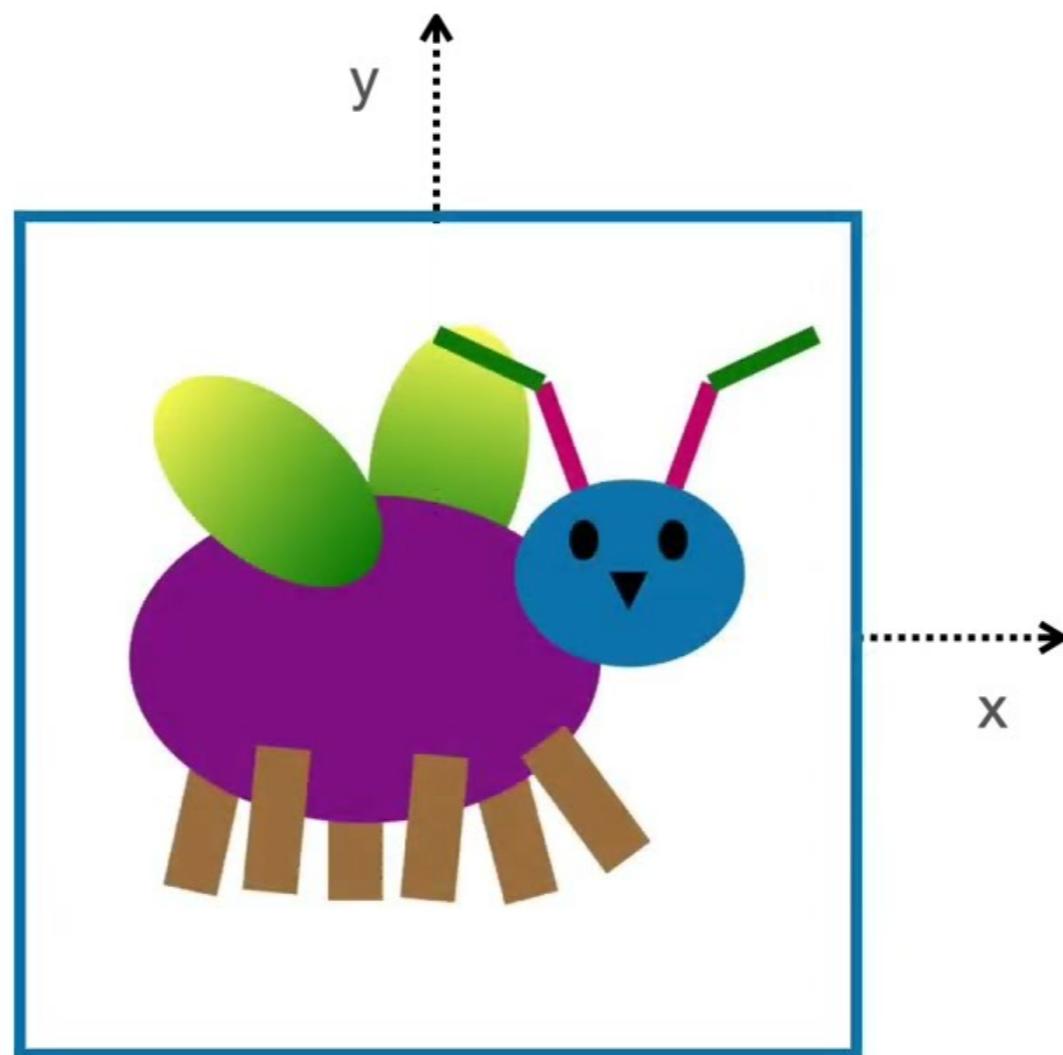
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s \cos \theta & -s \sin \theta & t_x \\ s \sin \theta & s \cos \theta & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Family of 2D warps

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
	$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$			

Think of as change of basis and 2D translation

Cool video (from Steve Seitz at UW)



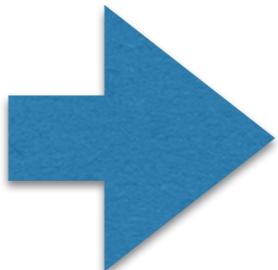
https://www.youtube.com/watch?v=AheATd_15Is&list=PLWfDJ5nla8UpwShx-lzLJqcp575fKpsSO&index=1

Family of 2D warps

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

“homogenous 2D coordinates”

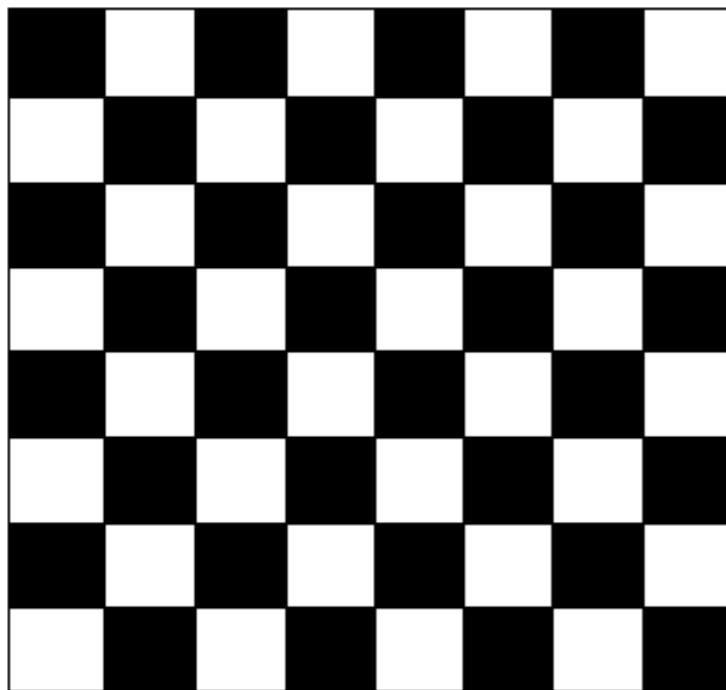


$$x' = \frac{ax + by + c}{gx + hy + i}$$

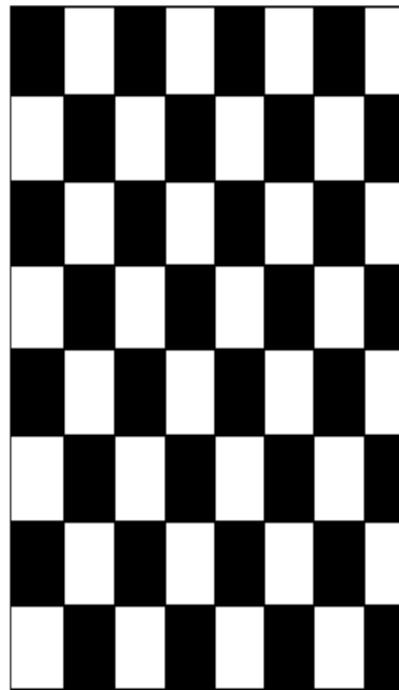
$$y' = \frac{dx + ey + f}{gx + hy + i}$$

(for now, think of as shorthand for...)

Important 3Dish property captured by 2D affine warps: *foreshortening*



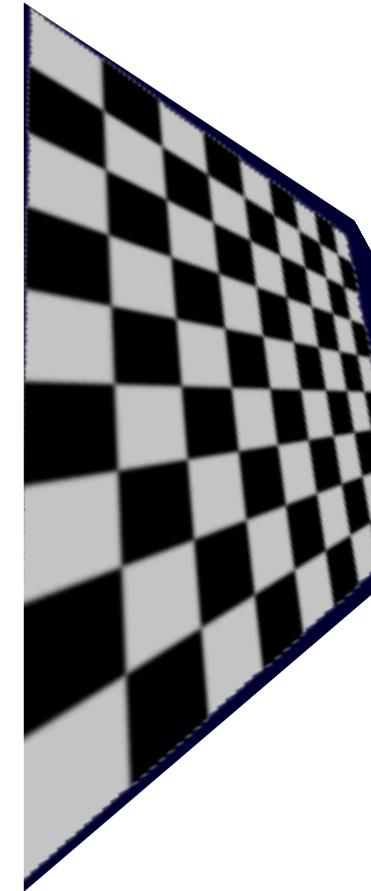
Fronto-parallel view



Foreshortened view

Rotation of far-away plane

Affine warp



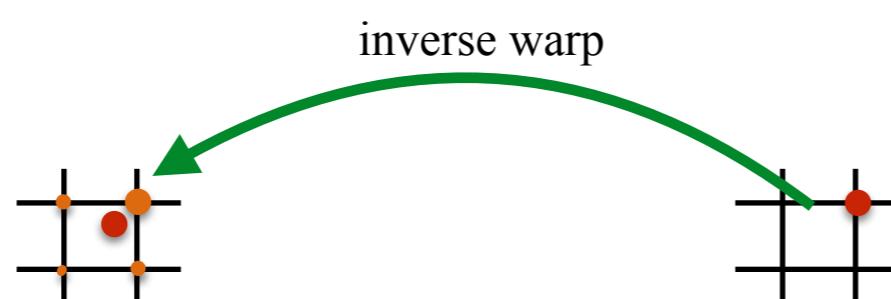
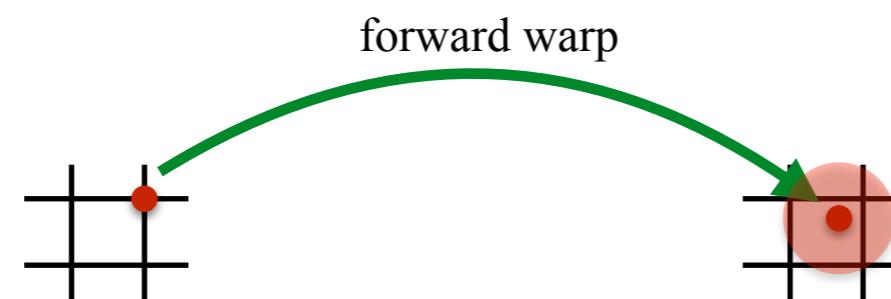
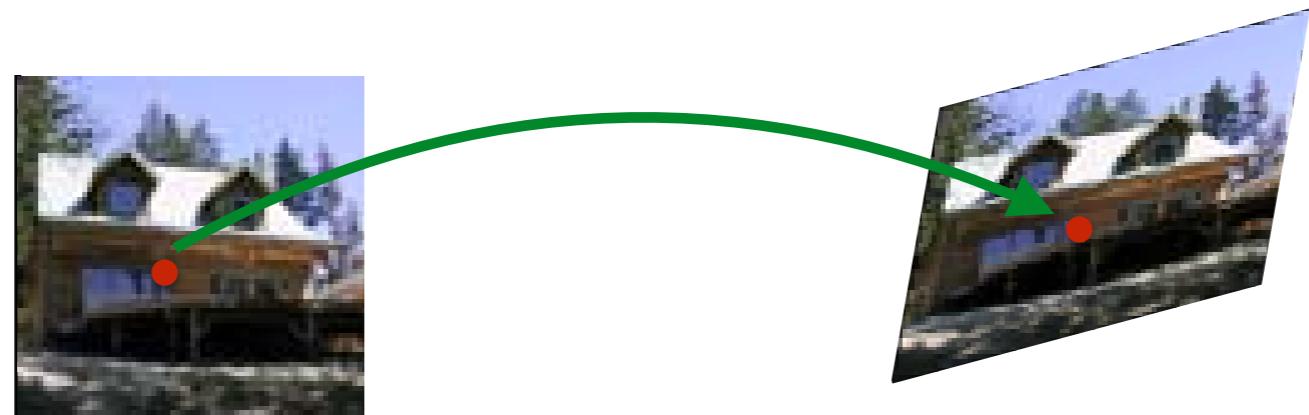
Perspective view

Rotation of close-by plane

Projective warp

(part of plane that gets further away becomes smaller)

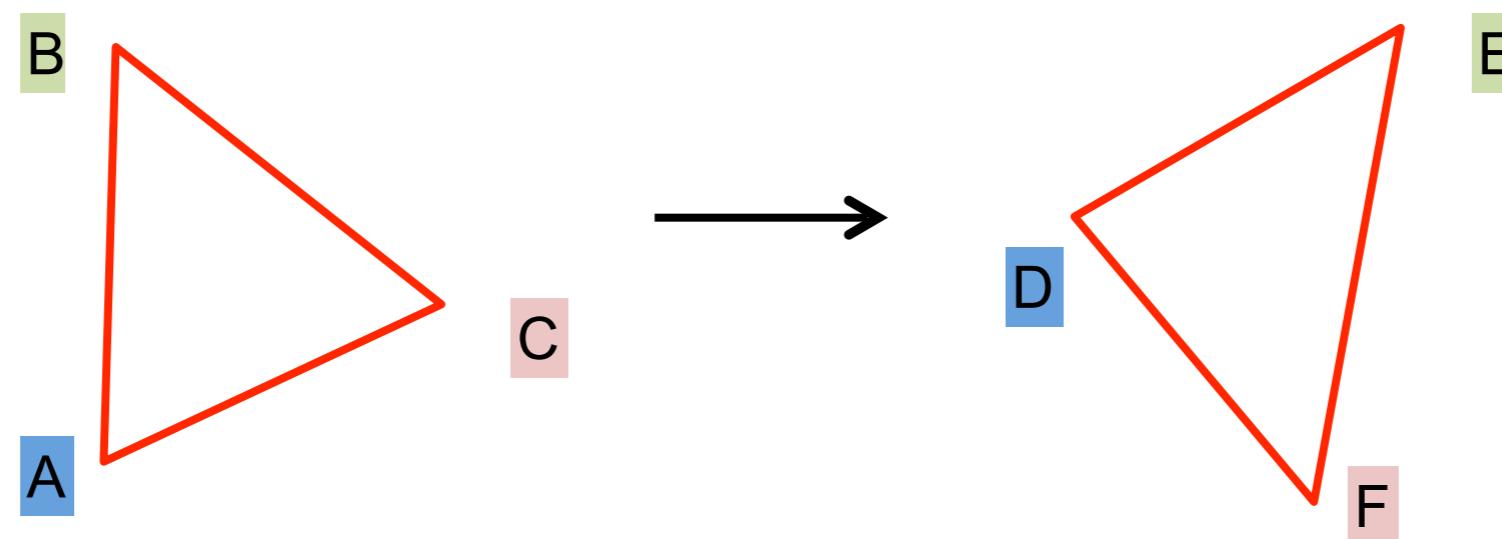
Forward vs inverse warps



[you'll use this in HW3!]

Estimating transformations from point correspondences

Suppose we have two triangles: ABC and DEF.

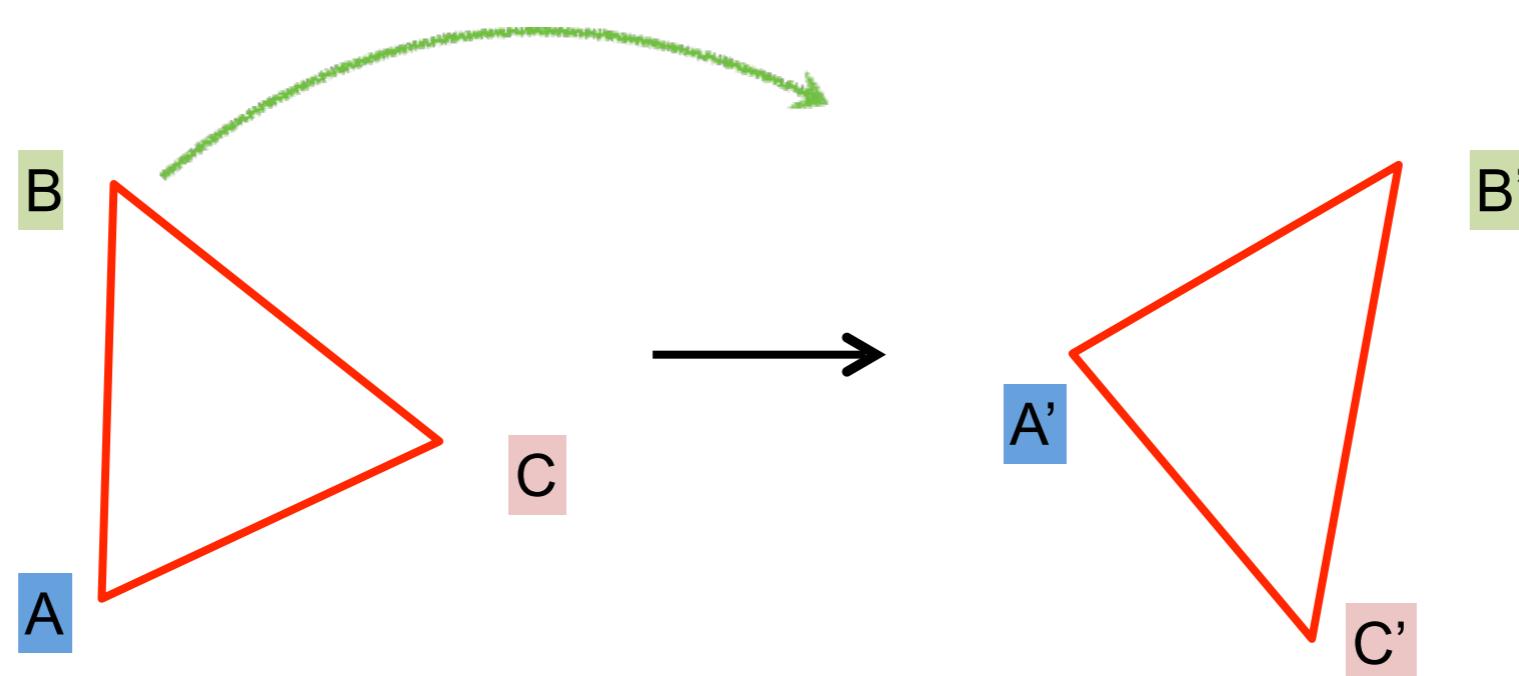


Can this be done with a euclidean / similarity / affine transformation?

Turns out any triangle can be warped to another triangle with an affine transform. Let's prove this useful fact.

Estimating transformations from point correspondences

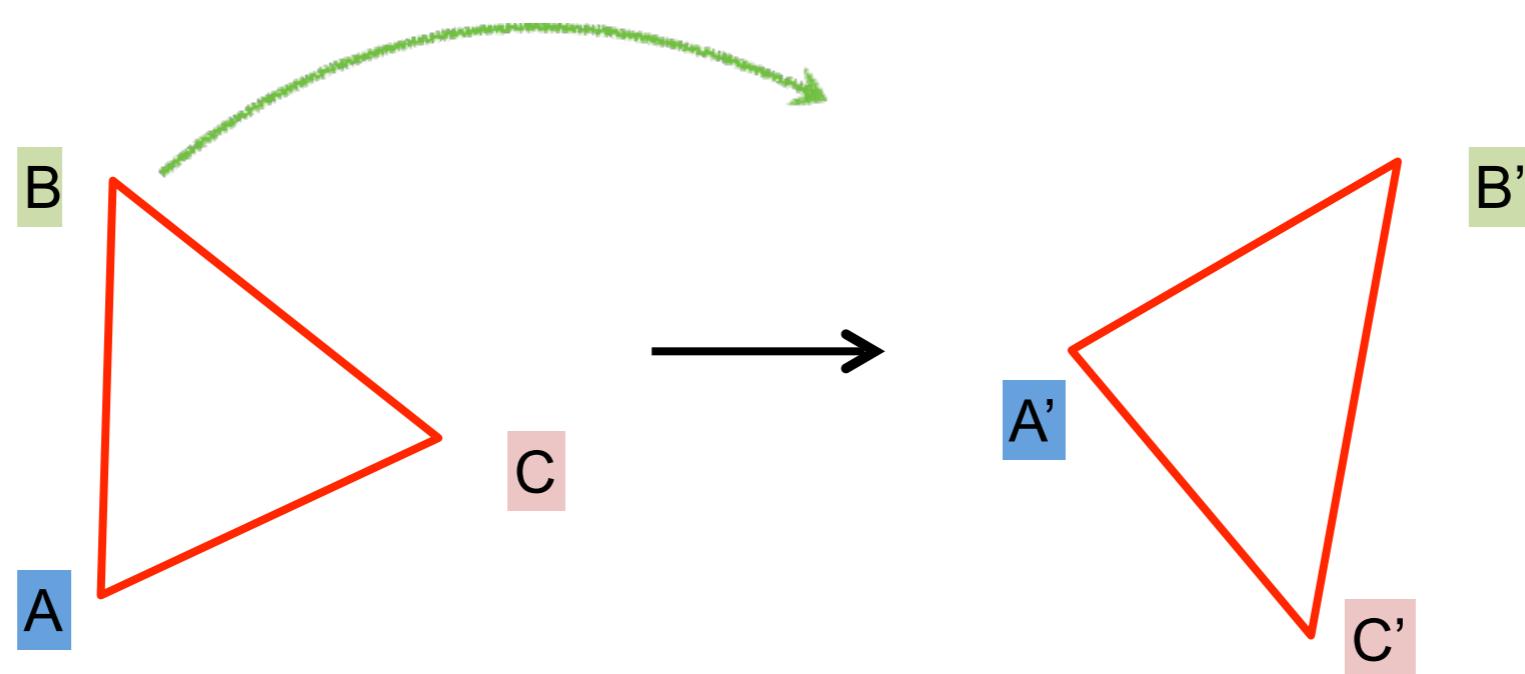
Suppose we have two triangles: ABC and A'B'C'.



$$\begin{bmatrix} A'_x \\ A'_y \\ 1 \end{bmatrix} = \begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} A_x \\ A_y \\ 1 \end{bmatrix}$$

Given similar equations for B and C, how to estimate a_1 through a_6 ?

Least Squares Error



Euclidean
(L2) norm

↑
predicted
location

↑
measured
location

$$\text{Error} = \|\text{Warp}(A, M) - A'\|^2 + \|\text{Warp}(B, M) - B'\|^2 + \|\text{Warp}(C, M) - C'\|^2$$

General form of linear least squares

(**Warning:** change of notation. \mathbf{x} is a vector of parameters!)

$$\begin{aligned} E_{\text{LLS}} &= \sum_i |a_i \mathbf{x} - b_i|^2 \\ &= \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \quad (\text{matrix form}) \end{aligned}$$

Determining unknown transformations

Affine transformation:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Why can we drop the last line?

Vectorize transformation parameters:
Stack equations from point correspondences:

$$\begin{bmatrix} x' \\ y' \\ x' \\ y' \\ \vdots \\ x' \\ y' \end{bmatrix} = \begin{bmatrix} x & y & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x & y & 1 \\ x & y & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x & y & 1 \\ \vdots & & & \vdots & & \\ x & y & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x & y & 1 \end{bmatrix} \begin{bmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \\ p_5 \\ p_6 \end{bmatrix}$$

$\underbrace{}$ $\underbrace{}$ $\underbrace{}$

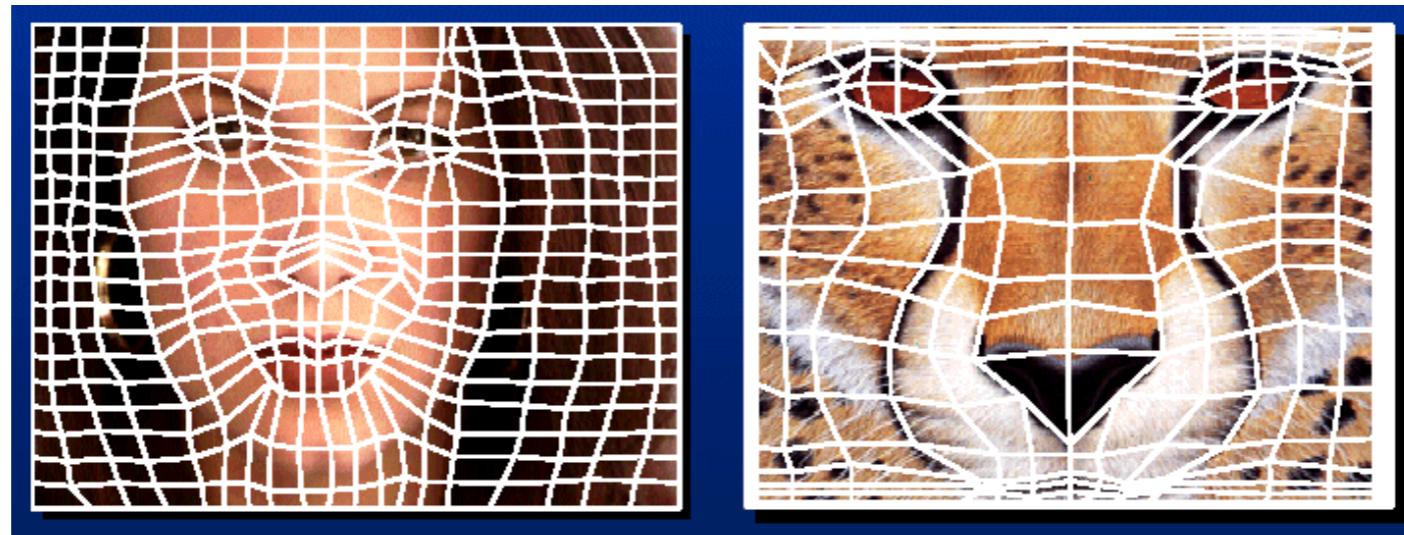
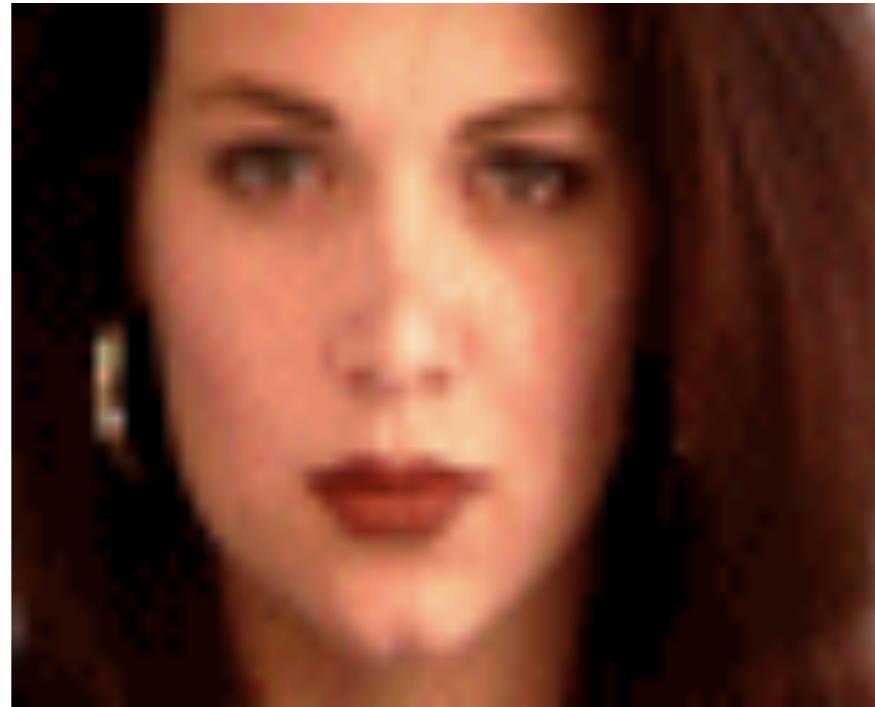
b A x

Notation in system form:

```
x = numpy.linalg.solve(A, b)
```

What's the minimum number of points needed to estimate an affine warp?

Image warping via *peicewise affine warps*



Piecewise affine warps (cut each quadrilateral into 2 triangles)

Example application: shape modeling



D'Arcy Thompson
“On Growth and Form” 1915

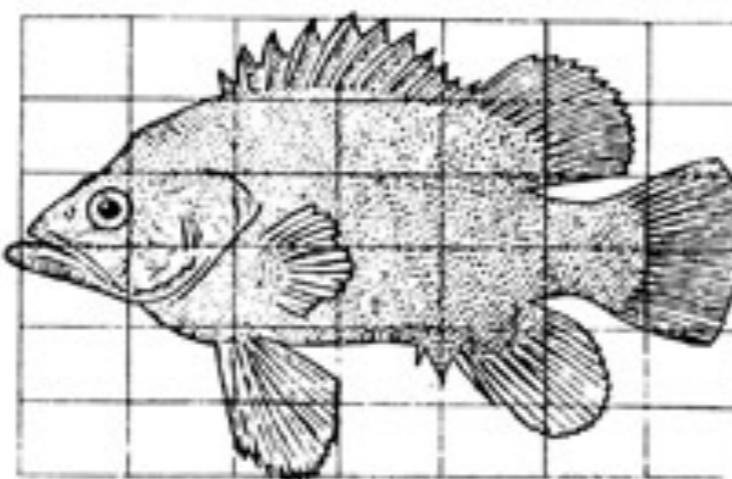


Fig. 150. *Polyprion*.

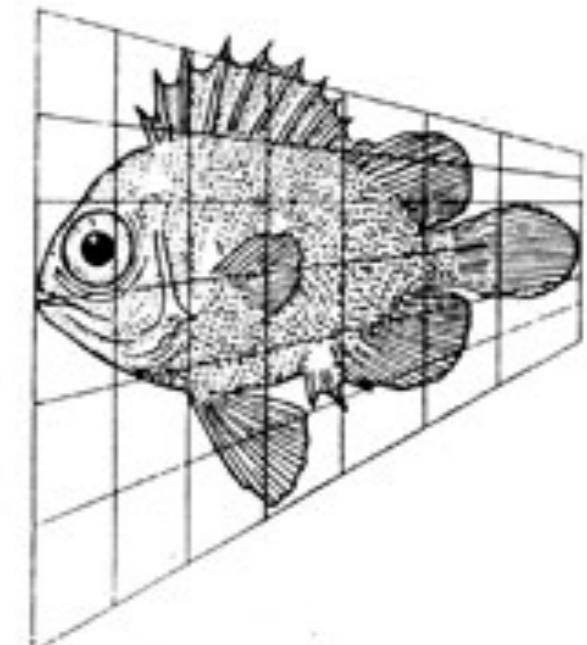


Fig. 151. *Pseudopriacanthus altus*.

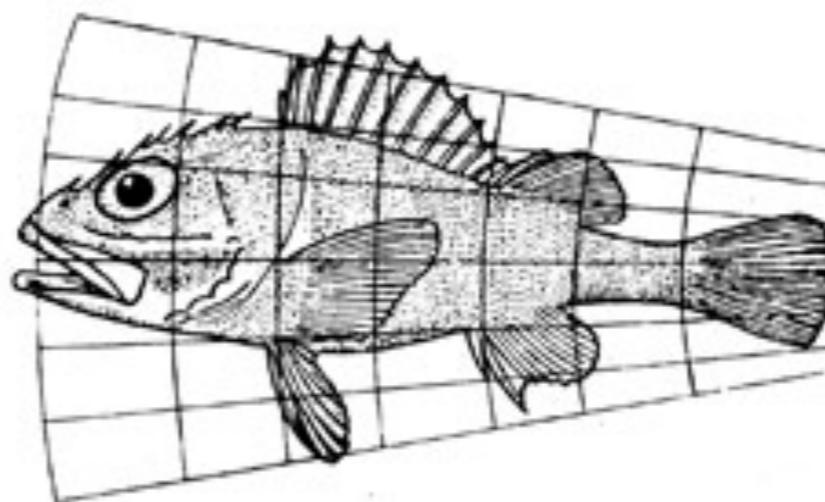


Fig. 152. *Scorpaena* sp.

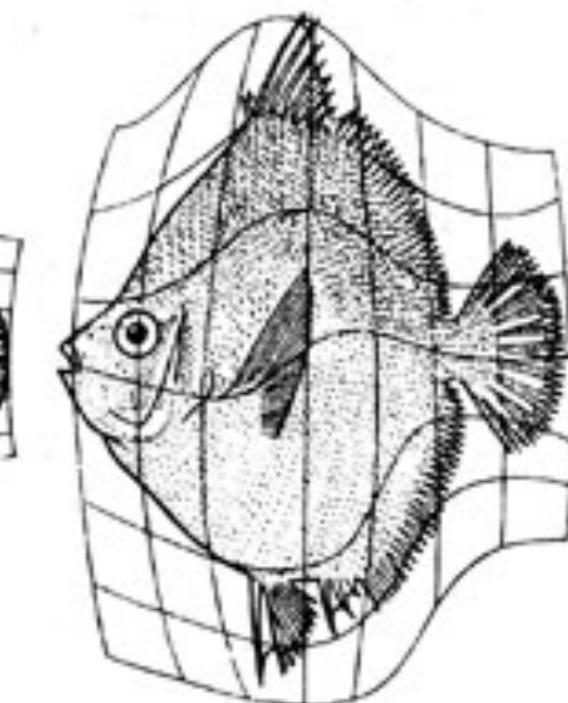
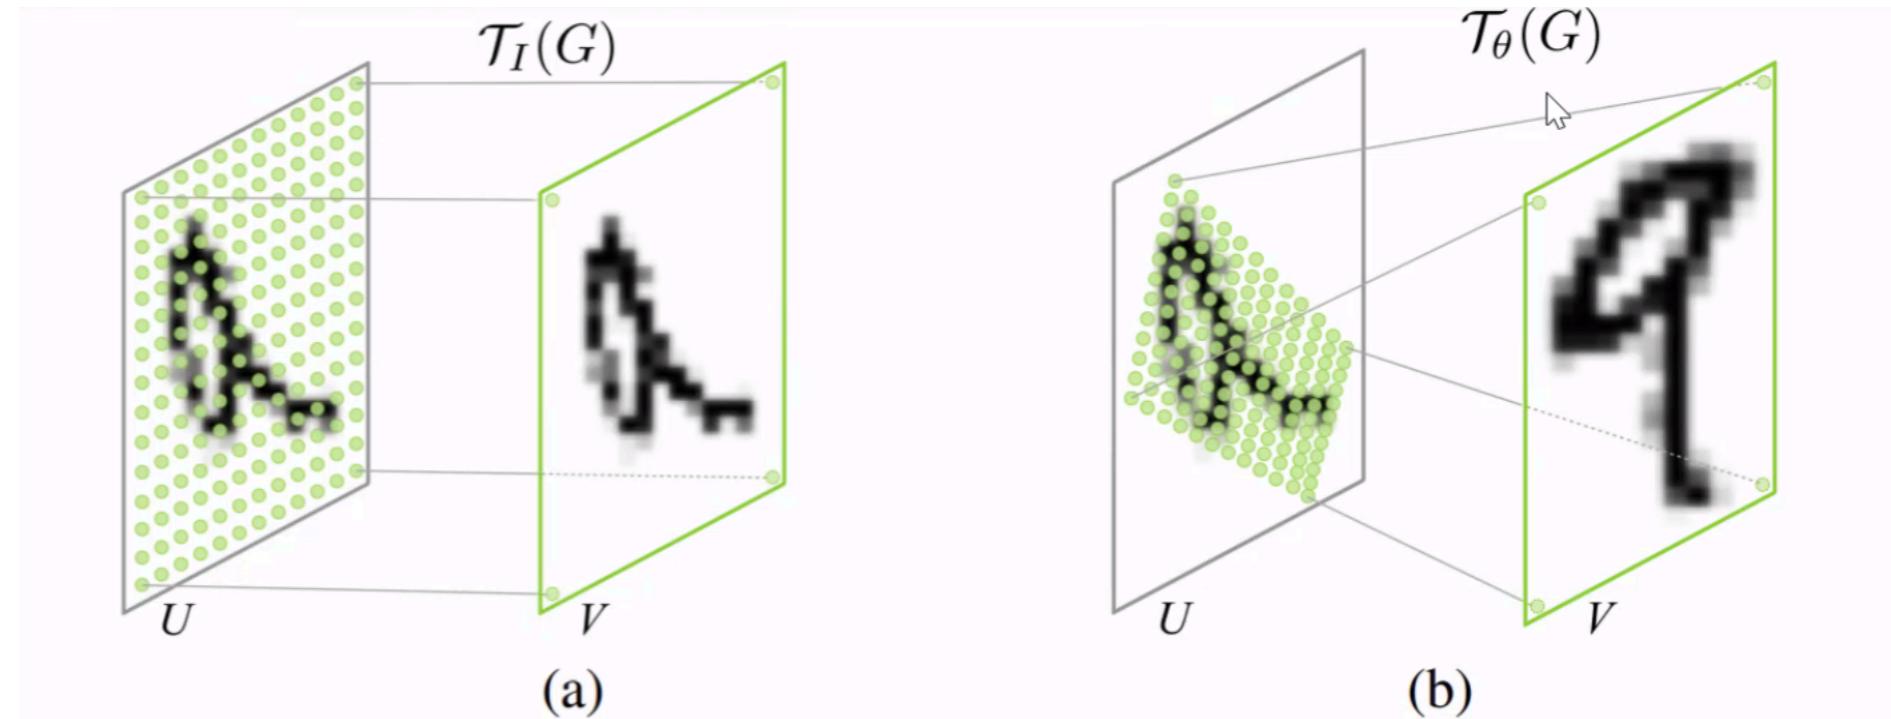
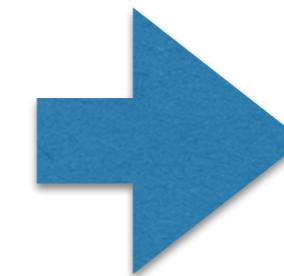
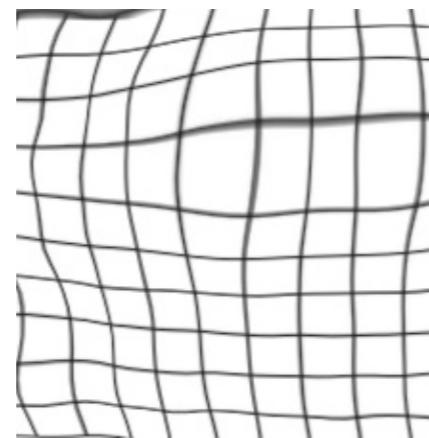
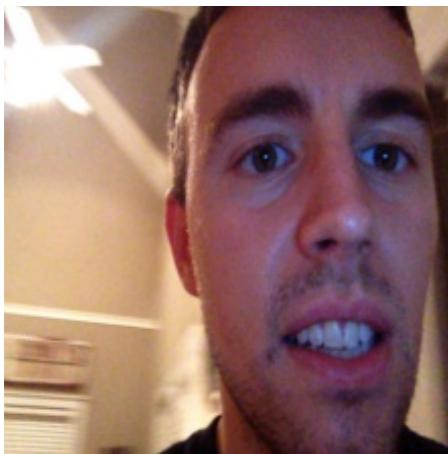


Fig. 153. *Antigonia capros*.

Application: spatial transformer networks



Neural networks that explicitly warp input images (or even feature maps)



classify eye_gaze direction

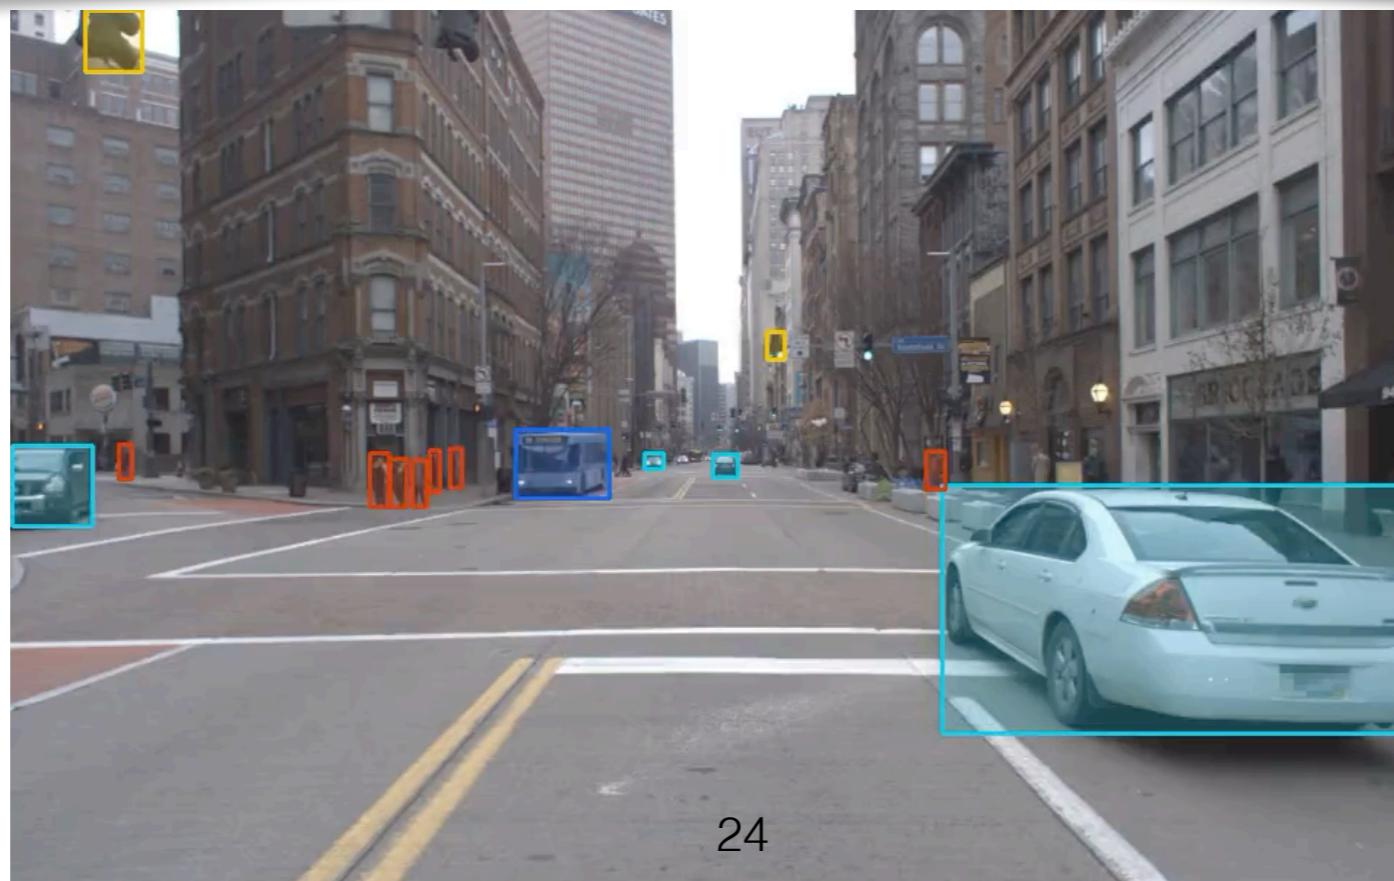
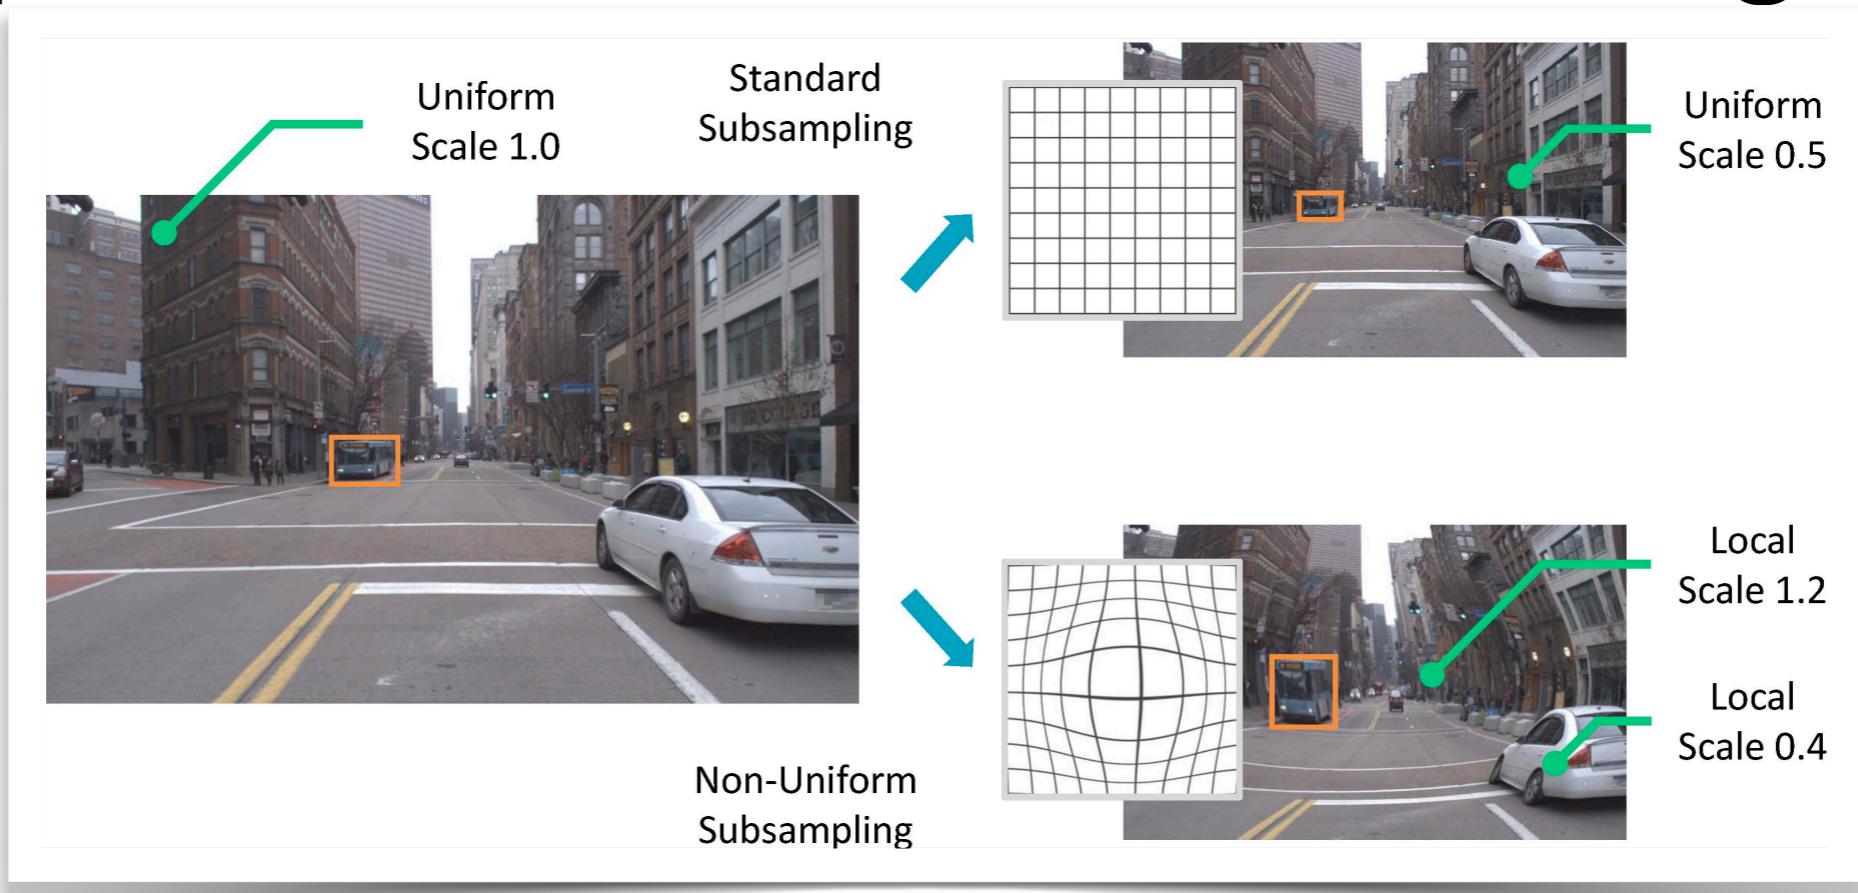
Spatial Transformer Networks

Max Jaderberg Karen Simonyan Andrew Zisserman Koray Kavukcuoglu
Google DeepMind, London, UK
{jaderberg,simonyan,zisserman,korayk}@google.com

Learning to Zoom: a Saliency-Based Sampling Layer for Neural Networks

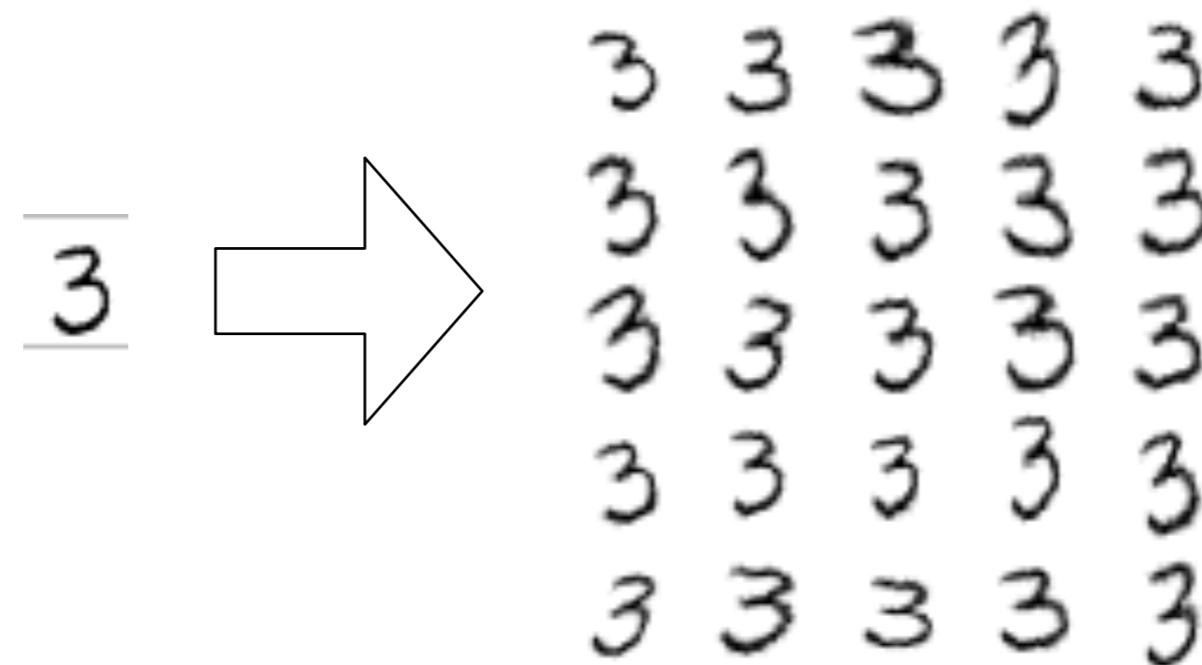
Adrià Recasens^{*1}, Petr Kellnhofer^{*1}, Simon Stent², Wojciech Matusik¹, and Antonio Torralba¹

Application: “attentional imaging”



ICCV 21

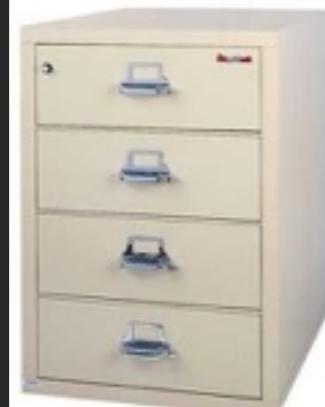
Application: data augmentation



Hypothesis: we'll see even more of this in the future!

Data augmentation with projective warps

Cabinet



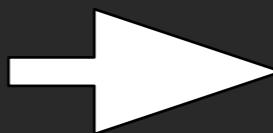
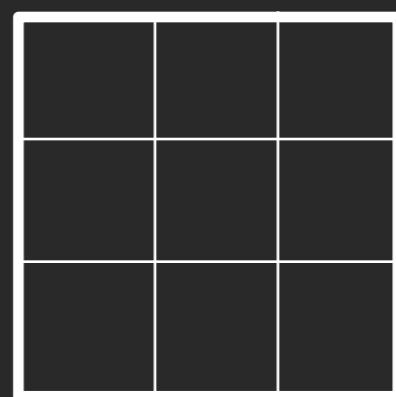
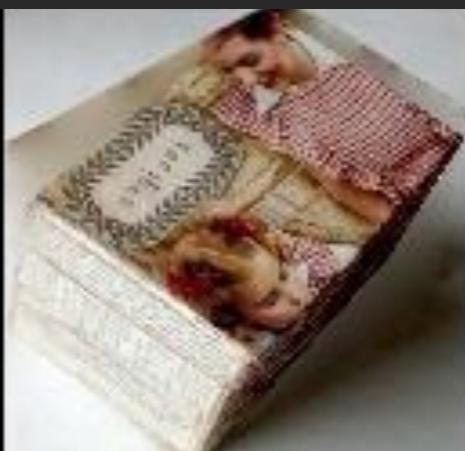
Cargo Container



Dryer

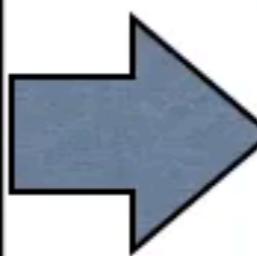


Cardboard Box

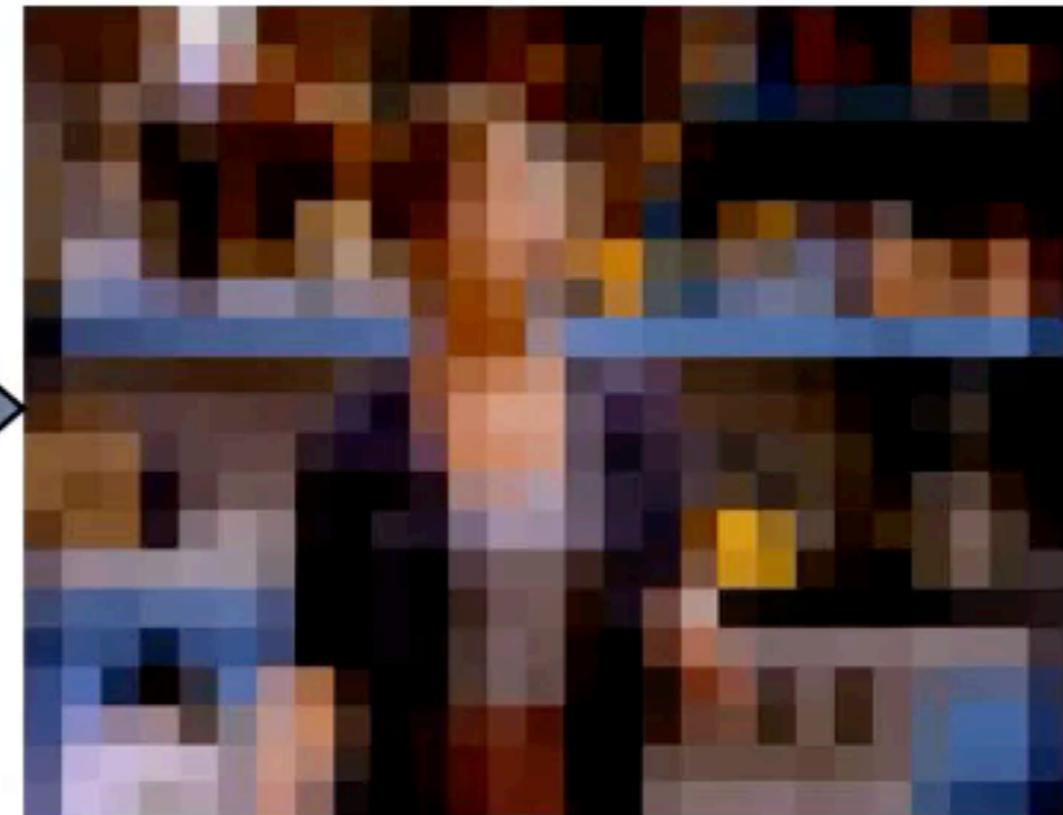


test video

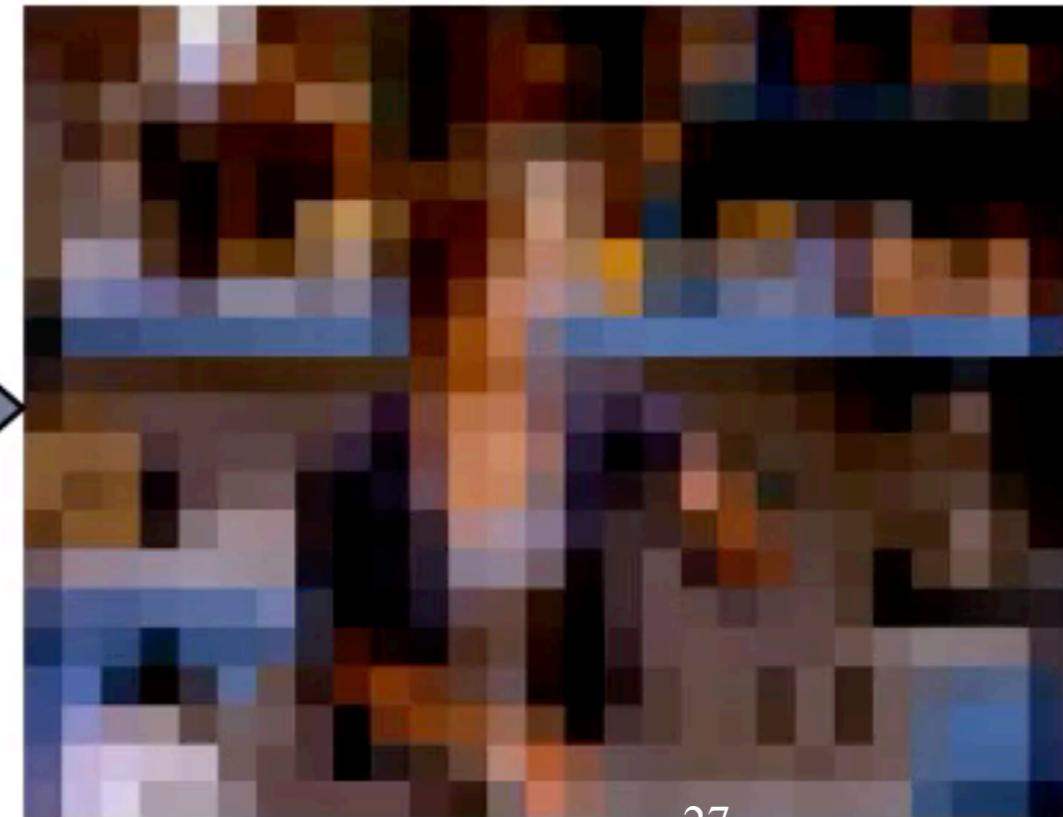
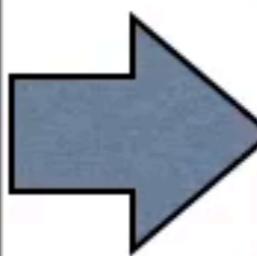
original



low resolution



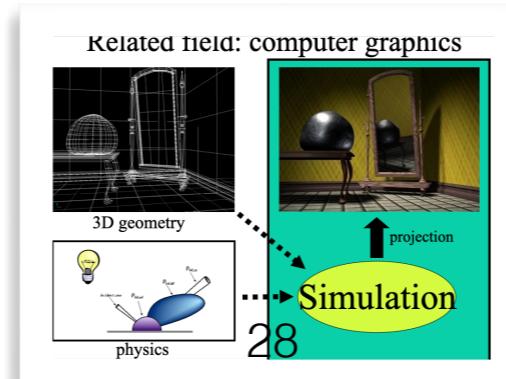
best match



Application: Tracking (HW2)!



Really instance of vision as inverse graphics or “analysis by synthesis”; interpret the world by *optimizing* for best fit of model

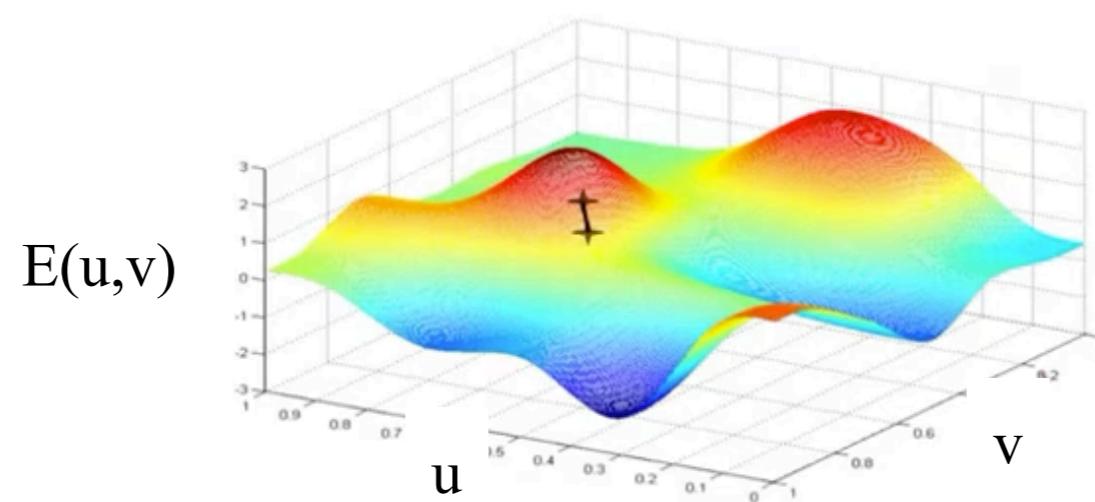


Lucas Kanade Alignment

$$\min_{u,v} E(u,v) \quad \text{where} \quad E(u,v) = \sum_{(x,y) \in W} [I(x+u, y+v) - T(x, y)]^2$$



template



What is the transformation that is being searched over?

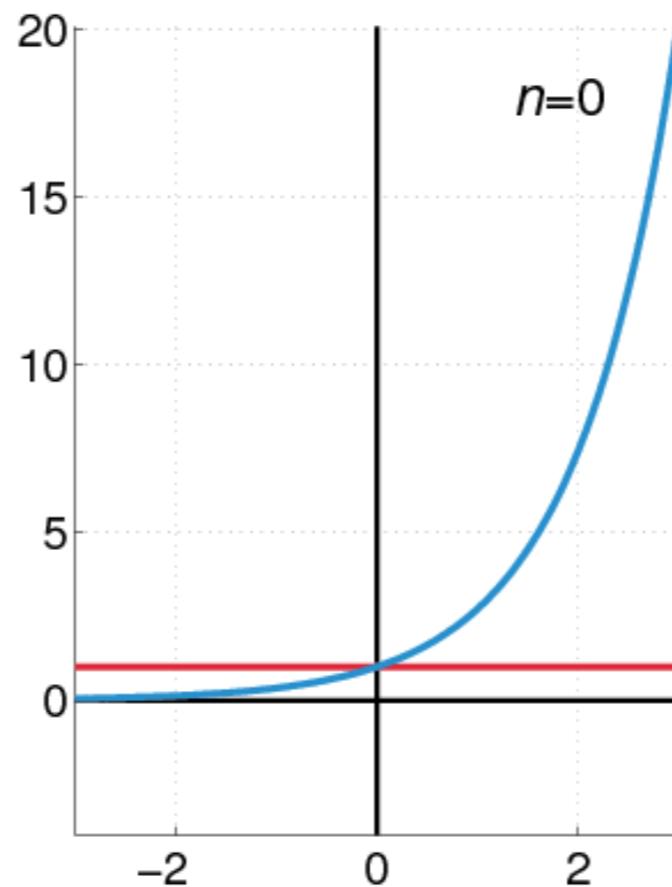
Andrew Ng

What kind of optimization problem is this?

Background: nonlinear optimization

Apply standard linear optimization tricks after using Taylor series approximation

$$f(u + \Delta u) = f(u) + \frac{\partial f(u)}{\partial u} \Delta u + \frac{\partial^2 f(u)}{\partial u^2} \Delta u^2 + \text{Higher Order Terms}$$



How do we generalize the above to a **multivariate** function $f(u,v)$?

Generalize first derivative to **gradient vector** and second derivative to **Hessian matrix**

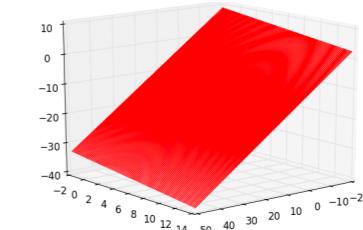
Background: nonlinear optimization

$$\min_{\mathbf{u}} f(\mathbf{u}) \quad \text{where } \mathbf{u} = (u, v)$$

1. First order method (gradient descent)

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \alpha \begin{bmatrix} \frac{\partial f}{\partial u} \\ \frac{\partial f}{\partial v} \end{bmatrix}$$

“fit a plane and take a step”

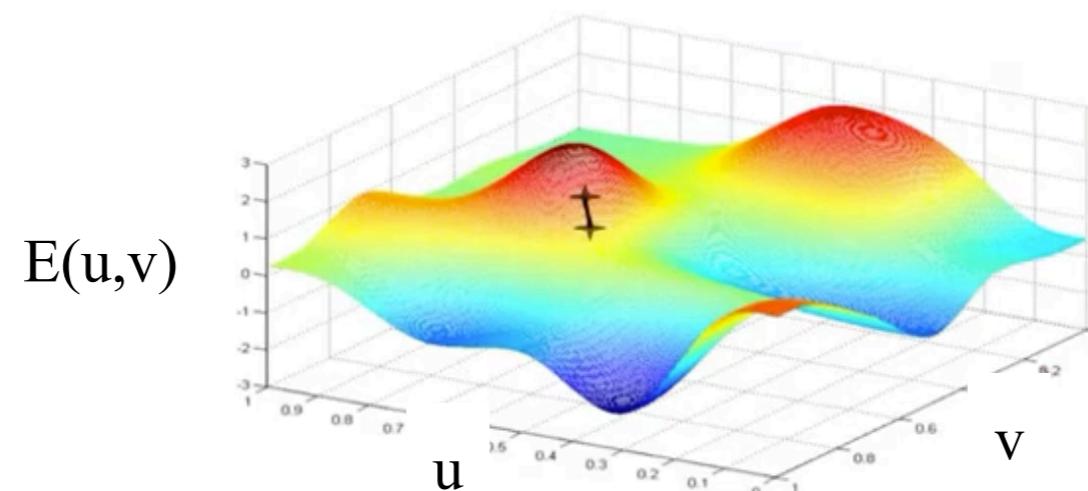
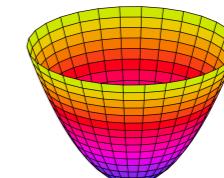


2. Second order method (Newton's method)

https://en.wikipedia.org/wiki/Newton's_method_in_optimization

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \begin{bmatrix} \frac{\partial^2 f}{\partial u \partial u} & \frac{\partial^2 f}{\partial u \partial v} \\ \frac{\partial^2 f}{\partial v \partial u} & \frac{\partial^2 f}{\partial v \partial v} \end{bmatrix}^{-1} \begin{bmatrix} \frac{\partial f}{\partial u} \\ \frac{\partial f}{\partial v} \end{bmatrix}$$

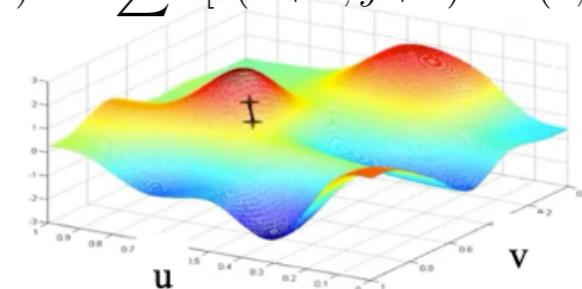
“fit a paraboloid and jump to minimum”



“Background”: nonlinear least squares

$$E(u, v) = \frac{1}{2} \left\| \begin{bmatrix} f(u, v) \\ g(u, v) \end{bmatrix} \right\|^2 = \frac{1}{2} \|\mathbf{f}(\mathbf{u})\|^2$$

$$E(u, v) = \sum [I(x + u, y + v) - T(x, y)]^2$$



1. First order method (gradient descent)

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \alpha \begin{bmatrix} \frac{\partial f}{\partial u} & \frac{\partial f}{\partial v} \\ \frac{\partial g}{\partial u} & \frac{\partial g}{\partial v} \end{bmatrix}^T \begin{bmatrix} f(u, v) \\ g(u, v) \end{bmatrix}$$

1.5 Gauss-Newton optimization

https://en.wikipedia.org/wiki/Gauss–Newton_algorithm

$$\begin{bmatrix} f(u + \Delta u, v + \Delta v) \\ g(u + \Delta u, v + \Delta v) \end{bmatrix} \approx \begin{bmatrix} f(u, v) \\ g(u, v) \end{bmatrix} + \begin{bmatrix} \frac{\partial f}{\partial u} & \frac{\partial f}{\partial v} \\ \frac{\partial g}{\partial u} & \frac{\partial g}{\partial v} \end{bmatrix} \begin{bmatrix} \Delta u \\ \Delta v \end{bmatrix}$$

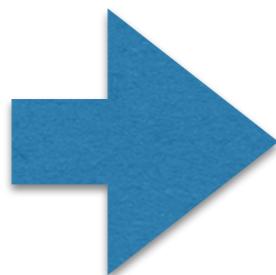
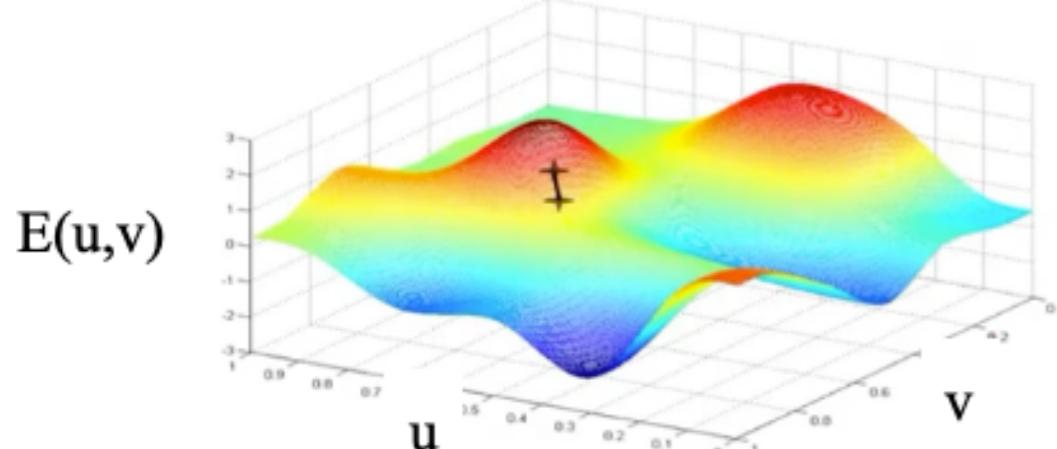
$$\mathbf{f}(\mathbf{u} + \Delta \mathbf{u}) \approx \mathbf{f}(\mathbf{u}) + \mathbf{J}(\mathbf{u}) \Delta \mathbf{u}$$

$$\mathbf{u} := \mathbf{u} + \arg \min_{\Delta \mathbf{u}} \frac{1}{2} \|\mathbf{f}(\mathbf{u}) + \mathbf{J}(\mathbf{u}) \Delta \mathbf{u}\|^2$$

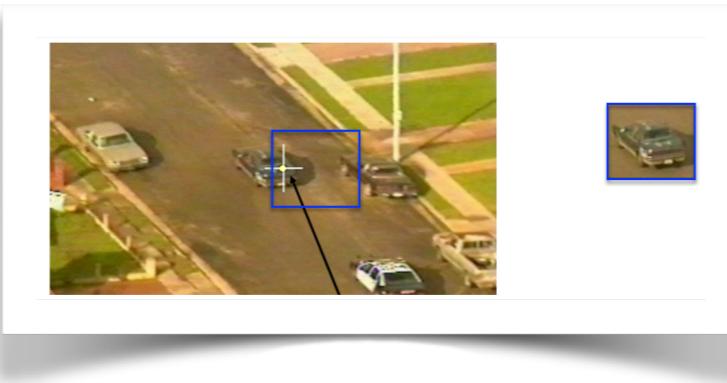
Take derivative of error wrt $\Delta u, \Delta v$ and set equal to 0;
get the performance of second-order optimizers with first-order compute!

Where we are headed...

$$\min_{u,v} \sum_{(x,y) \in W} [I(x+u, y+v) - T(x, y)]^2$$



Lucas Kanade alignment



$$\min_{u,v} E(u,v) = \sum_{(x,y) \in W} [I(x+u, y+v) - T(x, y)]^2$$

Start with initial guess of (u, v) . Linearize image about this guess

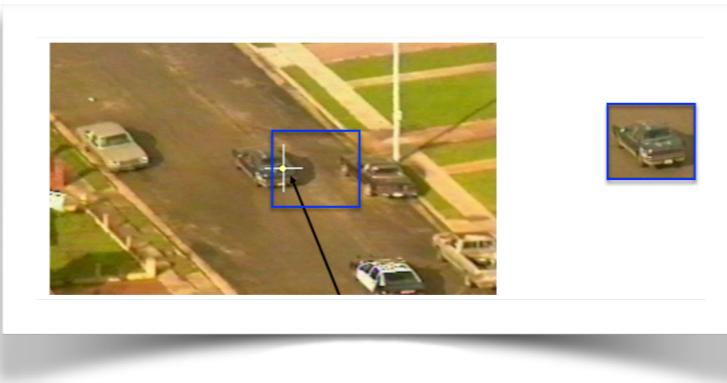
$$I(x + u + \Delta u, y + v + \Delta v) \approx I(x + u, y + v) + \frac{\partial I(x + u, y + v)}{\partial x} \Delta u + \frac{\partial I(x + u, y + v)}{\partial y} \Delta v$$

$$I(\mathbf{x} + \mathbf{u} + \Delta \mathbf{u}) \approx I(\mathbf{x} + \mathbf{u}) + \nabla I(\mathbf{x} + \mathbf{u})^T \Delta \mathbf{u}$$

Plug back into error term and solve for $\Delta \mathbf{u}$'s that minimizes the squared error

$$\min_{\Delta \mathbf{u}} \sum_{\mathbf{x}} [I(\mathbf{x} + \mathbf{u}) + \nabla I(\mathbf{x} + \mathbf{u})^T \Delta \mathbf{u} - T(\mathbf{x})]^2$$

Lucas Kanade alignment



$$\min_{\Delta \mathbf{u}} \sum_{\mathbf{x}} [I(\mathbf{x} + \mathbf{u}) + \nabla I(\mathbf{x} + \mathbf{u})^T \Delta \mathbf{u} - T(\mathbf{x})]^2$$

Treat \mathbf{u} as fixed wrt to optimization. Define $\tilde{I}(\mathbf{x}) = I(\mathbf{x} + \mathbf{u})$.

$$\min_{\Delta u, \Delta v} \sum_{x,y} [\tilde{I}(x,y) + \tilde{I}_x(x,y)\Delta u + \tilde{I}_y(x,y)\Delta v - T(x,y)]^2$$

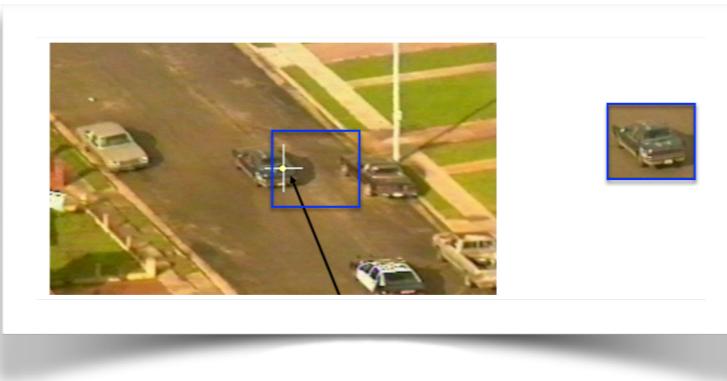
Define $D(x,y) = T(x,y) - \tilde{I}(x,y)$.

$$\min_{\Delta u, \Delta v} \sum_{x,y} [\tilde{I}_x(x,y)\Delta u + \tilde{I}_y(x,y)\Delta v - D(x,y)]^2$$

Take derivatives and set them equal to 0

$$\frac{\partial}{\partial \Delta u} = 2 \sum_{x,y} [\tilde{I}_x(x,y)\Delta u + \tilde{I}_y(x,y)\Delta v - D(x,y)] \tilde{I}_x(x,y) = 0$$

Lucas Kanade alignment



Form matrix equation

$$\begin{bmatrix} \cdot & \cdot \\ \cdot & \cdot \\ \cdot & \cdot \end{bmatrix}$$

$$\begin{bmatrix} \Delta u \\ \Delta v \end{bmatrix} =$$

$$\begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix}$$

Aside: how is the A matrix different than a Hessian matrix?

$$\left[\frac{\partial I(x,y)}{\partial x} \right] \left[\frac{\partial I(x,y)}{\partial x} \right] \neq \frac{\partial^2 I(x,y)}{\partial x \partial x}$$

“Ax=b”

“x=A⁻¹b”

Lucas Kanade Alignment

$$\min_{u,v} \sum_{(x,y) \in W} [I(x+u, y+v) - T(x, y)]^2$$

initialization



template



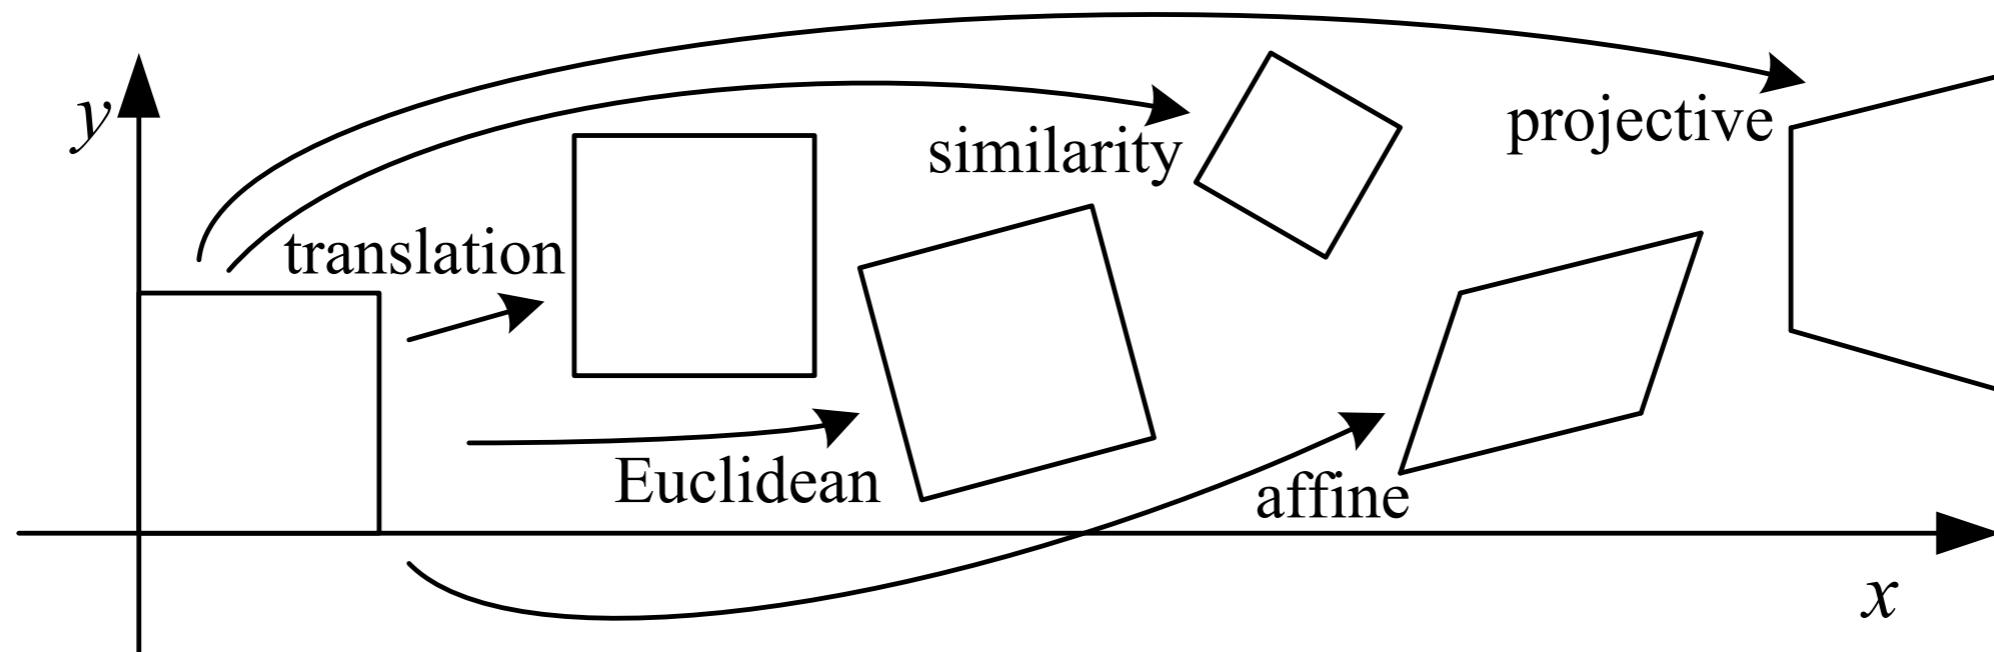
HW: Tracking by iterative template alignment



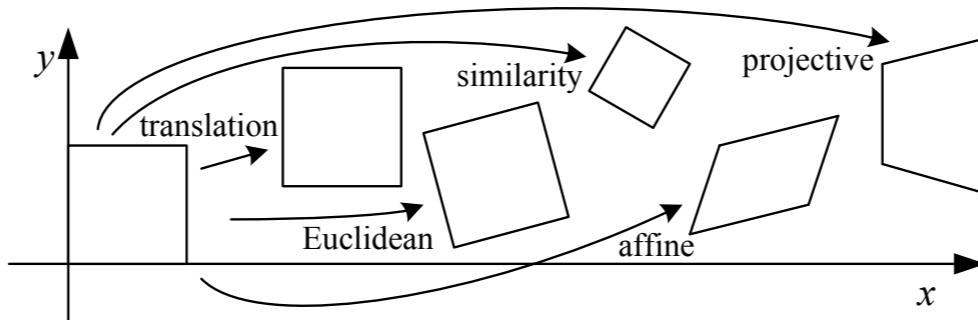
Start with template from first frame and repeat:

1. Align template to new frame with Lucas Kanade
2. Update template with new frame (or not?)

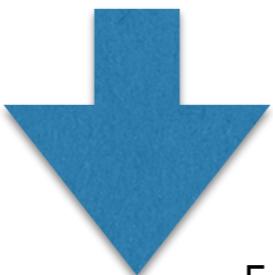
What about other warps?



Replace $[u, v]$ with a warping function $W(x, y; p)$



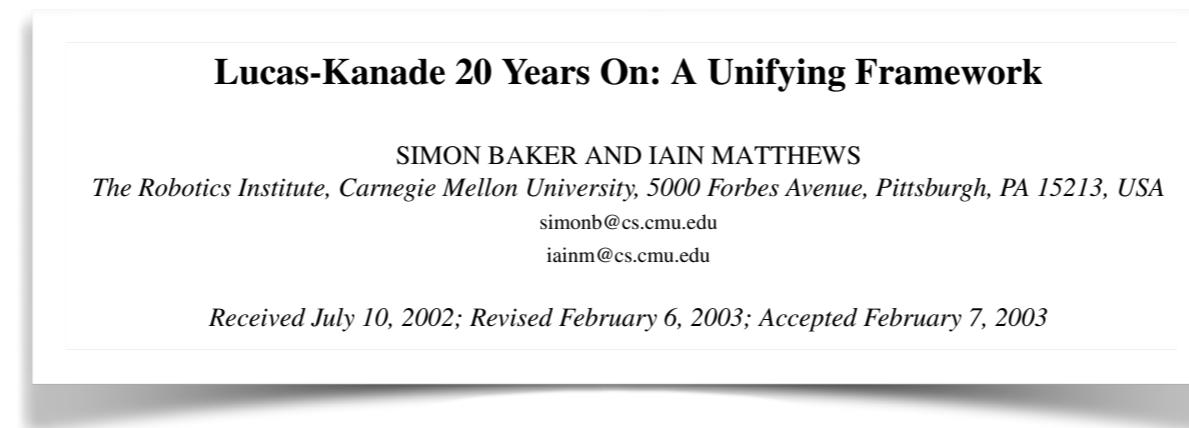
$$E(\mathbf{p}) = \sum_{\mathbf{x}} [I(W(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$



e.g., for translation warps: $W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} 1 & 0 & p_1 \\ 0 & 1 & p_2 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$

e.g., for affine warps: $W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$

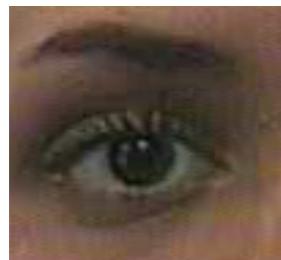
Seminal reference



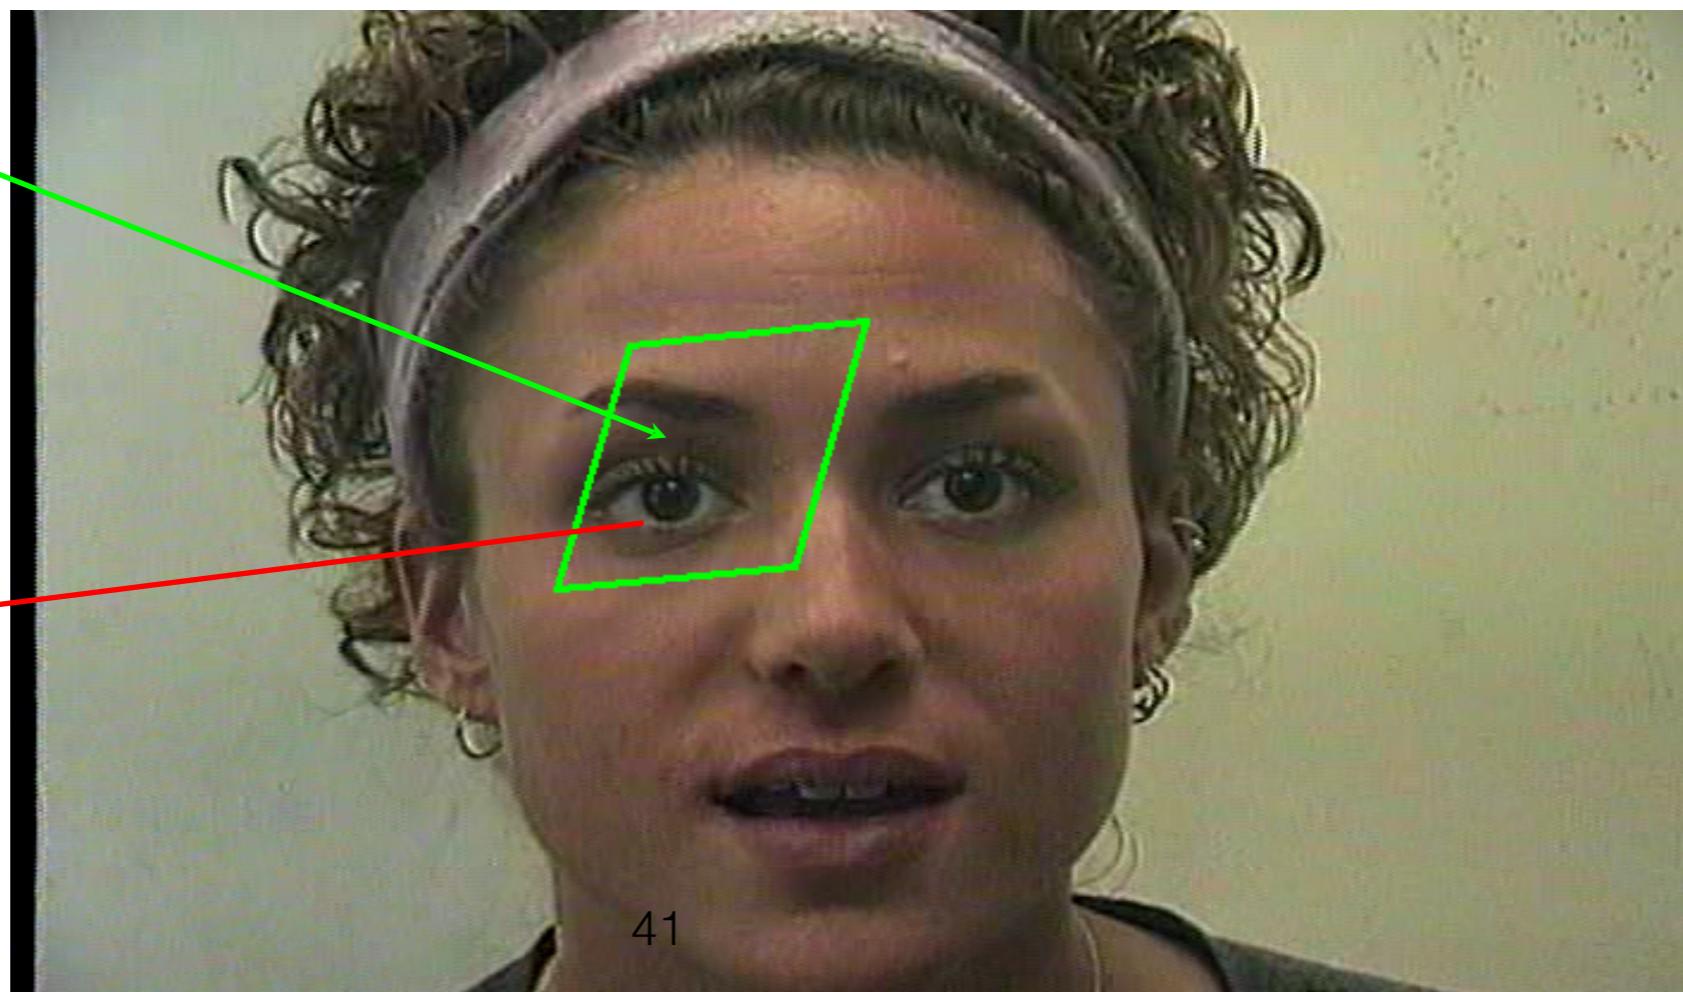
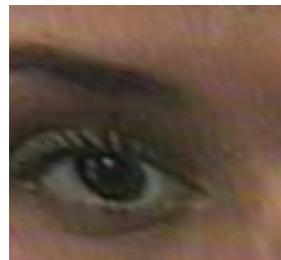
(will follow notation)

Warp image to get $I(W(x; p))$; then compute $[I(W(x; p)) - T(x)]^2$

Template, $T(x)$



Warped, $I(W(x;p))$



Nonlinear least squares

$$E(\mathbf{p}) = \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Apply taylor-series expansion to vector-valued function \mathbf{W}
(for simplicity, assume 2D rotation)

$$\mathbf{W}(\mathbf{x}; p) = \begin{bmatrix} \cos(p) & -\sin(p) \\ \sin(p) & \cos(p) \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

$$W_x = \cos(p)x - \sin(p)y$$

$$W_y = \sin(p)x + \cos(p)y$$

Make use of chain rule to compute derivative of image wrt $p = \theta$:

$$\frac{\partial I(W(\mathbf{x}; p))}{\partial p} = \frac{\partial I(W(\mathbf{x}; p))}{\partial x} \frac{\partial W_x(\mathbf{x}; p)}{\partial p} + \frac{\partial I(W(\mathbf{x}; p))}{\partial y} \frac{\partial W_y(\mathbf{x}; p)}{\partial p}$$

Taylor expansion of warped image

1-parameter warp (e.g., rotation)

$$I(W(\mathbf{x}; p + \Delta p)) \approx I(W(\mathbf{x}; p)) + \begin{bmatrix} \frac{\partial I(W(\mathbf{x}; p))}{\partial x} & \frac{\partial I(W(\mathbf{x}; p))}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial W_x(\mathbf{x}; p)}{\partial p} \\ \frac{\partial W_y(\mathbf{x}; p)}{\partial p} \end{bmatrix} \Delta p$$

current warped image current warped image gradient jacobian of warp parameter update

multi-parameter warp (e.g., affine)

$$I(W(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) \approx I(W(\mathbf{x}; \mathbf{p})) + \begin{bmatrix} \frac{\partial I(W(\mathbf{x}; \mathbf{p}))}{\partial x} & \frac{\partial I(W(\mathbf{x}; \mathbf{p}))}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial W_x(\mathbf{x}; \mathbf{p})}{\partial p_1} & \frac{\partial W_x(\mathbf{x}; \mathbf{p})}{\partial p_2} \\ \frac{\partial W_y(\mathbf{x}; \mathbf{p})}{\partial p_1} & \frac{\partial W_y(\mathbf{x}; \mathbf{p})}{\partial p_2} \end{bmatrix} \begin{bmatrix} \Delta p_1 \\ \Delta p_2 \end{bmatrix}$$

$\nabla I(W(\mathbf{x}; \mathbf{p})))$ $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$ $\Delta \mathbf{p}$

Example: jacobian of affine warp

$$W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} 1 + p_1 & p_3 & p_5 \\ p_2 & 1 + p_4 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \rightarrow \quad \begin{aligned} W_x &= (1 + p_1)x + p_3y + p_5 \\ W_y &= p_2x + (1 + p_4)y + p_6 \end{aligned}$$

$$\frac{\partial W}{\partial \mathbf{p}} = \begin{bmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \cdots \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \cdots \end{bmatrix} \quad \rightarrow \quad \frac{\partial W}{\partial \mathbf{p}} = \begin{bmatrix} x & 0 & y & 0 & 1 & 0 \\ 0 & x & 0 & y & 0 & 1 \end{bmatrix}$$

Above affine parameterization is better conditioned for optimization because all-zero parameters default to the identity transformation

Aside: low-dimensional parameterizations of warps

$$W(\mathbf{x}; p)$$



Figure 1: Example face image annotated with landmarks

Active Shape Models —Their Training and Application

T. F. COOTES, C. J. TAYLOR, D. H. COOPER, AND J. GRAHAM*

Department of Medical Biophysics, University of Manchester, Oxford Road, Manchester M13 9PT, England

Received July 29, 1992; accepted April 12, 1994



Apply SVD (or PCA) to collection of (x,y) facial landmarks to learn *eigenshapes*

$$\mathbf{s} = [x_1 \quad y_2 \quad x_2 \quad y_2 \dots]$$

$$\mathbf{s} = \mathbf{s}_0 + p_1 \mathbf{s}_1 + p_2 \mathbf{s}_2 + \dots$$

$$Pr(\mathbf{s}) = N(\mathbf{s}; \mu, \Sigma)$$

Back to the big-picture

$$\begin{aligned} & \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2 \\ & \approx \sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2 \end{aligned}$$

Gradient Descent Solution

Least squares problem: Minimize to solve for $\Delta\mathbf{p}$

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Solution,

$$\Delta \mathbf{p} = H^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

Gradient

Approx hessian

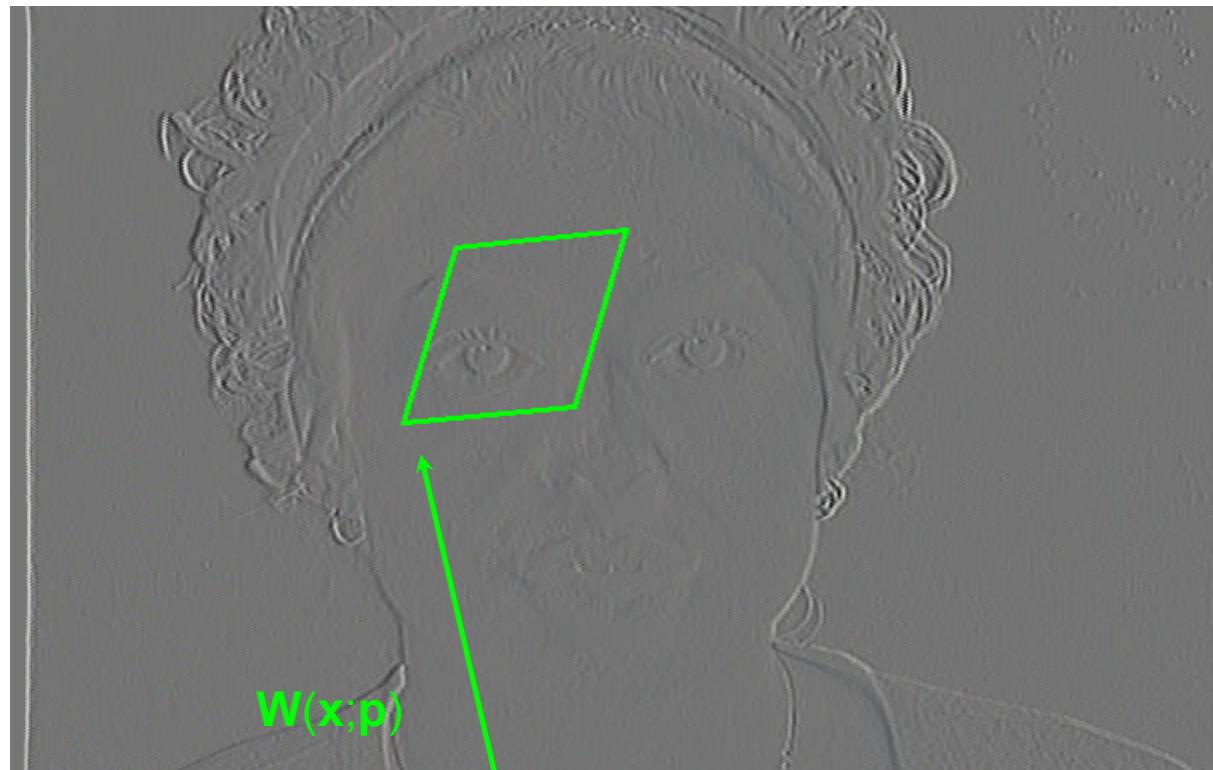
$$H = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

Error Image

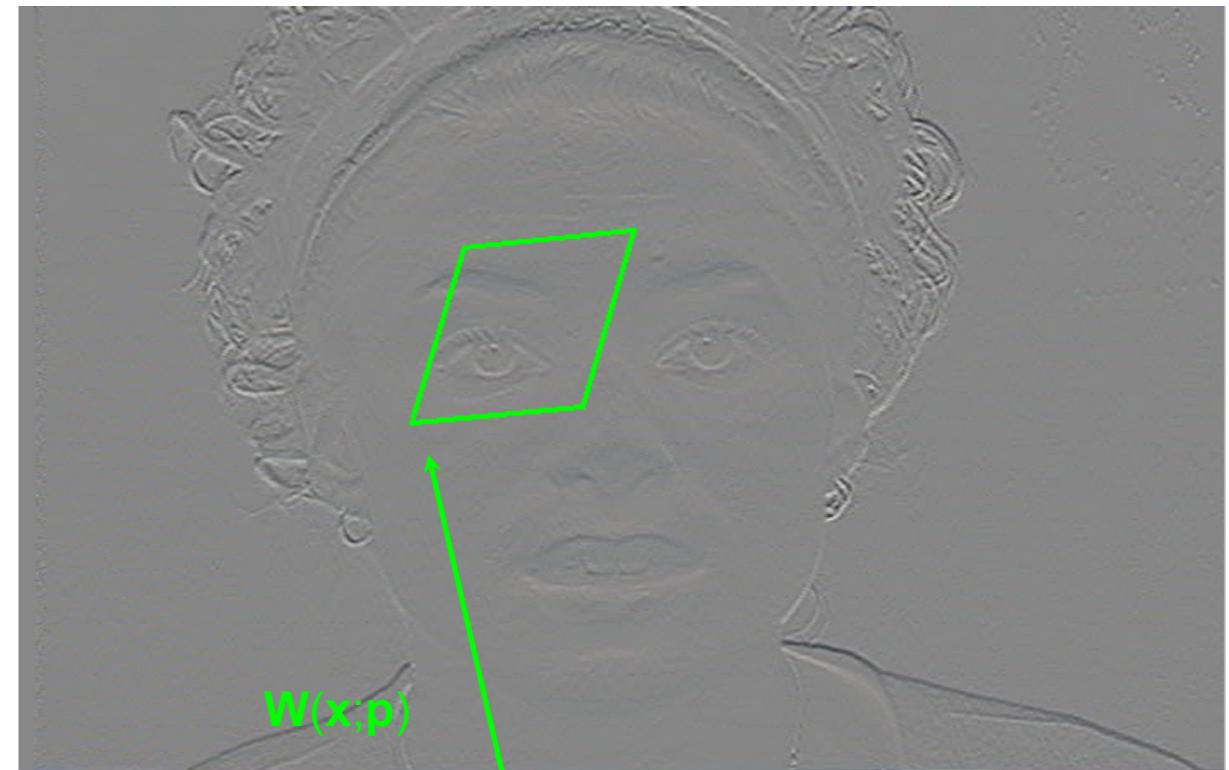
Jacobian

Gradient Images

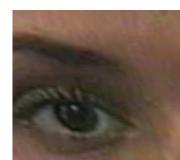
- Compute image gradient ∇I



∇I_x



$I(W(x;p))$



∇I_y



$$W([x, y]; P) = \begin{pmatrix} 1 + p_1 & p_3 & p_5 \\ p_2 & 1 + p_4 & p_6 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

Jacobian

- Compute Jacobian

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{pmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \dots & \frac{\partial W_x}{\partial p_n} \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \dots & \frac{\partial W_y}{\partial p_n} \end{pmatrix}$$

(e.g., for affine warps)

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{pmatrix} (1 + p_1) \cdot x & + & p_3 \cdot y & + & p_5 \\ p_2 \cdot x & + & (1 + p_4) \cdot y & + & p_6 \end{pmatrix} = \begin{pmatrix} 1 + p_1 & p_3 & p_5 \\ p_2 & 1 + p_4 & p_6 \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{pmatrix} x & 0 & y & 0 & 1 & 0 \\ 0 & x & 0 & y & 0 & 1 \end{pmatrix} = \begin{img alt="A 6x6 grid showing the Jacobian matrix for an affine warp. The grid is divided into 6 columns. The first column has values x, 0, y, 0, 1, 0. The second column has values 0, x, 0, y, 0, 1. The third column has values x, 0, y, 0, 1, 0. The fourth column has values 0, y, 0, 1, 0, x. The fifth column has values y, 0, 1, 0, x, 0. The sixth column has values 0, 1, 0, x, 0, y. The grid shows a checkerboard pattern of black and white squares, with some squares being gray to represent intermediate values." data-bbox="525 625 840 754}\end{img}$$

Lucas-Kanade Algorithm

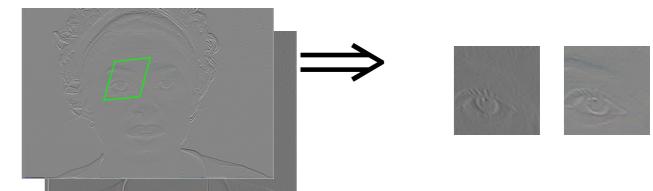
1. Warp I with $\mathbf{W}(\mathbf{x}; \mathbf{p}) \Rightarrow I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$



1. Compute error image $T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$



1. Warp image ∇I with $\mathbf{W}(\mathbf{x}; \mathbf{p})$



1. Evaluate Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$



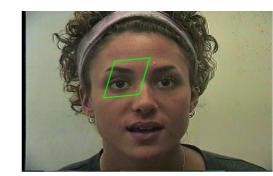
1. Compute Approx Hessian

$$H = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

1. Compute $\Delta \mathbf{p}$

$$\Delta \mathbf{p} = \sum_{\mathbf{x}} H^{-1} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

1. Update parameters $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$



Important extensions

Lucas-Kanade 20 Years On: A Unifying Framework

SIMON BAKER AND IAIN MATTHEWS

The Robotics Institute, Carnegie Mellon University, 5000 Forbes Avenue, Pittsburgh, PA 15213, USA

simonb@cs.cmu.edu

iainm@cs.cmu.edu

image



template



Key insight: exploit flexibility of warping the image *or* warping the template

Overview

Additive warp: $\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$ $\mathbf{p} \leftarrow \mathbf{p} + \Delta\mathbf{p}$.

Compositional warps: $\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$ $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p})$,

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{pmatrix} (1 + p_1) \cdot x & + & p_3 \cdot y & + & p_5 \\ p_2 \cdot x & + & (1 + p_4) \cdot y & + & p_6 \end{pmatrix} \quad \text{replace "x" with "(1 + \Delta p_1)x + \Delta p_3y + \Delta p_5" and "y" with ...}$$

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p}) =$$

$$\begin{pmatrix} (1 + p_1) \cdot ((1 + \Delta p_1) \cdot x + \Delta p_3 \cdot y + \Delta p_5) + p_3 \cdot (\Delta p_2 \cdot x + (1 + \Delta p_4) \cdot y + \Delta p_6) + p_5 \\ p_2 \cdot ((1 + \Delta p_1) \cdot x + \Delta p_3 \cdot y + \Delta p_5) + (1 + p_4) \cdot (\Delta p_2 \cdot x + (1 + \Delta p_4) \cdot y + \Delta p_6) + p_6 \end{pmatrix}$$

Overview

Additive warp: $\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$ $\mathbf{p} \leftarrow \mathbf{p} + \Delta\mathbf{p}$.

Compositional warp: $\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$ $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p})$,

Inverse compositional warp: $\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}))]^2$ $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p})^{-1}$.

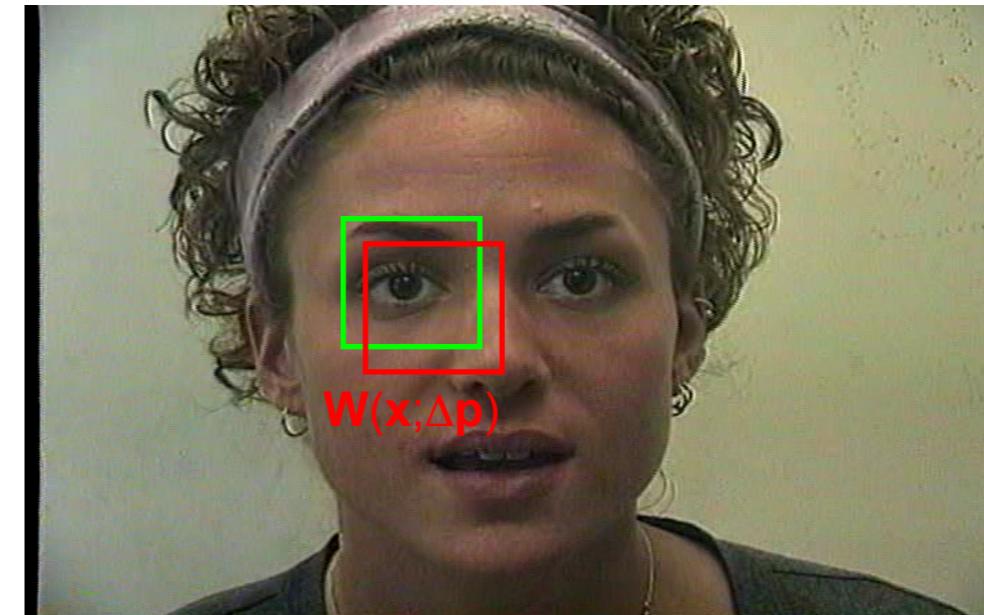
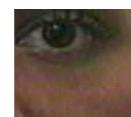


Key idea: compute update on template (not image). But apply *inverse* update to image. Because the template is never warped, its Jacobian and approxHessian (and its inverse!) can be precomputed.

Forward and Inverse Compositional

- Forwards compositional

$T(x)$

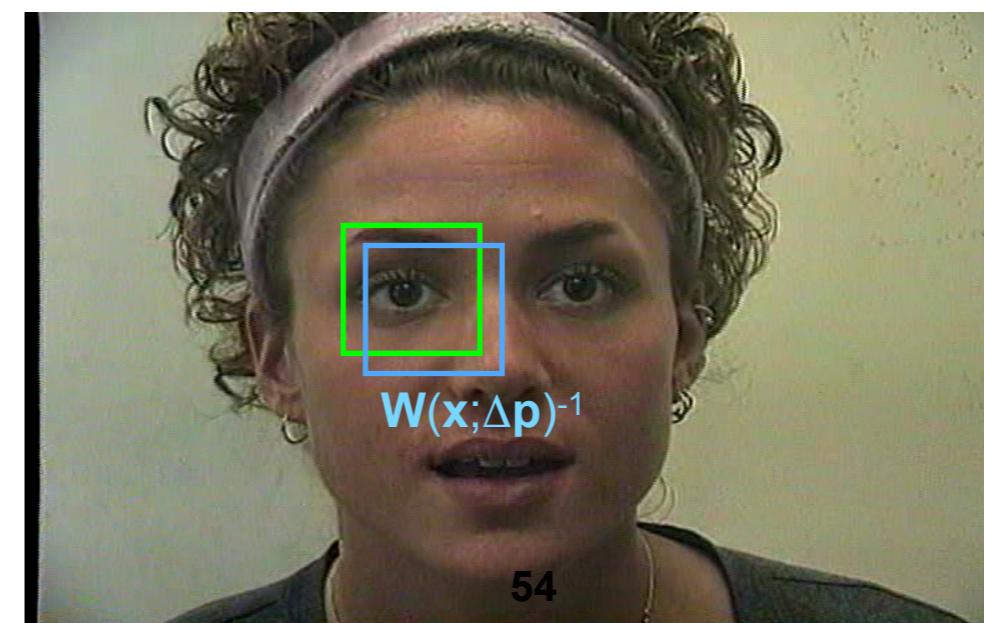


$I(x)$

- Inverse compositional



$W(x; \Delta p)$



54

$I(x)$

Inverse Compositional

Minimize: $\sum_{\mathbf{x}} [T(\mathbf{W}(\mathbf{x}; \Delta \mathbf{p})) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]^2$



Taylor approx

of template: $T(\mathbf{W}(\mathbf{x}; \Delta \mathbf{p})) \approx T(\mathbf{W}(\mathbf{x}; \mathbf{0})) + \nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p}$

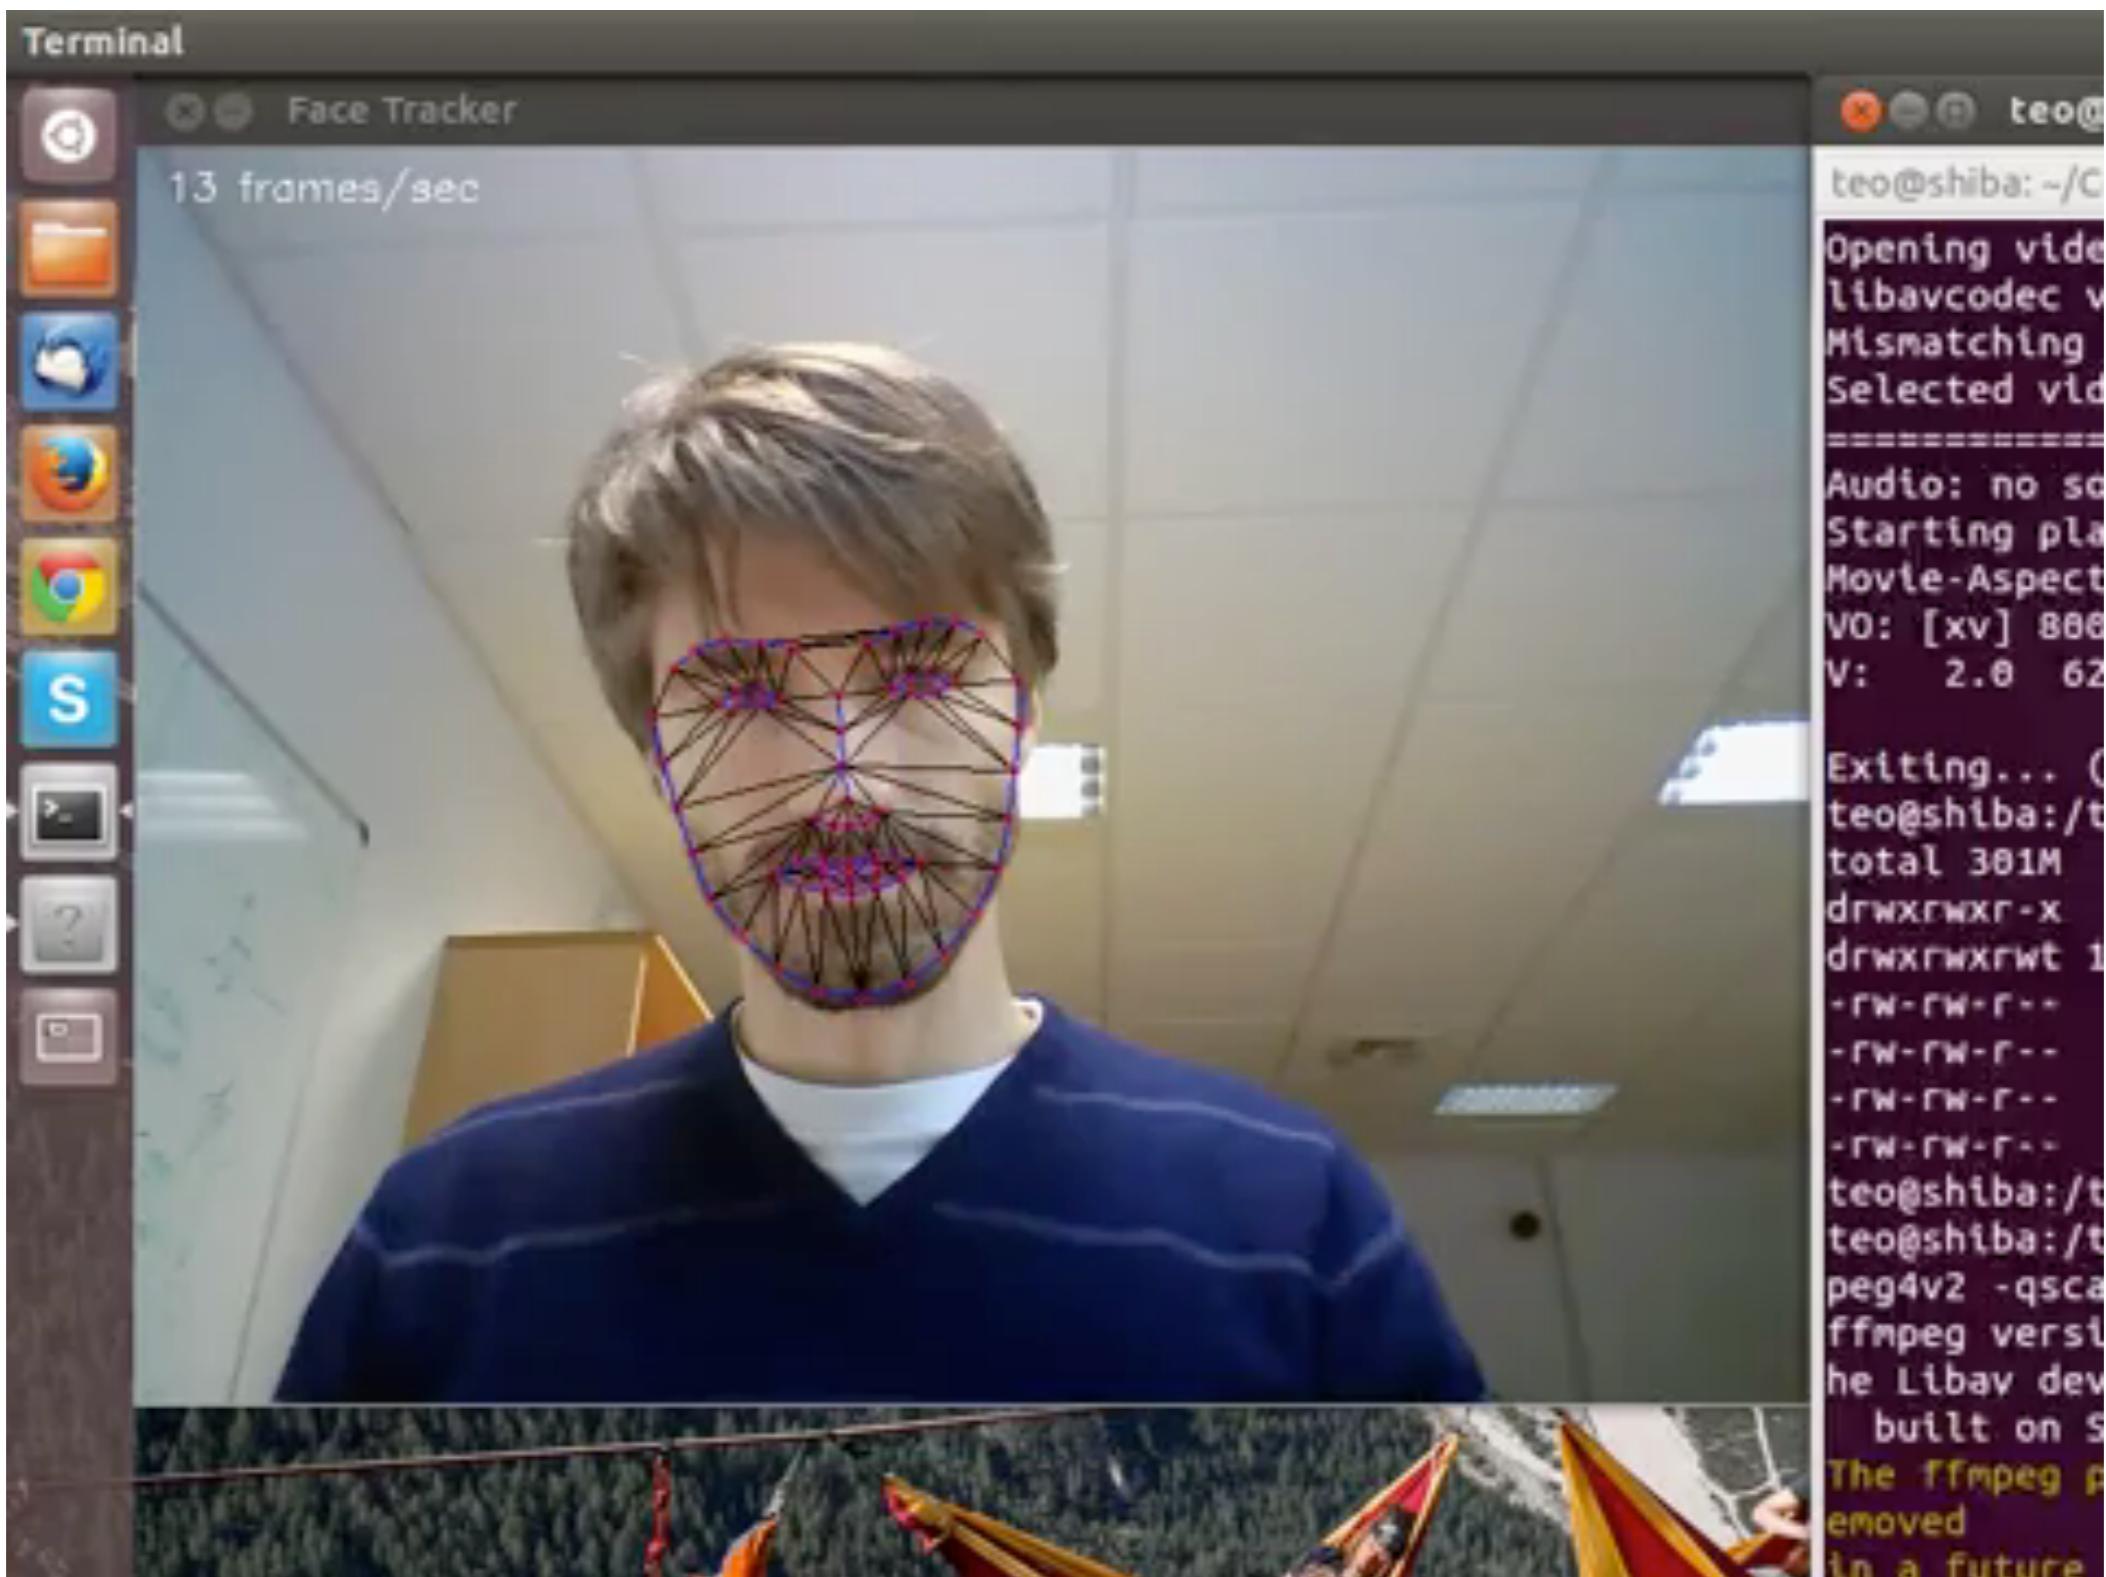
$$\text{Soln: } \Delta \mathbf{p} = - \sum_{\mathbf{x}} H^{-1} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

$$H = \sum_{\mathbf{x}} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

Precompute

Update: $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta \mathbf{p})^{-1}$

Peicewise affine-tracking



Contemporary extensions

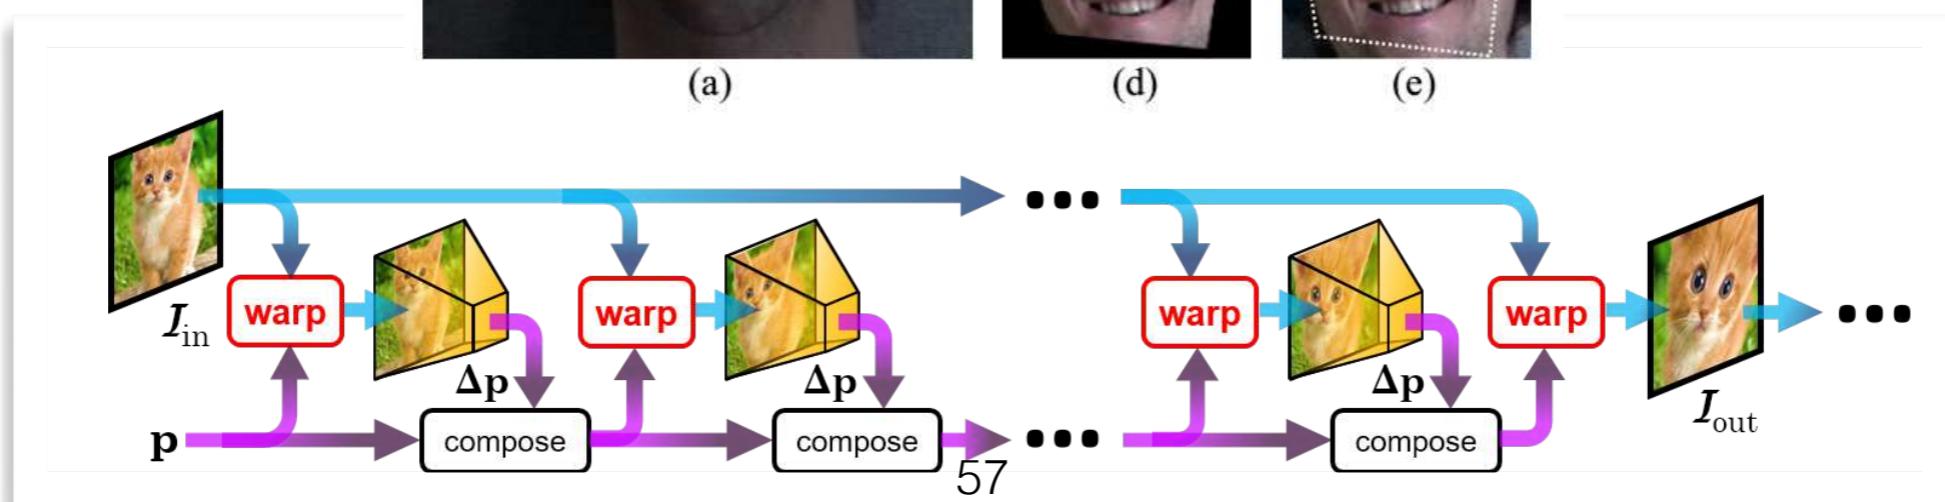
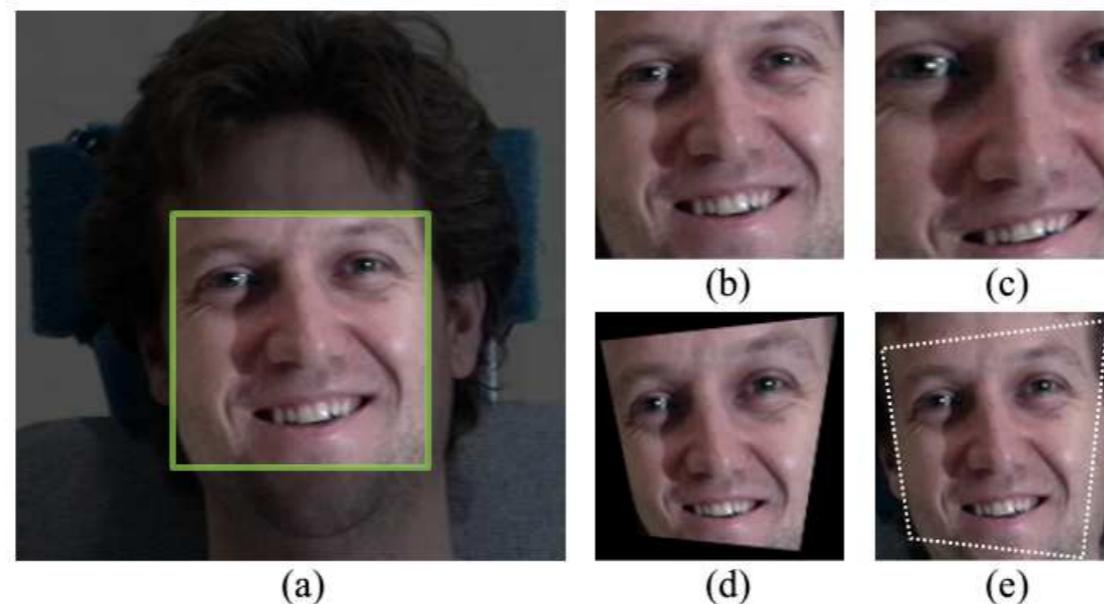
Inverse Compositional Spatial Transformer Networks

Chen-Hsuan Lin Simon Lucey

The Robotics Institute

Carnegie Mellon University

chenhsul@andrew.cmu.edu slucey@cs.cmu.edu



A look back

- Geometric warps
 - Foreshortening-vs-projective
 - Peicewise affine
 - Fwd-vs-inverse warps
- Lucas Kanade
 - Gauss Newton optimiztion of nonlinear least squares
 - Inverse composition

