

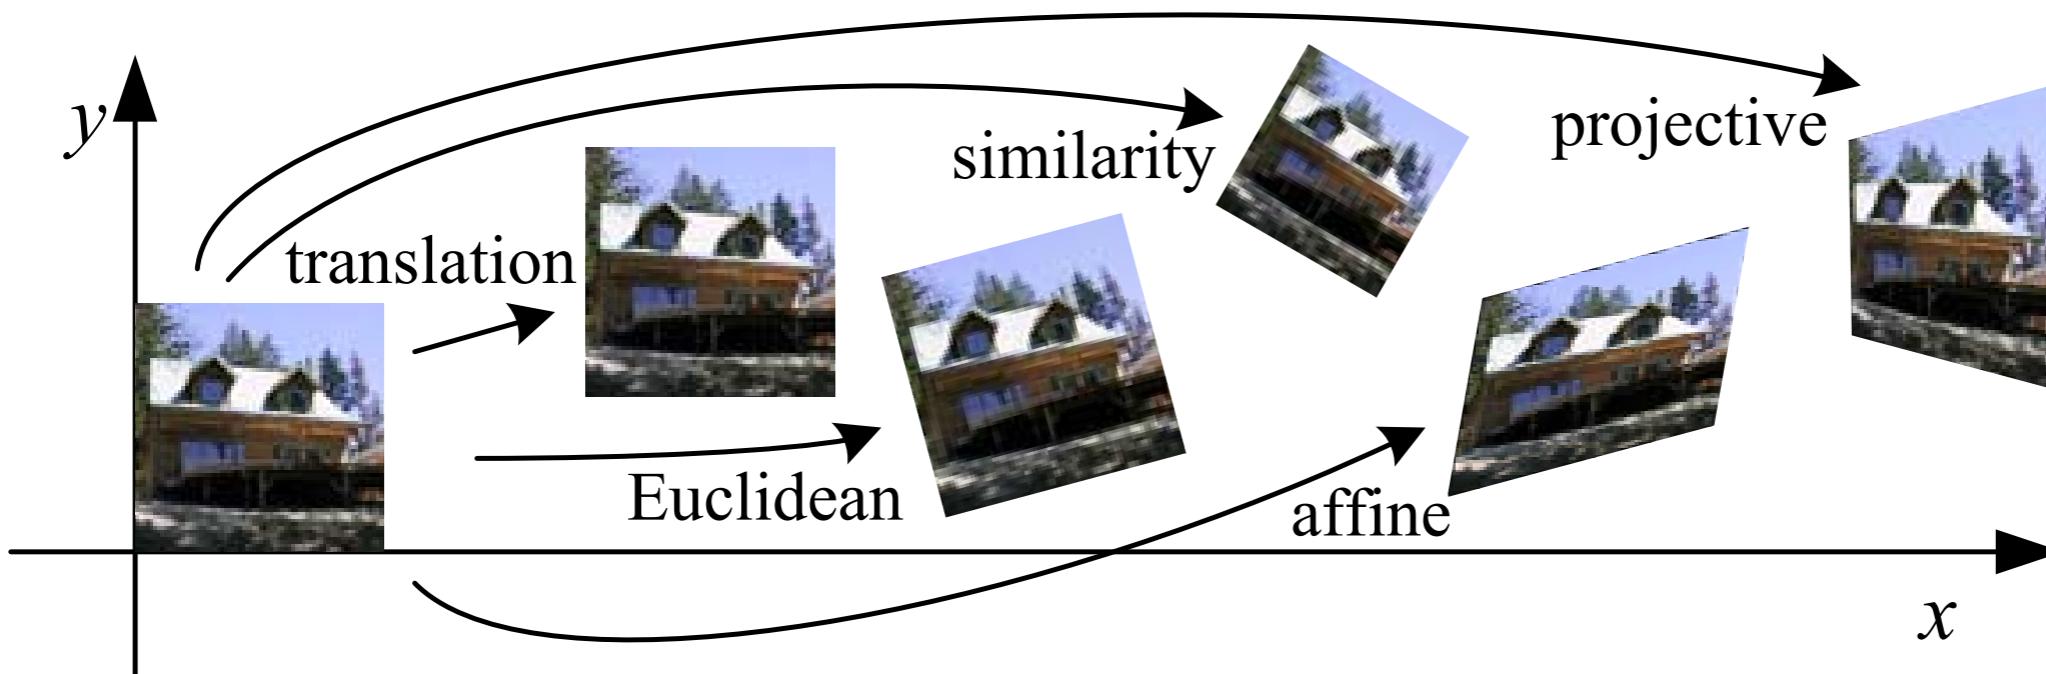
Image Alignment

(material for HW2)

Recall:

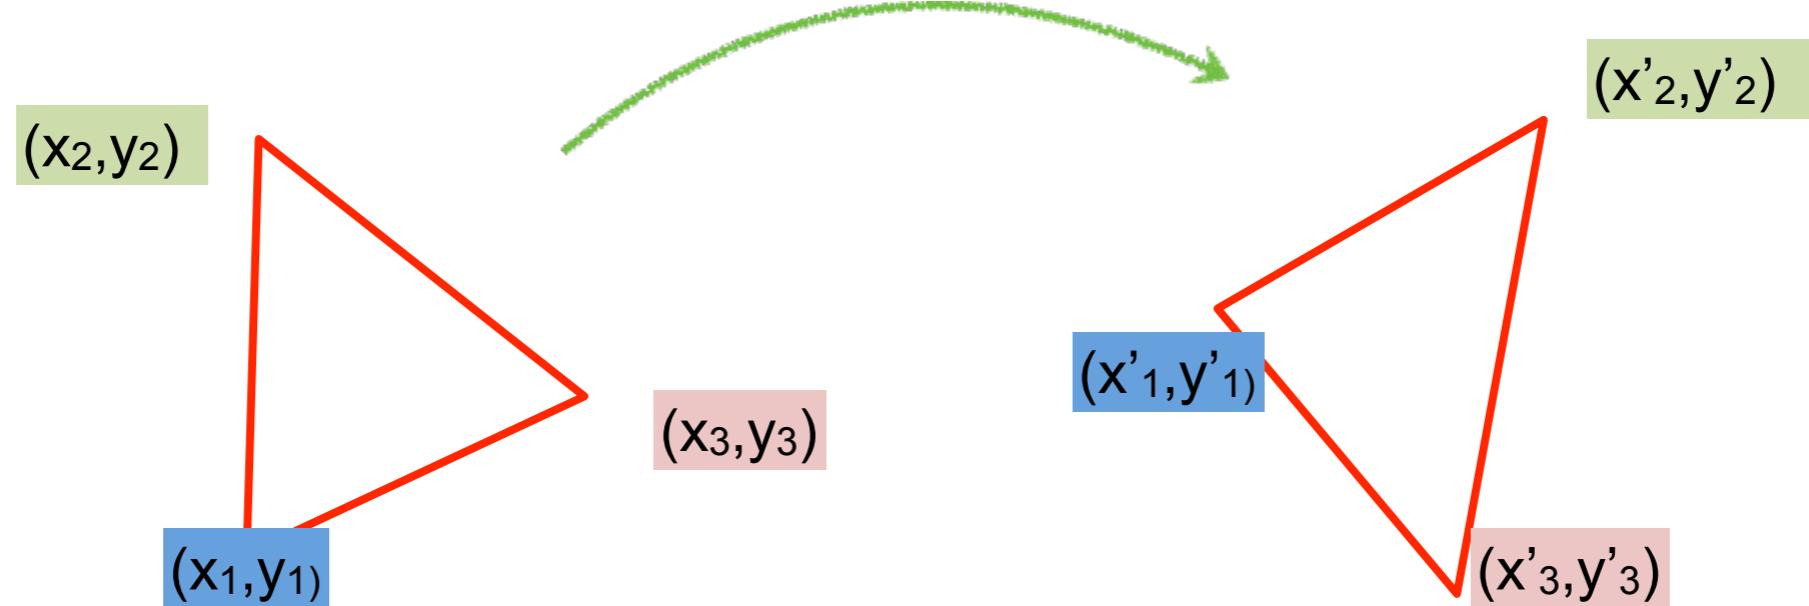
family of 2D transformations representable by a matrix

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	



Estimating transformations from point correspondences

Suppose we have two triangles:

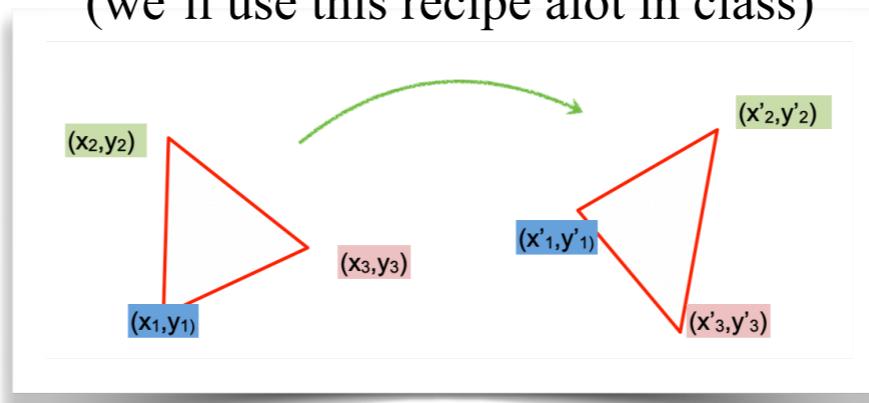


$$\begin{bmatrix} x'_1 \\ y'_1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$$

Given point correspondences, how do we estimate affine parameters (a, b, c, d, e, f)?

Estimating motion from correspondences

(we'll use this recipe a lot in class)



Affine transformation:

$$\begin{aligned} ax + by + c &= x' \\ sx + ey + f &= y' \end{aligned} \quad \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix}$$

Write for all three points (1,2,3):

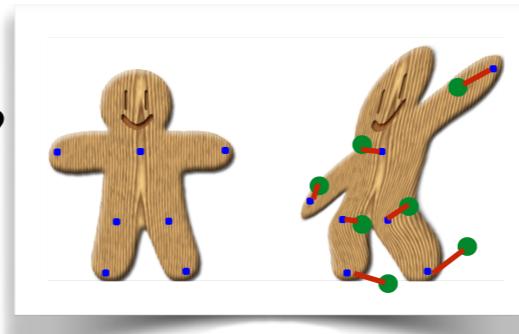
$$\begin{aligned} ax_1 + by_1 + c &= x'_1 \\ dx_1 + ey_1 + f &= y'_1 \\ ax_2 + by_2 + c &= x'_2 \\ dx_2 + ey_2 + f &= y'_2 \\ ax_3 + by_3 + c &= x'_3 \\ dx_3 + ey_3 + f &= y'_3 \end{aligned} \quad \begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_1 & y_1 & 1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_2 & y_2 & 1 \\ x_3 & y_3 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_3 & y_3 & 1 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} x'_1 \\ y'_1 \\ x'_2 \\ y'_2 \\ x'_3 \\ y'_3 \end{bmatrix}$$
$$A \quad \mathbf{x} = \mathbf{b}$$
$$\mathbf{x} = \text{numpy.linalg.solve}(A, \mathbf{b})$$

$$\mathbf{x} = A^{-1}\mathbf{b}$$

When is A invertible?

If 3 points are not *co-linear*

What if we have more than 3 points?



Estimate parameters that minimize (squared) euclidean error with least squares

$$\mathbf{x} = \underset{\mathbf{X}}{\operatorname{argmin}} \|\hat{\mathbf{b}} - \mathbf{b}\|^2 \text{ where } \hat{\mathbf{b}} = A\mathbf{x}$$

$$\mathbf{x} = \text{numpy.linalg.lstsq}(A, \mathbf{b})$$

Image warping via piecewise affine warps



Piecewise affine warps (cut each quadrilateral into 2 triangles)

5

Compute intermediate vertex position by linear interpolation $\alpha \begin{bmatrix} x \\ y \end{bmatrix} + (1 - \alpha) \begin{bmatrix} x' \\ y' \end{bmatrix}$

Compute intermediate color by linear interpolation $\alpha \begin{bmatrix} R \\ G \\ B \end{bmatrix} + (1 - \alpha) \begin{bmatrix} R' \\ G' \\ B' \end{bmatrix}$

Example application: shape modeling



D'Arcy Thompson
“On Growth and Form” 1915

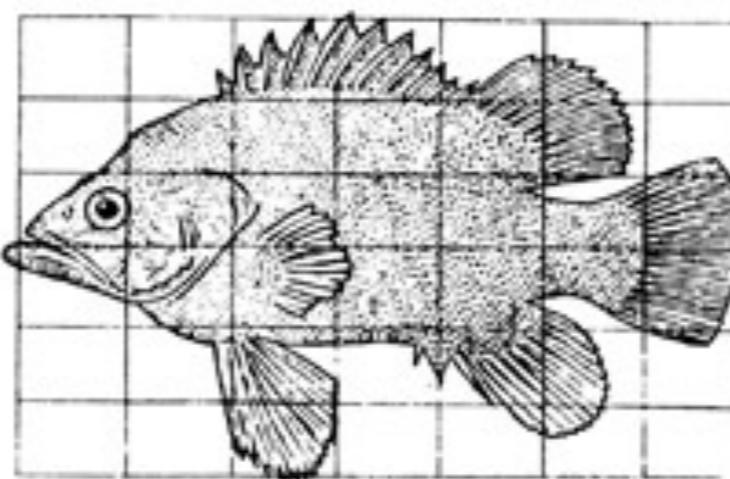


Fig. 150. *Polyprion*.

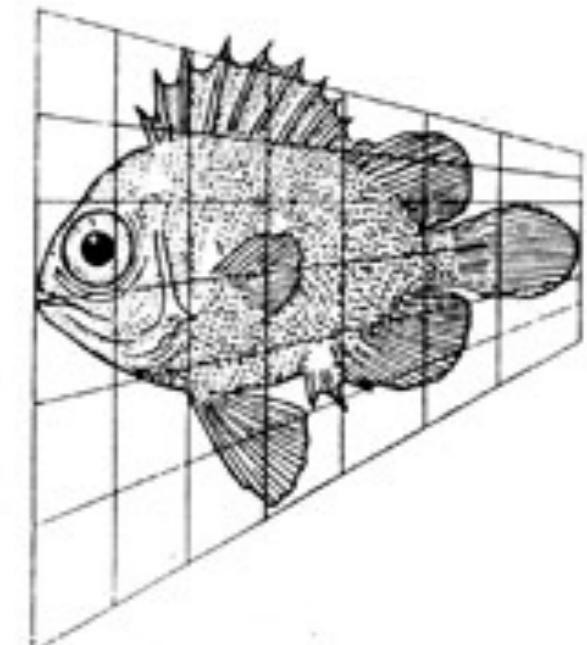


Fig. 151. *Pseudopriacanthus altus*.

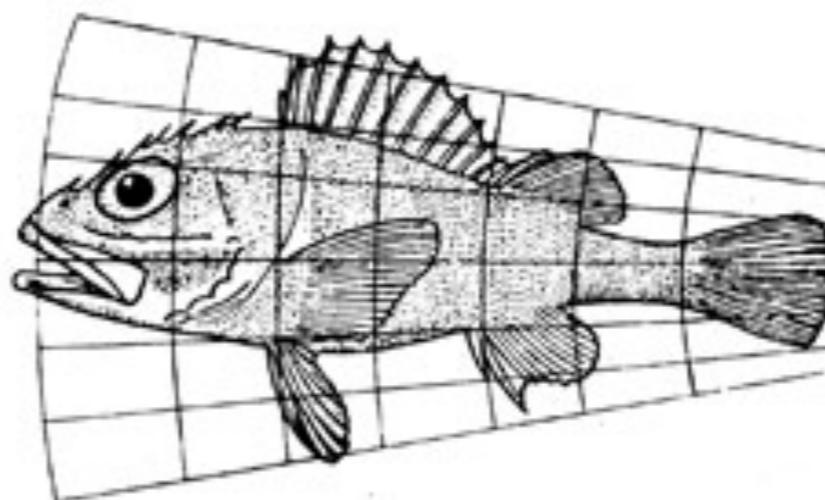


Fig. 152. *Scorpaena* sp.

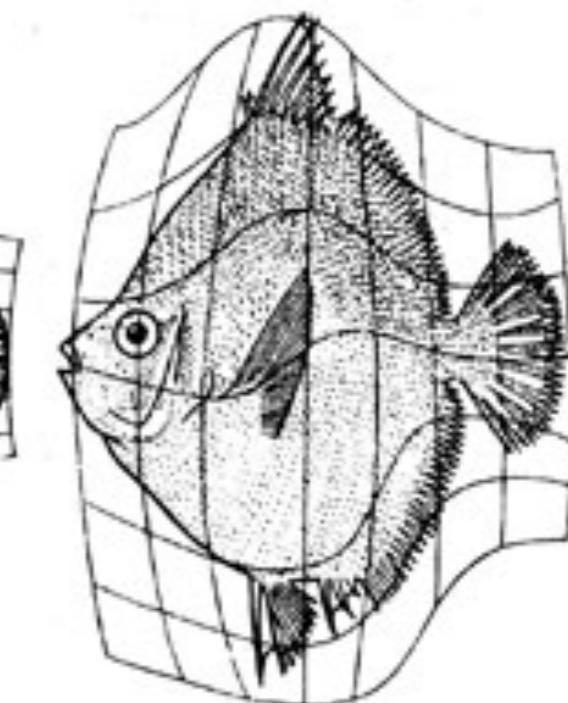
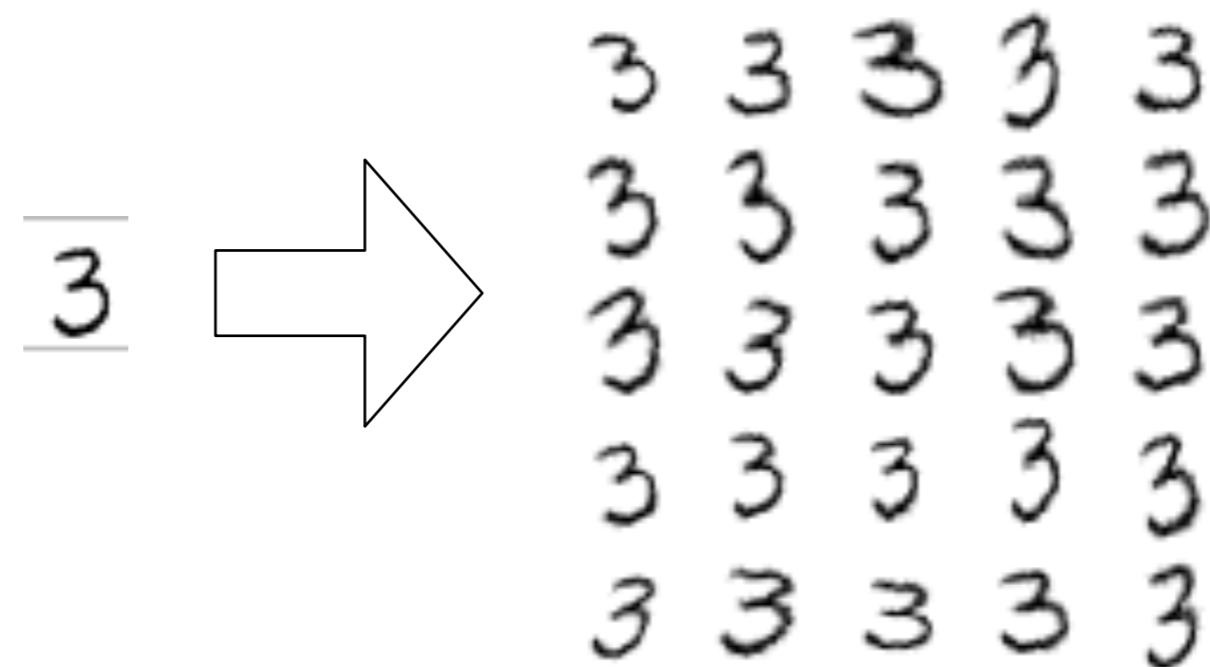


Fig. 153. *Antigonia capros*.

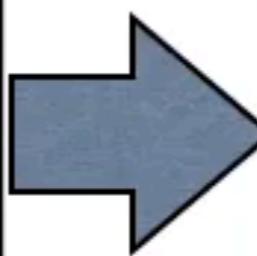
Application: data augmentation



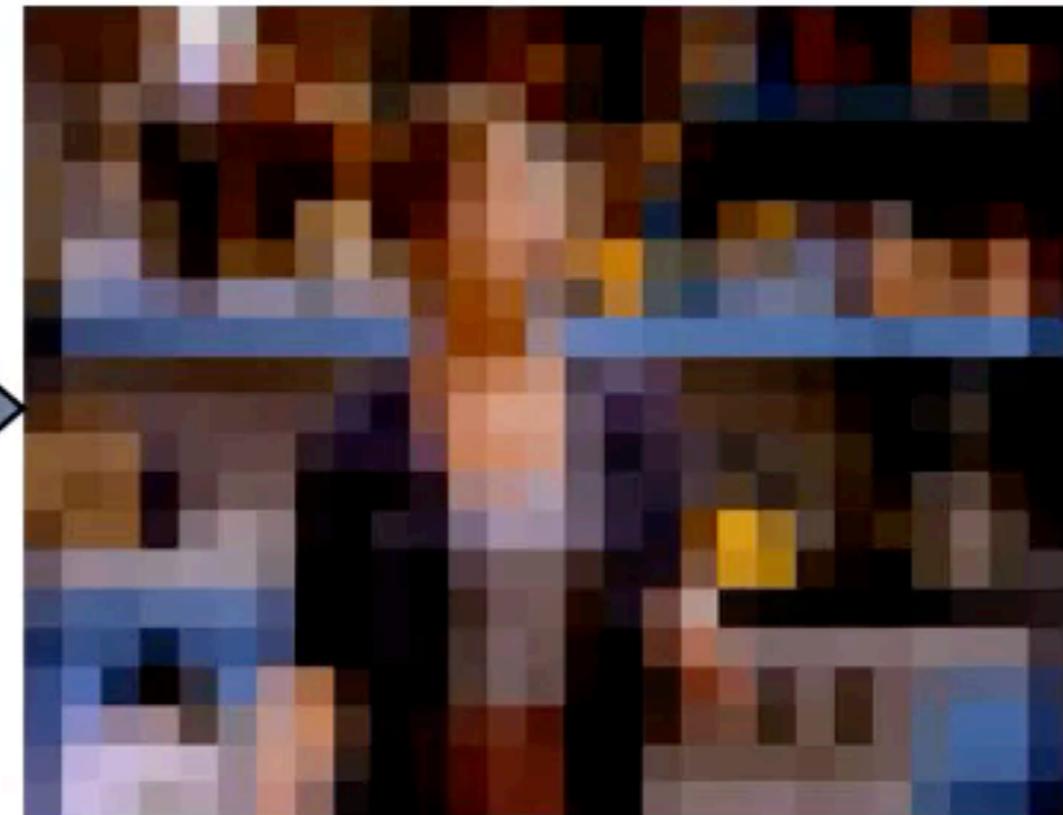
Hypothesis: we'll see even more of this in the future!

test video

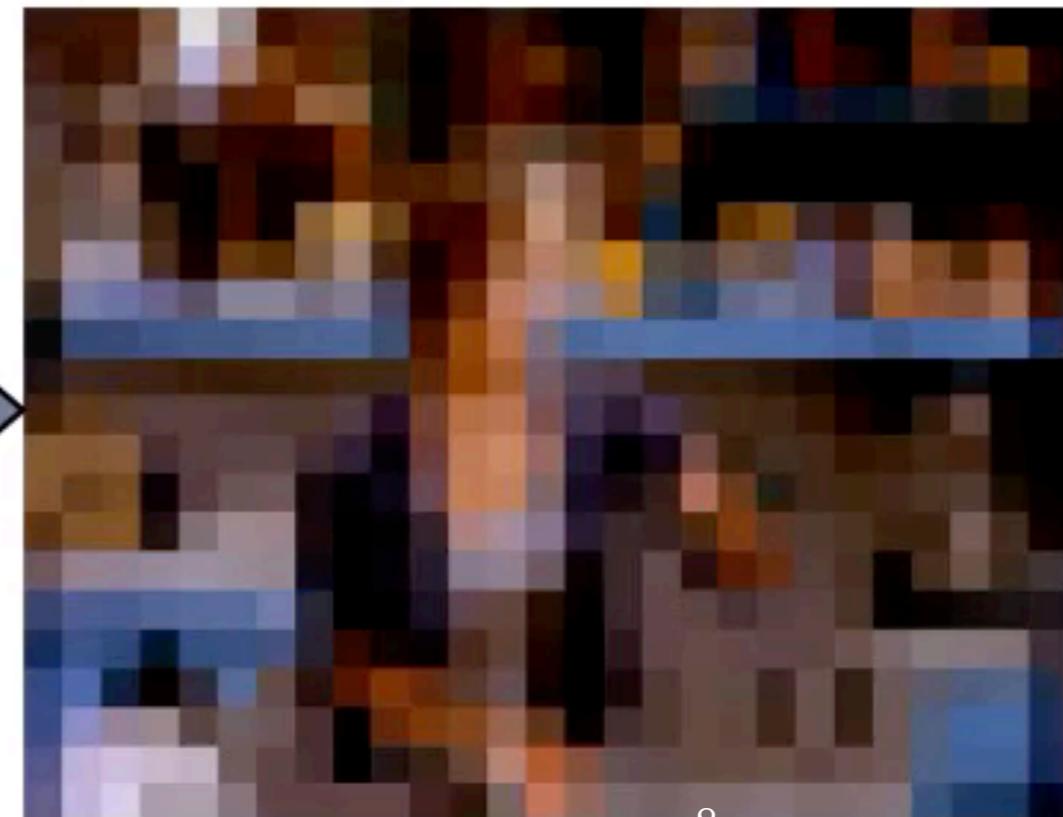
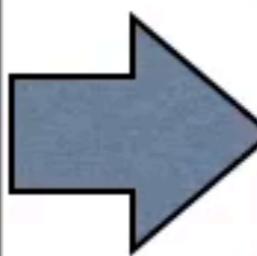
original



low resolution



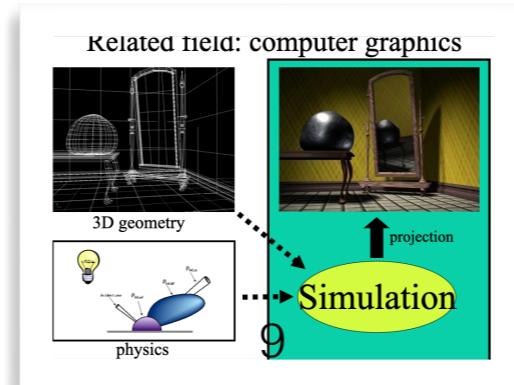
best match

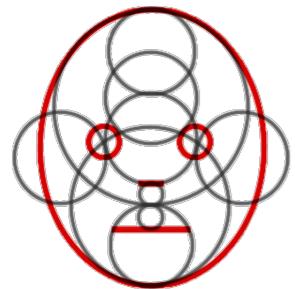
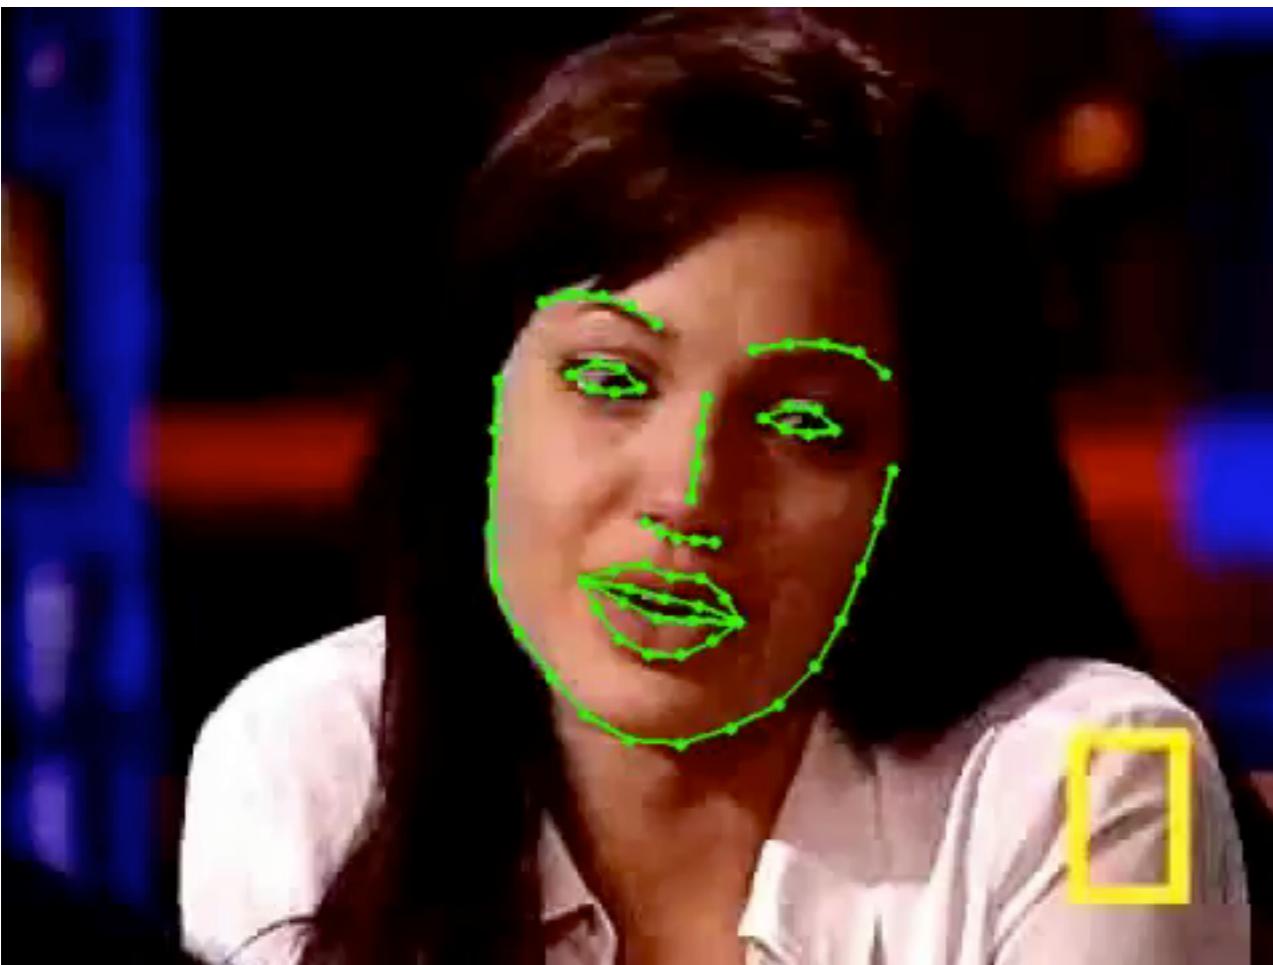


Application: Tracking (HW2)!



Really instance of vision as inverse graphics or “analysis by synthesis”; interpret the world by *optimizing* for best fit of model



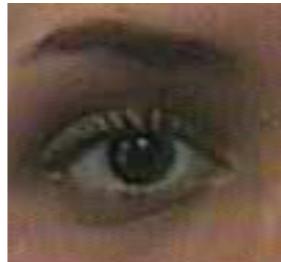


IntraFace





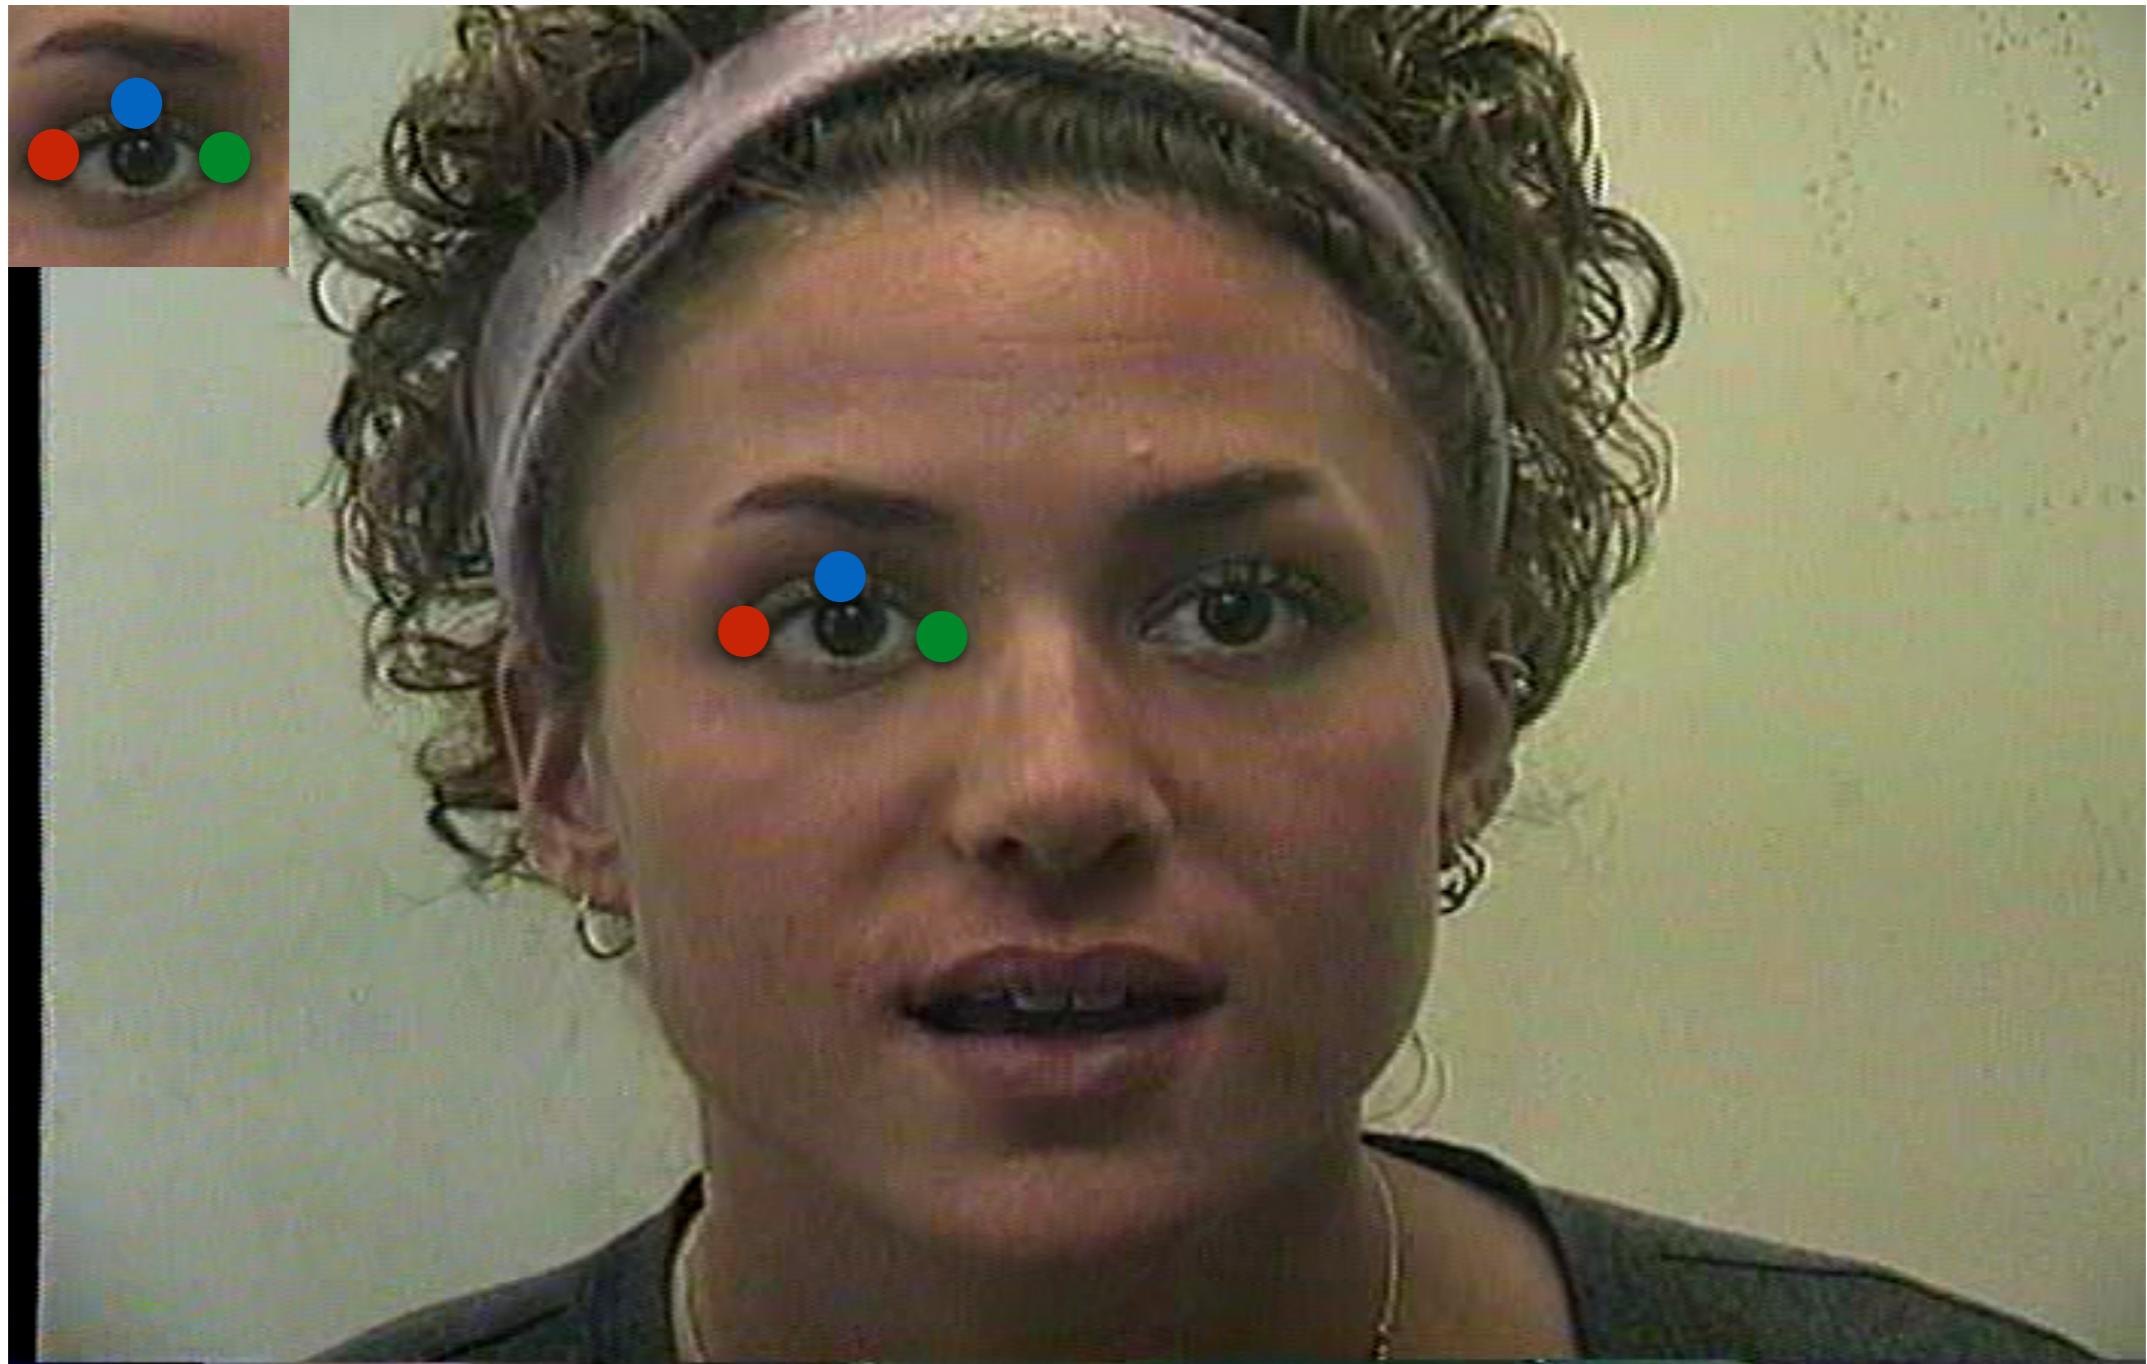
How can I find



in the image?

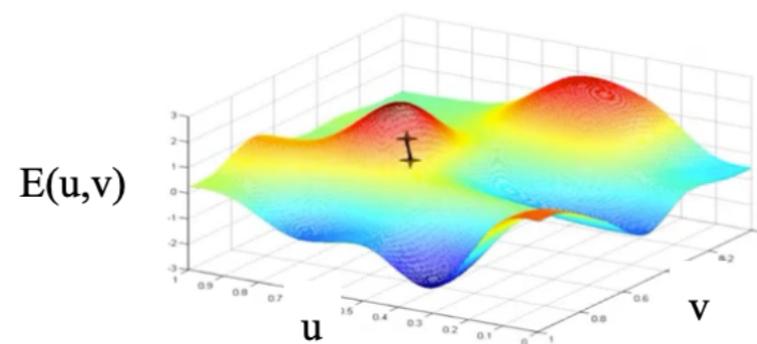
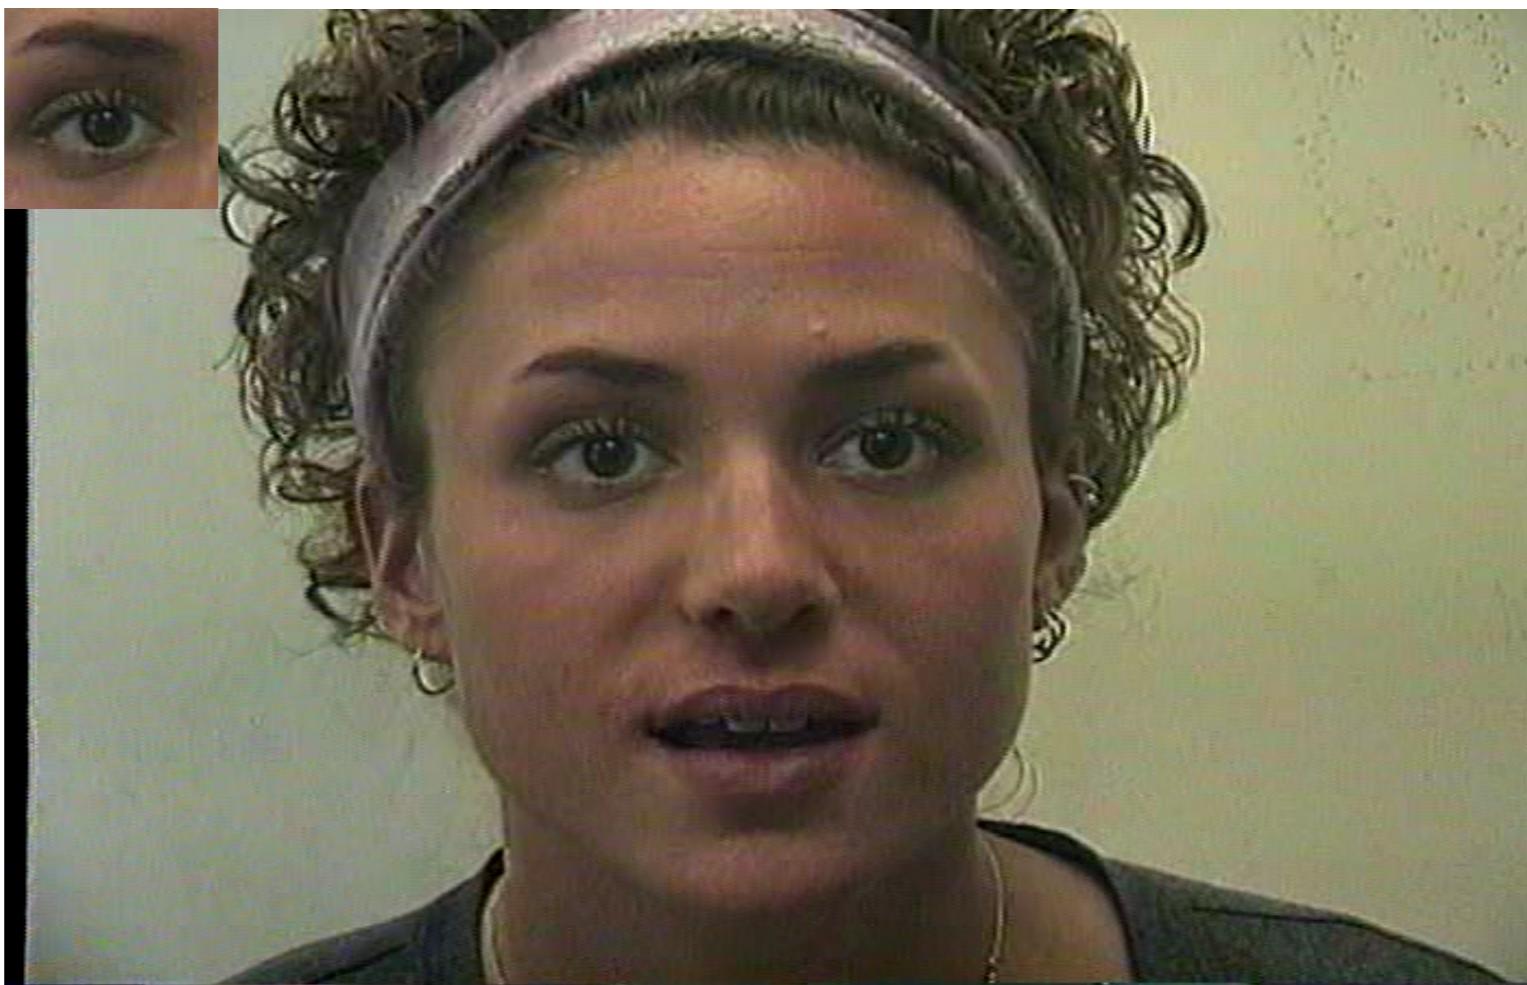


Idea #1: Least squares fit to corresponding points



Works great, but need to find corresponding points :(
(will explore automatic methods in HW3!)

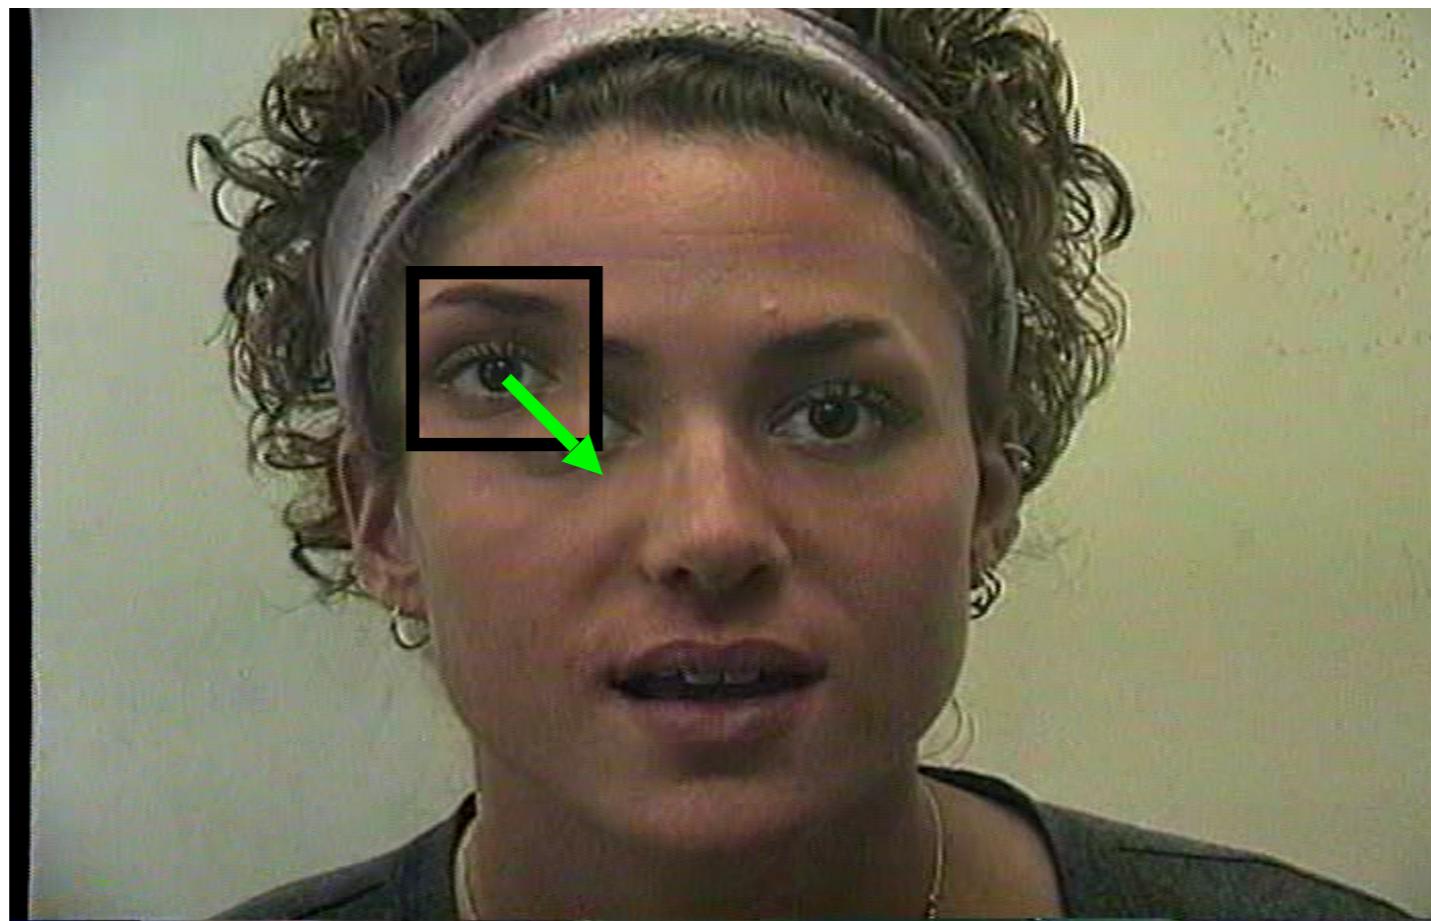
Idea #2: Template Matching by Filtering



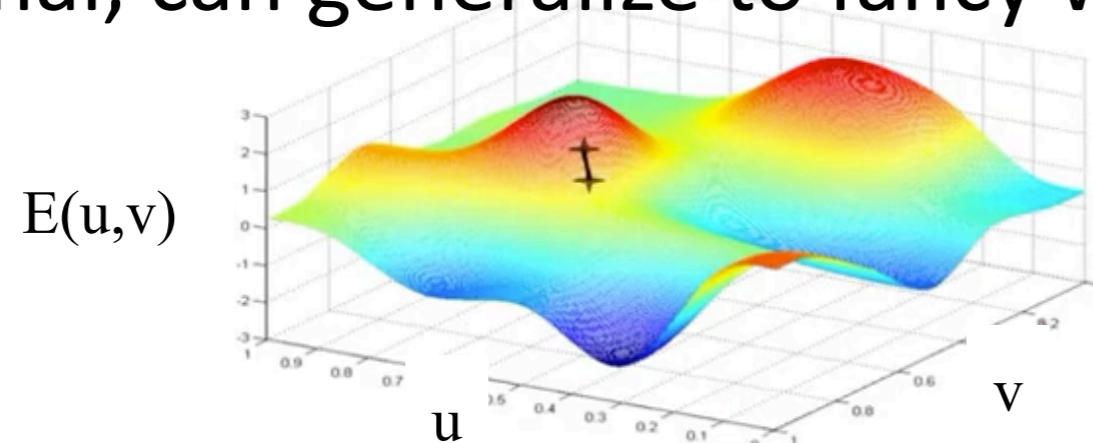
Global solution but slow

(hard to generalize beyond translation to affine warps)

Idea #3: Local search with gradient descent



Locally optimal, can generalize to fancy warps, but need a good init

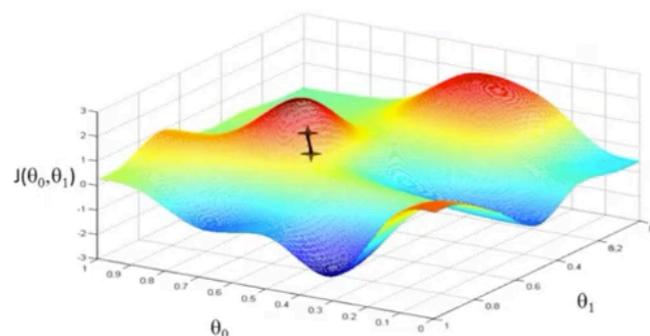


Where we are headed (HW2)



Tracking via template matching:
initialize bounding box in frame 1 and warp to match it to frame 2 (and repeat)

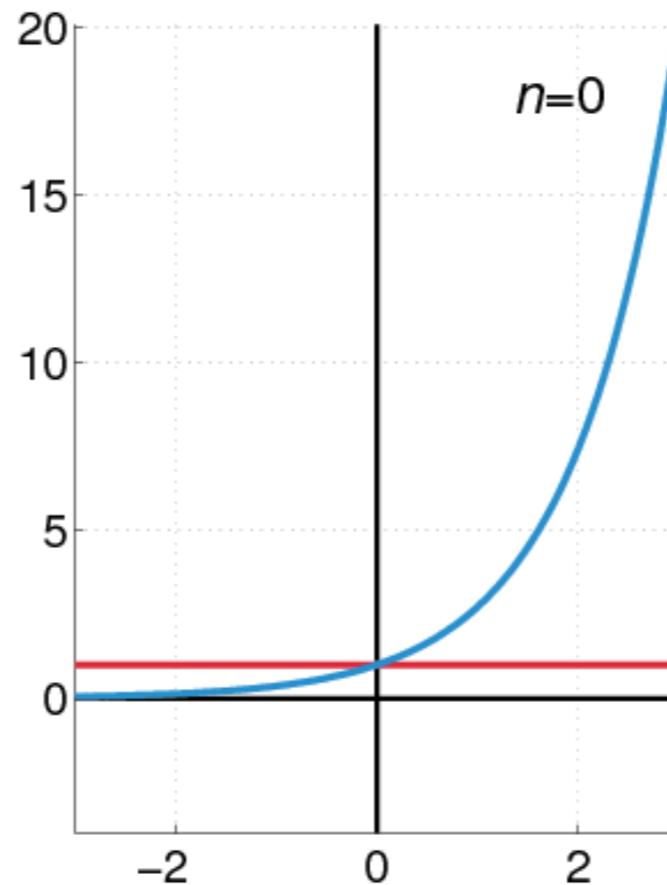
Turns out, we can do much better than gradient descent by opening up our “optimization toolbox”



Background: nonlinear optimization

Let's approximate a complicated function (or error surface) with a Taylor series

$$f(u + \Delta u) = f(u) + \frac{\partial f(u)}{\partial u} \Delta u + \frac{1}{2} \frac{\partial^2 f(u)}{\partial u \partial u} \Delta u^2 + \text{Higher Order Terms}$$



How do we generalize the above to a **multivariate** function $f(u,v)$?

Generalize first derivative to $\begin{bmatrix} \frac{\partial f}{\partial u} \\ \frac{\partial f}{\partial v} \end{bmatrix}$ gradient and second derivative to $\begin{bmatrix} \frac{\partial^2 f}{\partial u \partial u} & \frac{\partial^2 f}{\partial u \partial v} \\ \frac{\partial^2 f}{\partial v \partial u} & \frac{\partial^2 f}{\partial v \partial v} \end{bmatrix}$ Hessian

symmetric matrix!
Is it PSD?
only for convex
functions

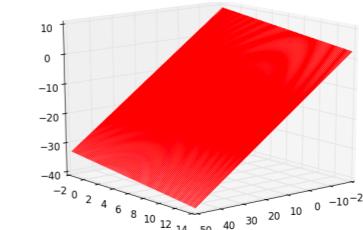
Background: nonlinear optimization

$$\min_{\mathbf{u}} f(\mathbf{u}) \quad \text{where } \mathbf{u} = (u, v)$$

1. First order method (gradient descent)

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \alpha \begin{bmatrix} \frac{\partial f}{\partial u} \\ \frac{\partial f}{\partial v} \end{bmatrix}$$

“fit a plane and take a step”

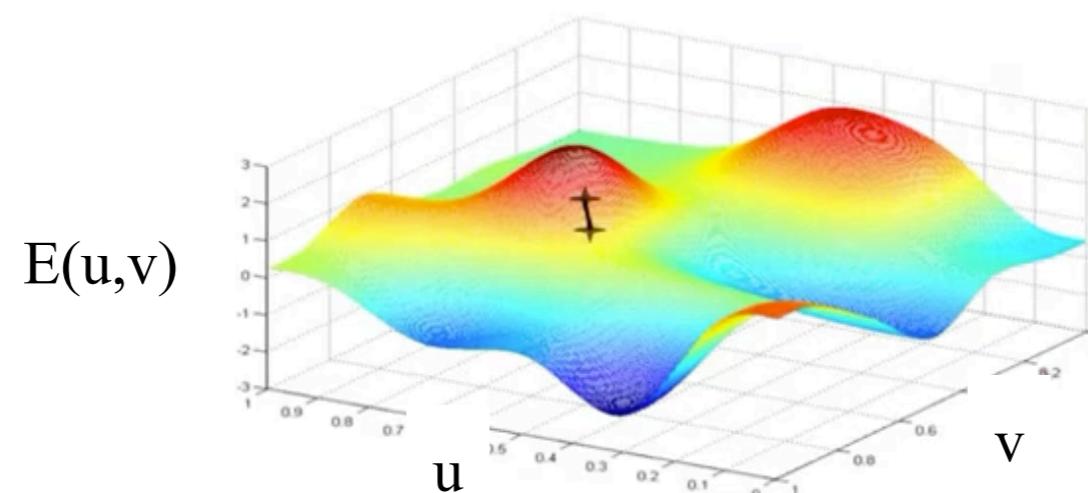
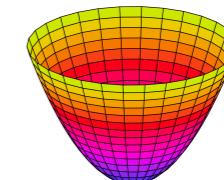


2. Second order method (Newton's method)

https://en.wikipedia.org/wiki/Newton's_method_in_optimization

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \begin{bmatrix} \frac{\partial^2 f}{\partial u \partial u} & \frac{\partial^2 f}{\partial u \partial v} \\ \frac{\partial^2 f}{\partial v \partial u} & \frac{\partial^2 f}{\partial v \partial v} \end{bmatrix}^{-1} \begin{bmatrix} \frac{\partial f}{\partial u} \\ \frac{\partial f}{\partial v} \end{bmatrix}$$

“fit a paraboloid and jump to minimum”



Some notation before we get into the math...

2D image transformation

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix}$$

2D image coordinate

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Parameters of the transformation

$$\mathbf{p} = \begin{bmatrix} p_1 \\ p_2 \\ \vdots \end{bmatrix}$$

Warped image

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(W_x(\mathbf{x}; \mathbf{p}), W_y(\mathbf{x}; \mathbf{p}))$$

Translation

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} x + p_1 \\ y + p_2 \end{bmatrix}$$

$$= \begin{bmatrix} 1 & 0 & p_1 \\ 0 & 1 & p_2 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(x + p_1, y + p_2)$$

Affine

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$= \begin{bmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

Find the warp parameters \mathbf{p} such that minimize SSD

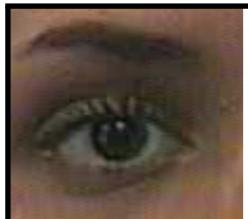
$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

$$\min_{(p_1, p_2, \dots)} \sum_{x,y} [I(p_1x + p_2y + p_3, p_4x + p_5y + p_6) - T(x, y)]^2$$

what pixels should be summed over?

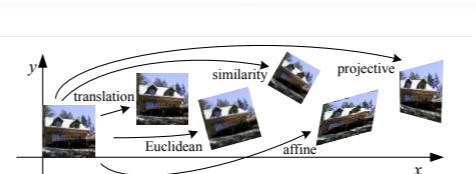
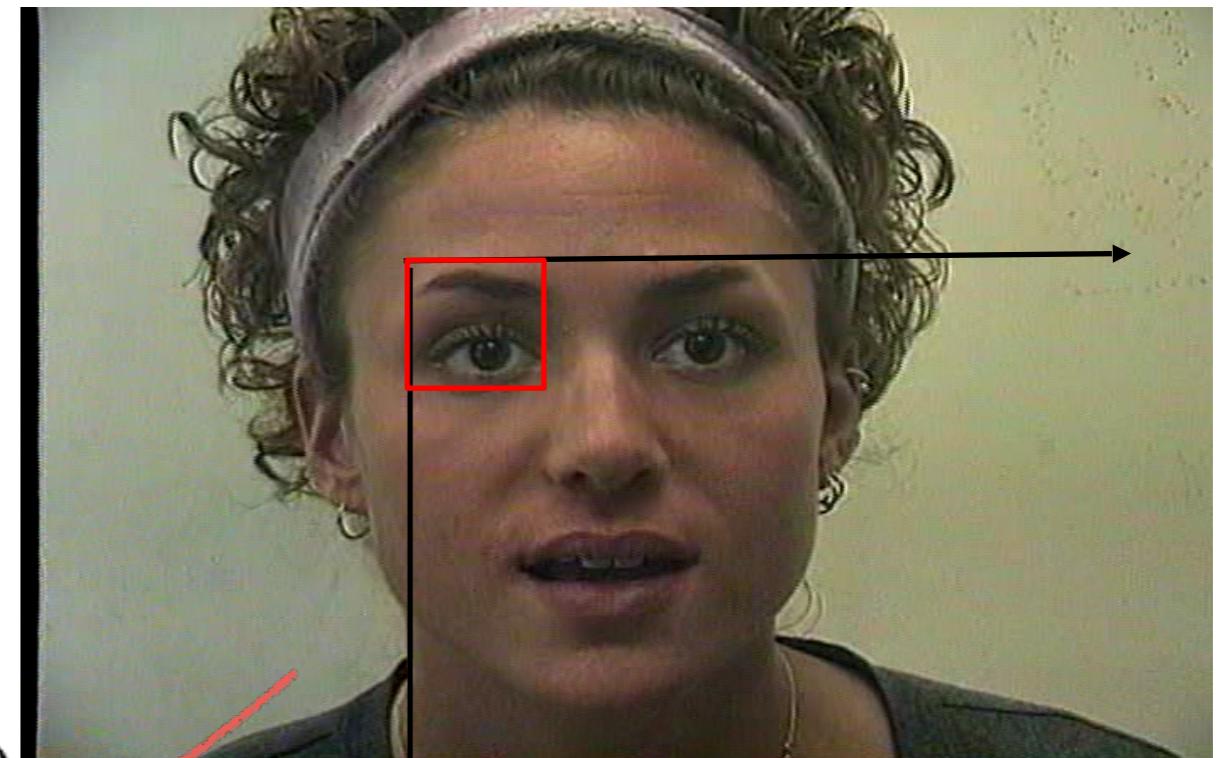
(only those that are non-zero
in both template and warped image)

$T(\mathbf{x})$



$\mathbf{W}(\mathbf{x}; \mathbf{p})$

$I(\mathbf{x})$



This is a non-linear least squares optimization!

What's the source of the nonlinearity? the image

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

$$\min_{(p_1, p_2, \dots)} \sum_{x,y} [I(p_1x + p_2y + p_3, p_4x + p_5y + p_6) - T(x, y)]^2$$

Assuming a good initialization of \mathbf{p} ,
linearize objective and update incrementally

Lucas-Kanade alignment

Lucas-Kanade 20 Years On: A Unifying Framework

SIMON BAKER AND IAIN MATTHEWS

The Robotics Institute, Carnegie Mellon University, 5000 Forbes Avenue, Pittsburgh, PA 15213, USA

simonb@cs.cmu.edu

iainm@cs.cmu.edu

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

$$\min_{(p_1, p_2, \dots)} \sum_{x,y} [I(p_1x + p_2y + p_3, p_4x + p_5y + p_6) - T(x, y)]^2$$

Assuming an initial guess \mathbf{p} , let's optimize over increment $\Delta \mathbf{p}$

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) - T(\mathbf{x})]^2$$

This is **still** a nonlinear least squares optimization wrt $\Delta \mathbf{p}$

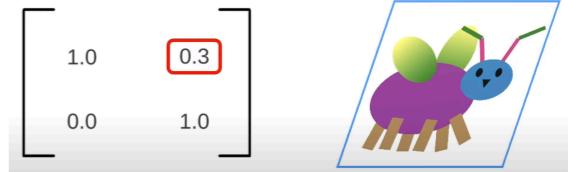
Key assumption: let's assume initial guess \mathbf{p} is good and increment $\Delta \mathbf{p}$ will be small
Then we can remove nonlinearity with a Taylor expansion

$$f(x, y) \approx f(a, b) + f_x(a, b)(x - a) - f_y(a, b)(y - b)$$

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

how does image change as I
tweak the warp parameters?



$$I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) \approx I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \frac{\partial I(\mathbf{W}(\mathbf{x}; \mathbf{p}))}{\partial \mathbf{p}} \Delta\mathbf{p}$$

chain rule

$$= I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \frac{\partial I(\mathbf{W}(\mathbf{x}; \mathbf{p}))}{\partial \mathbf{W}(\mathbf{x}; \mathbf{p})} \frac{\partial \mathbf{W}(\mathbf{x}; \mathbf{p})}{\partial \mathbf{p}} \Delta\mathbf{p}$$

$$= I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta\mathbf{p}$$

↑
↑
short-hand

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) - T(\mathbf{x})]^2$$

Multivariable Taylor Series Expansion

$$f(x, y) \approx f(a, b) + f_x(a, b)(x - a) + f_y(a, b)(y - b)$$

|

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

What are the unknowns here?

Now, the squared error is linear in the unknowns (suggests we can use least squares to solve for them!)

The Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$

(A matrix of partial derivatives)

Rate of change of the warp

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{bmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \dots & \frac{\partial W_x}{\partial p_N} \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \dots & \frac{\partial W_y}{\partial p_N} \end{bmatrix}$$

$$\frac{\partial W_x}{\partial p_1} = \quad \frac{\partial W_x}{\partial p_2} = \quad \dots$$

$$\frac{\partial W_y}{\partial p_1} = \quad \dots$$

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} =$$

The Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$

(A matrix of partial derivatives)

Rate of change of the warp

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{bmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \dots & \frac{\partial W_x}{\partial p_N} \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \dots & \frac{\partial W_y}{\partial p_N} \end{bmatrix}$$

$\frac{\partial W_x}{\partial p_1} = x$	$\frac{\partial W_x}{\partial p_2} = 0$	\dots
$\frac{\partial W_y}{\partial p_1} = 0$	\dots	
$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} =$	$\begin{bmatrix} x & 0 & y & 0 & 1 & 0 \\ 0 & x & 0 & y & 0 & 1 \end{bmatrix}$	

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

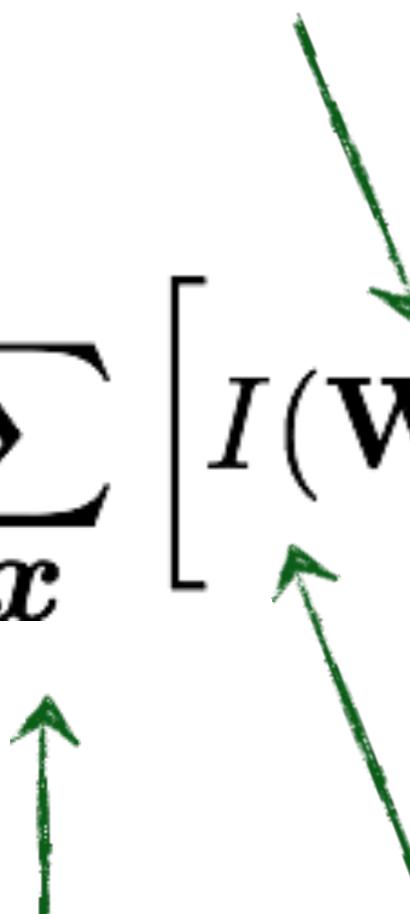


 pixel coordinate
 (2×1)

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$



 pixel coordinate (2 x 1) image intensity (scalar)

$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

warp function
(2 x 1)

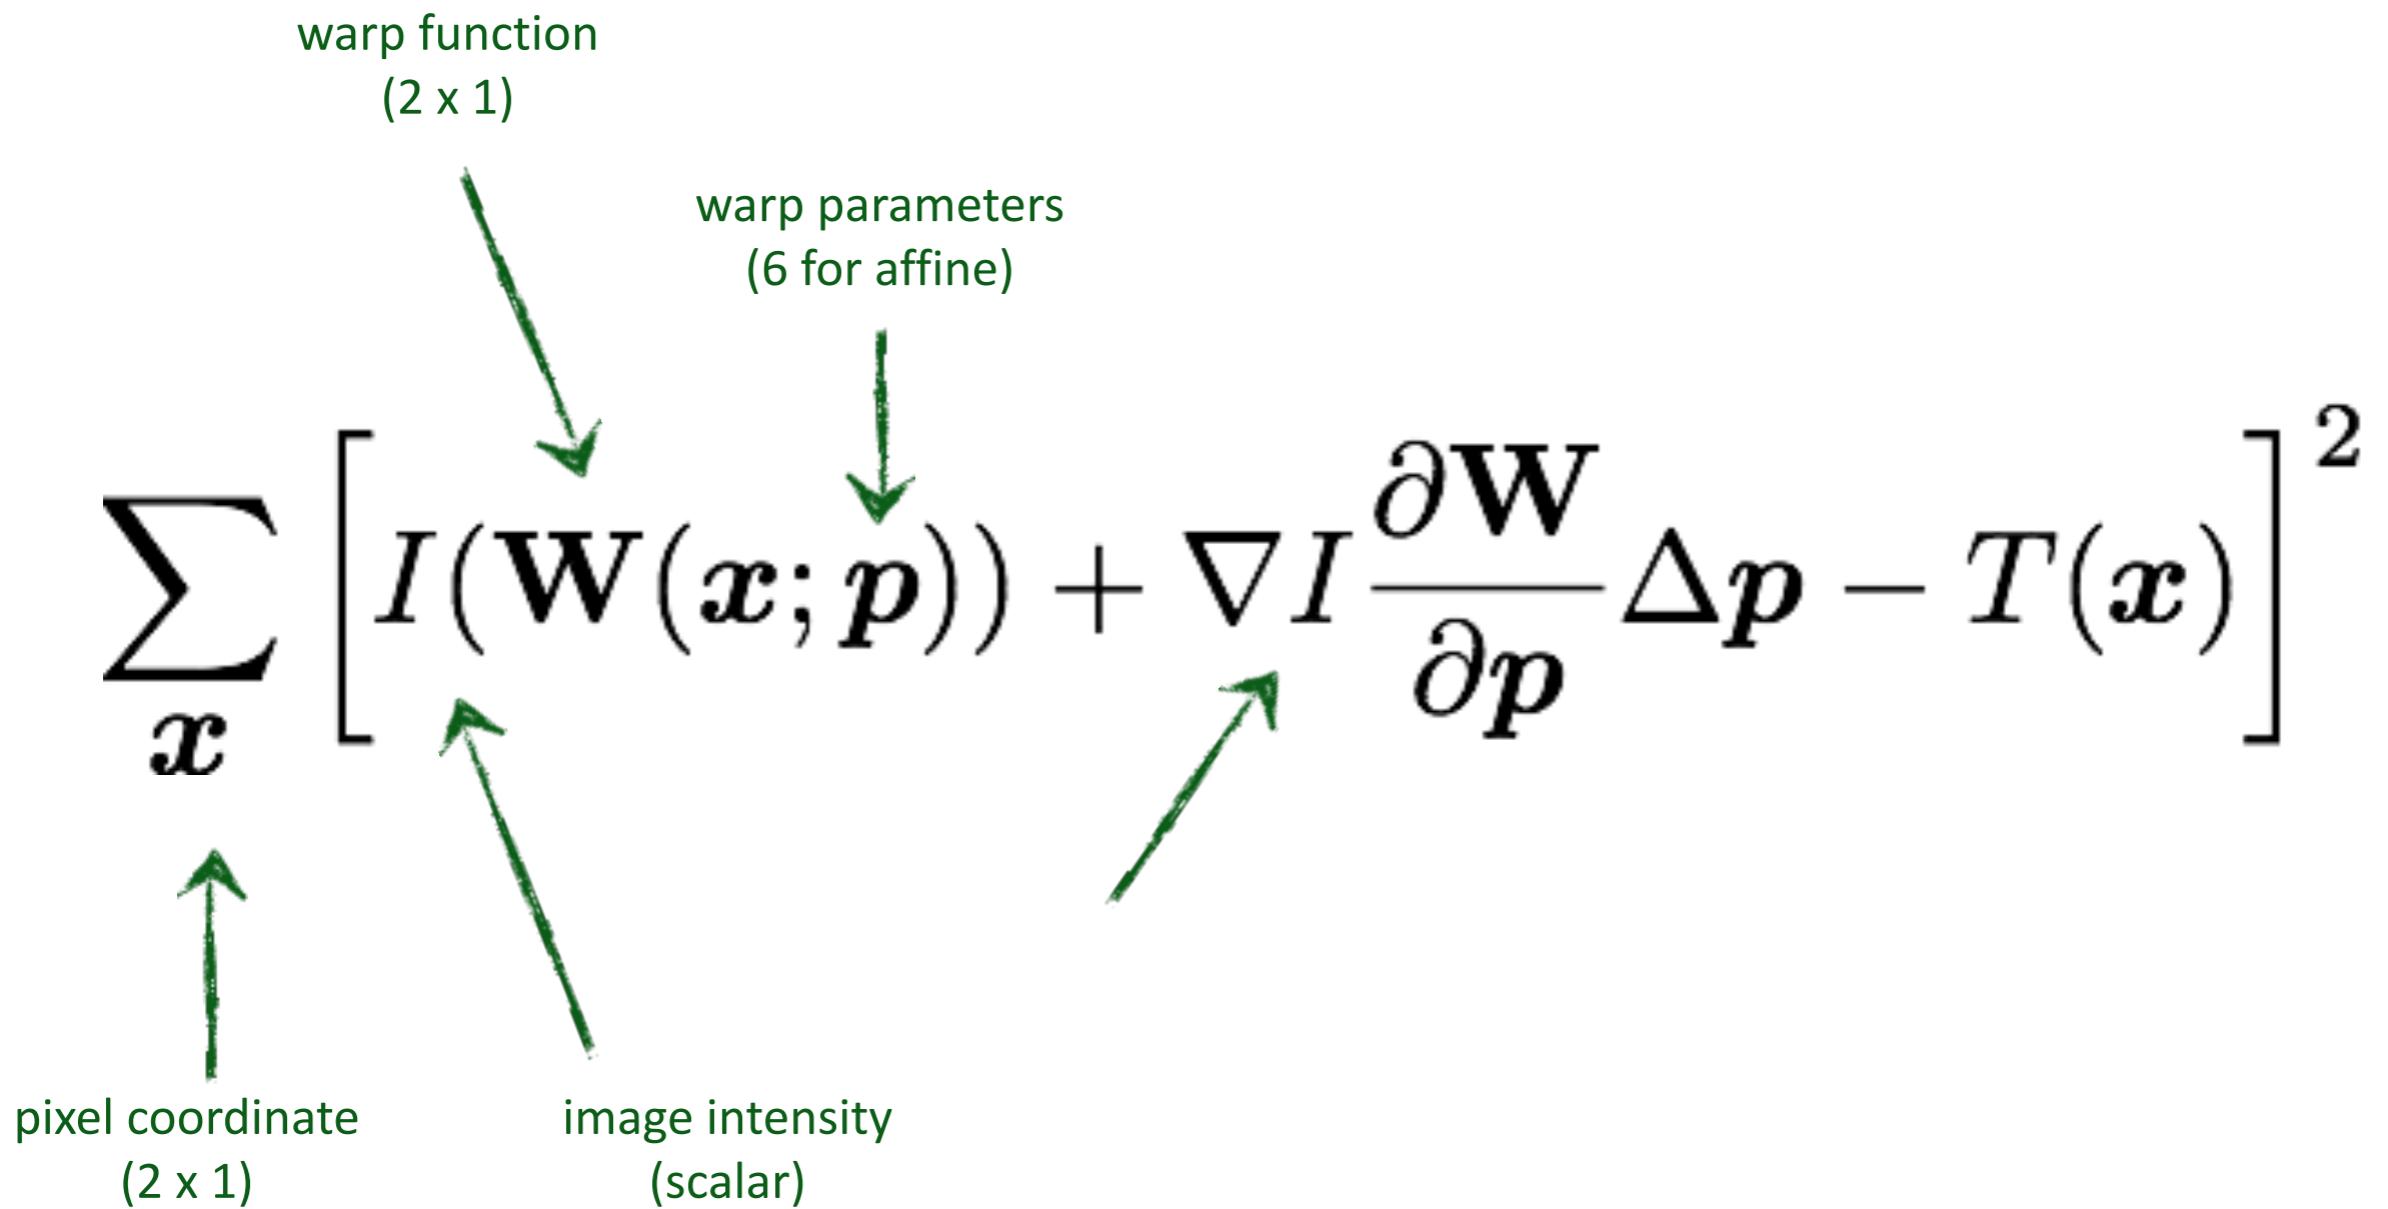
$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$

pixel coordinate
(2 x 1)

image intensity
(scalar)

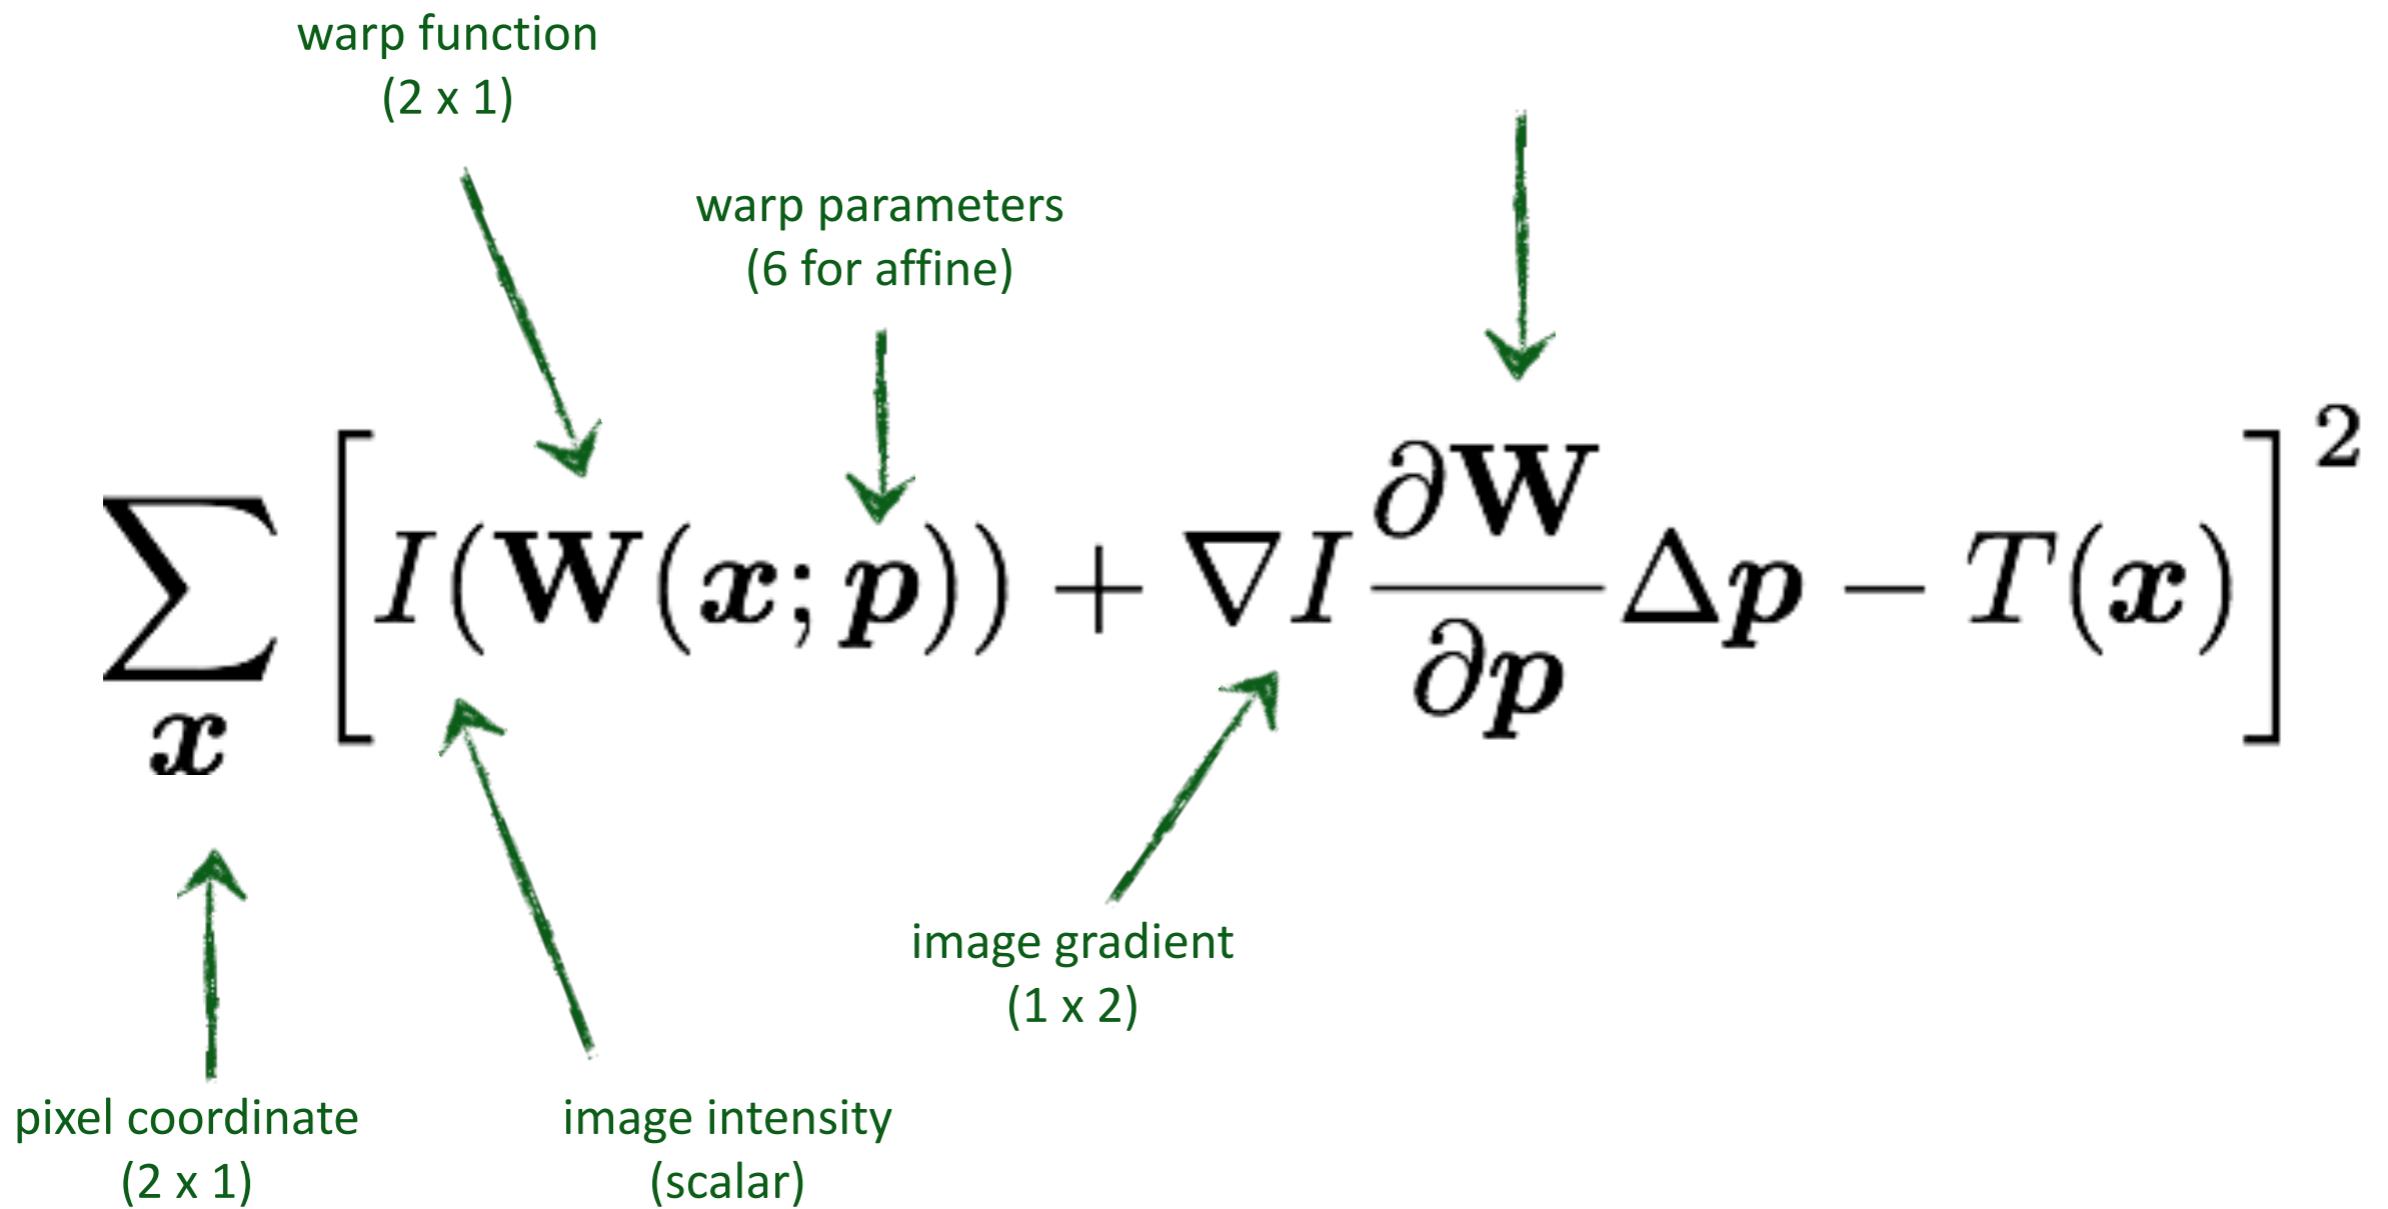
$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$



$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

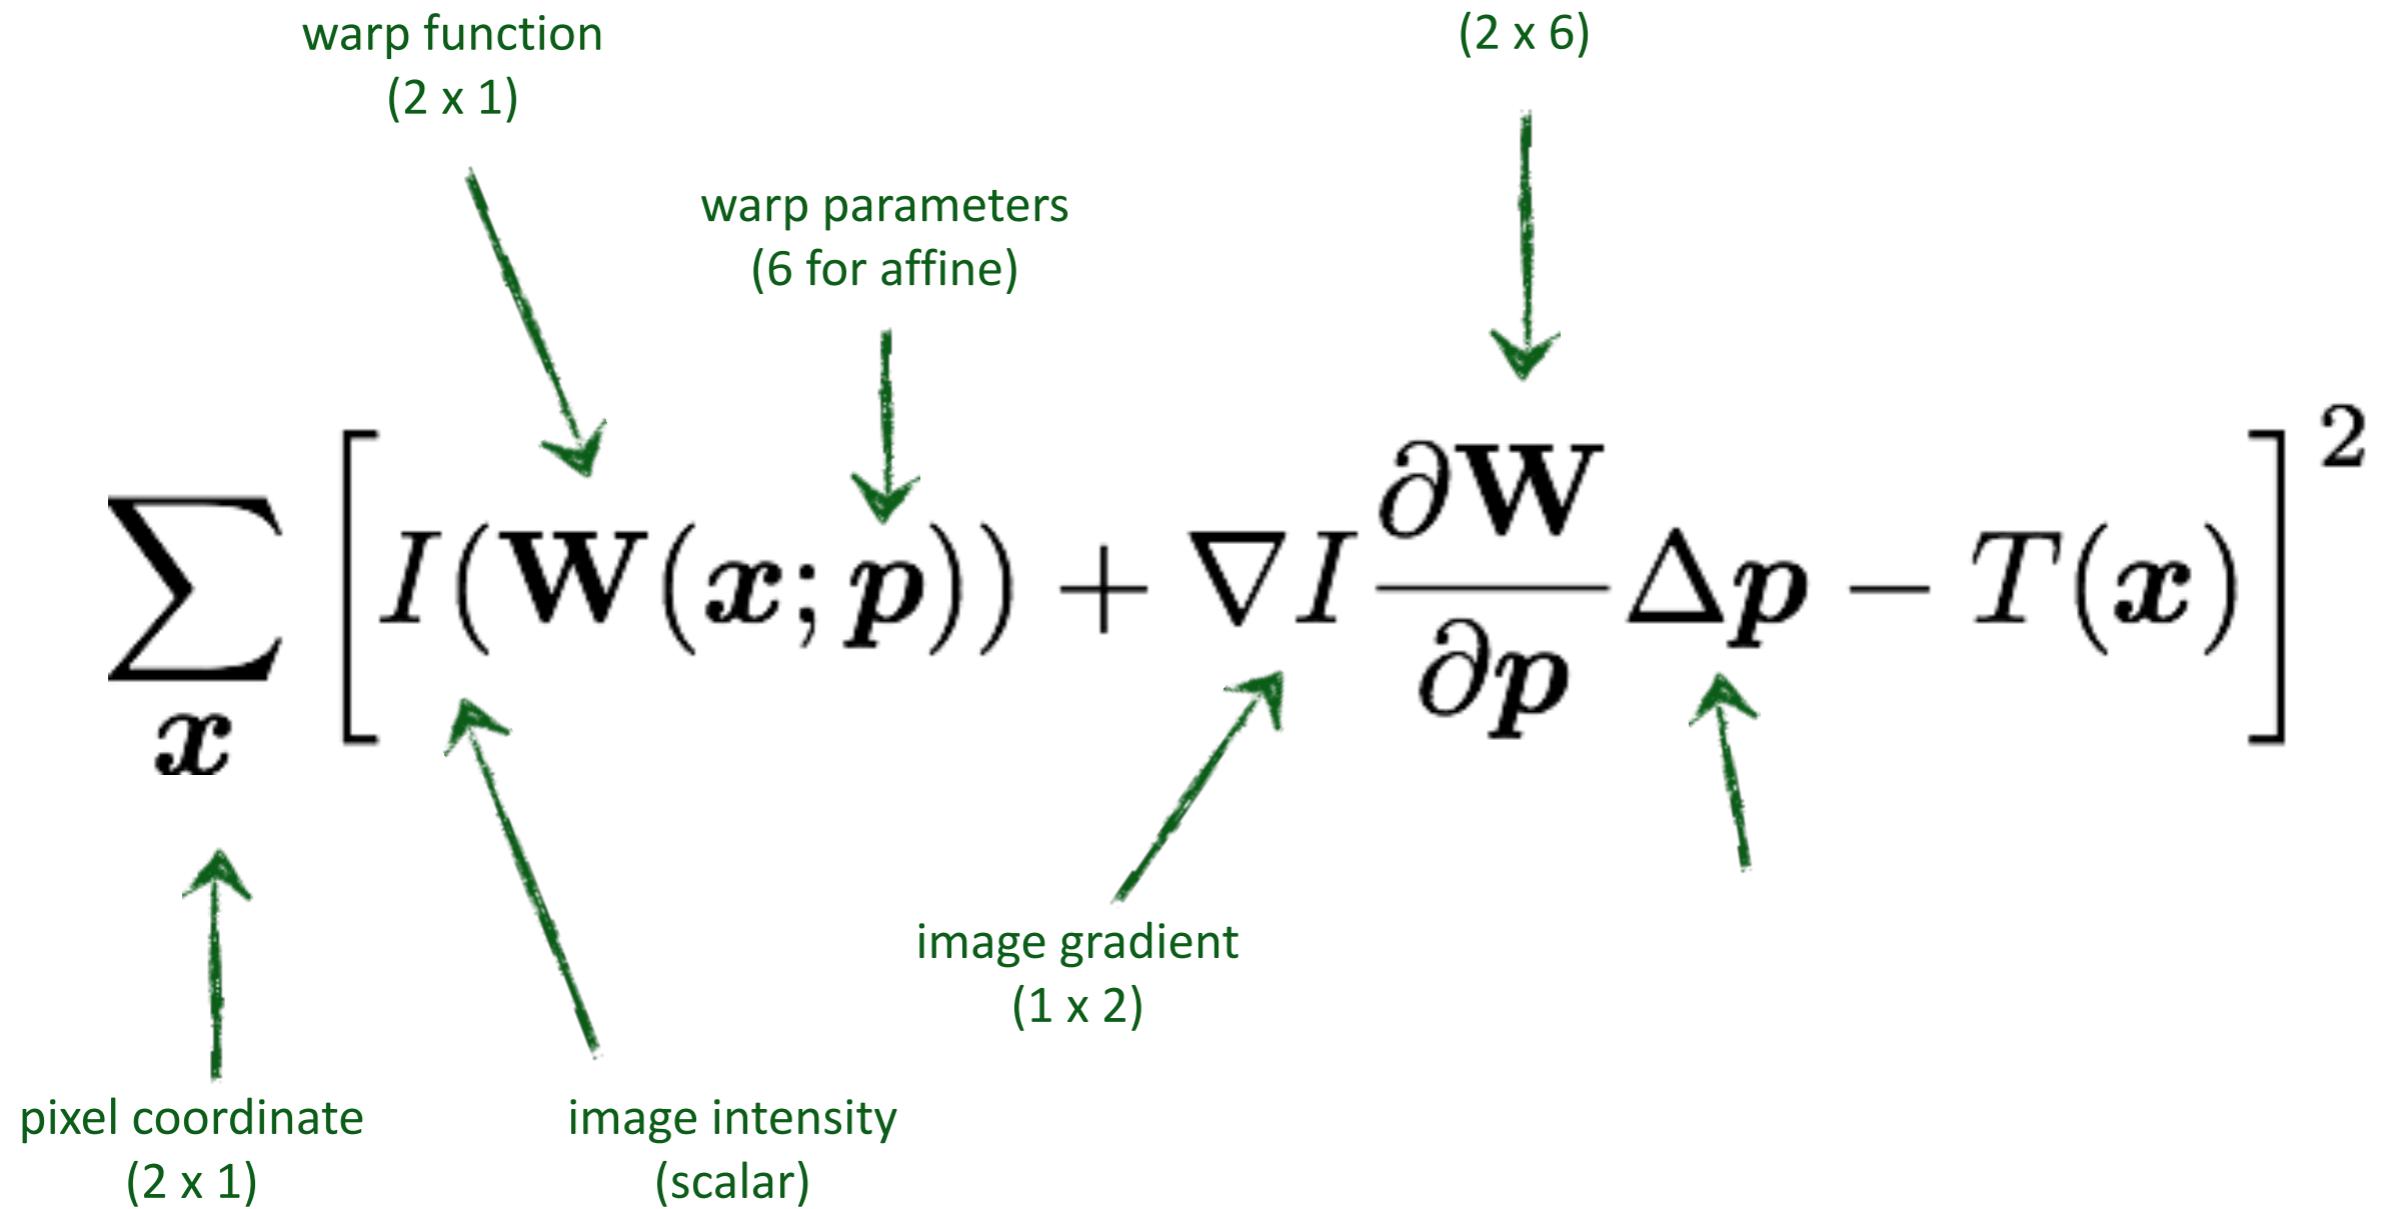
$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$



$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

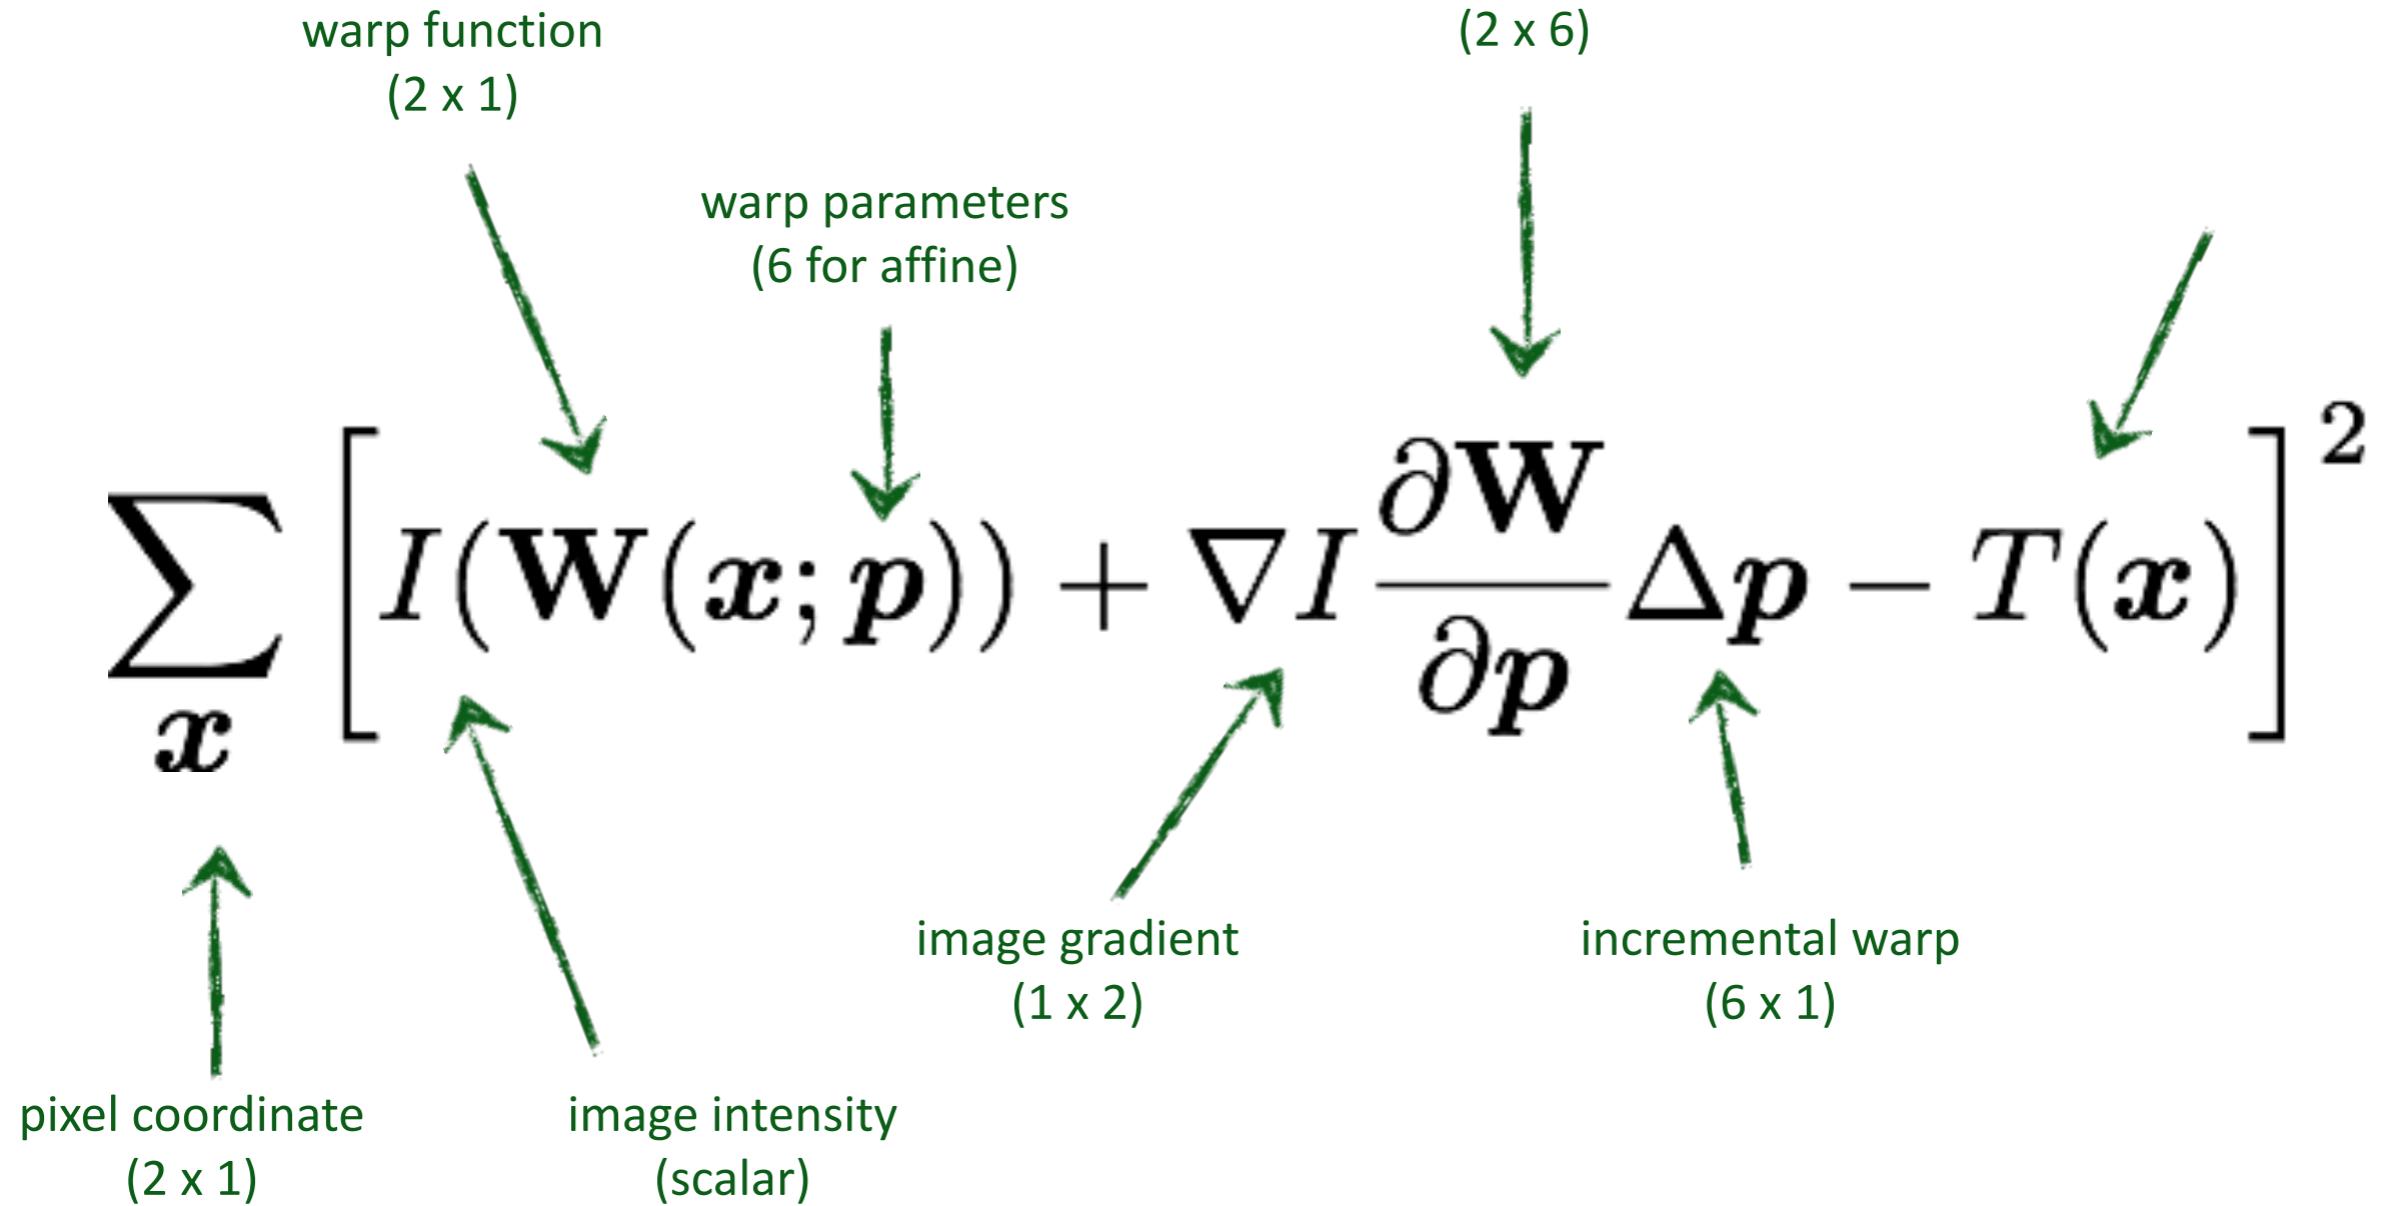
Partial derivatives of warp function
(2×6)



$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

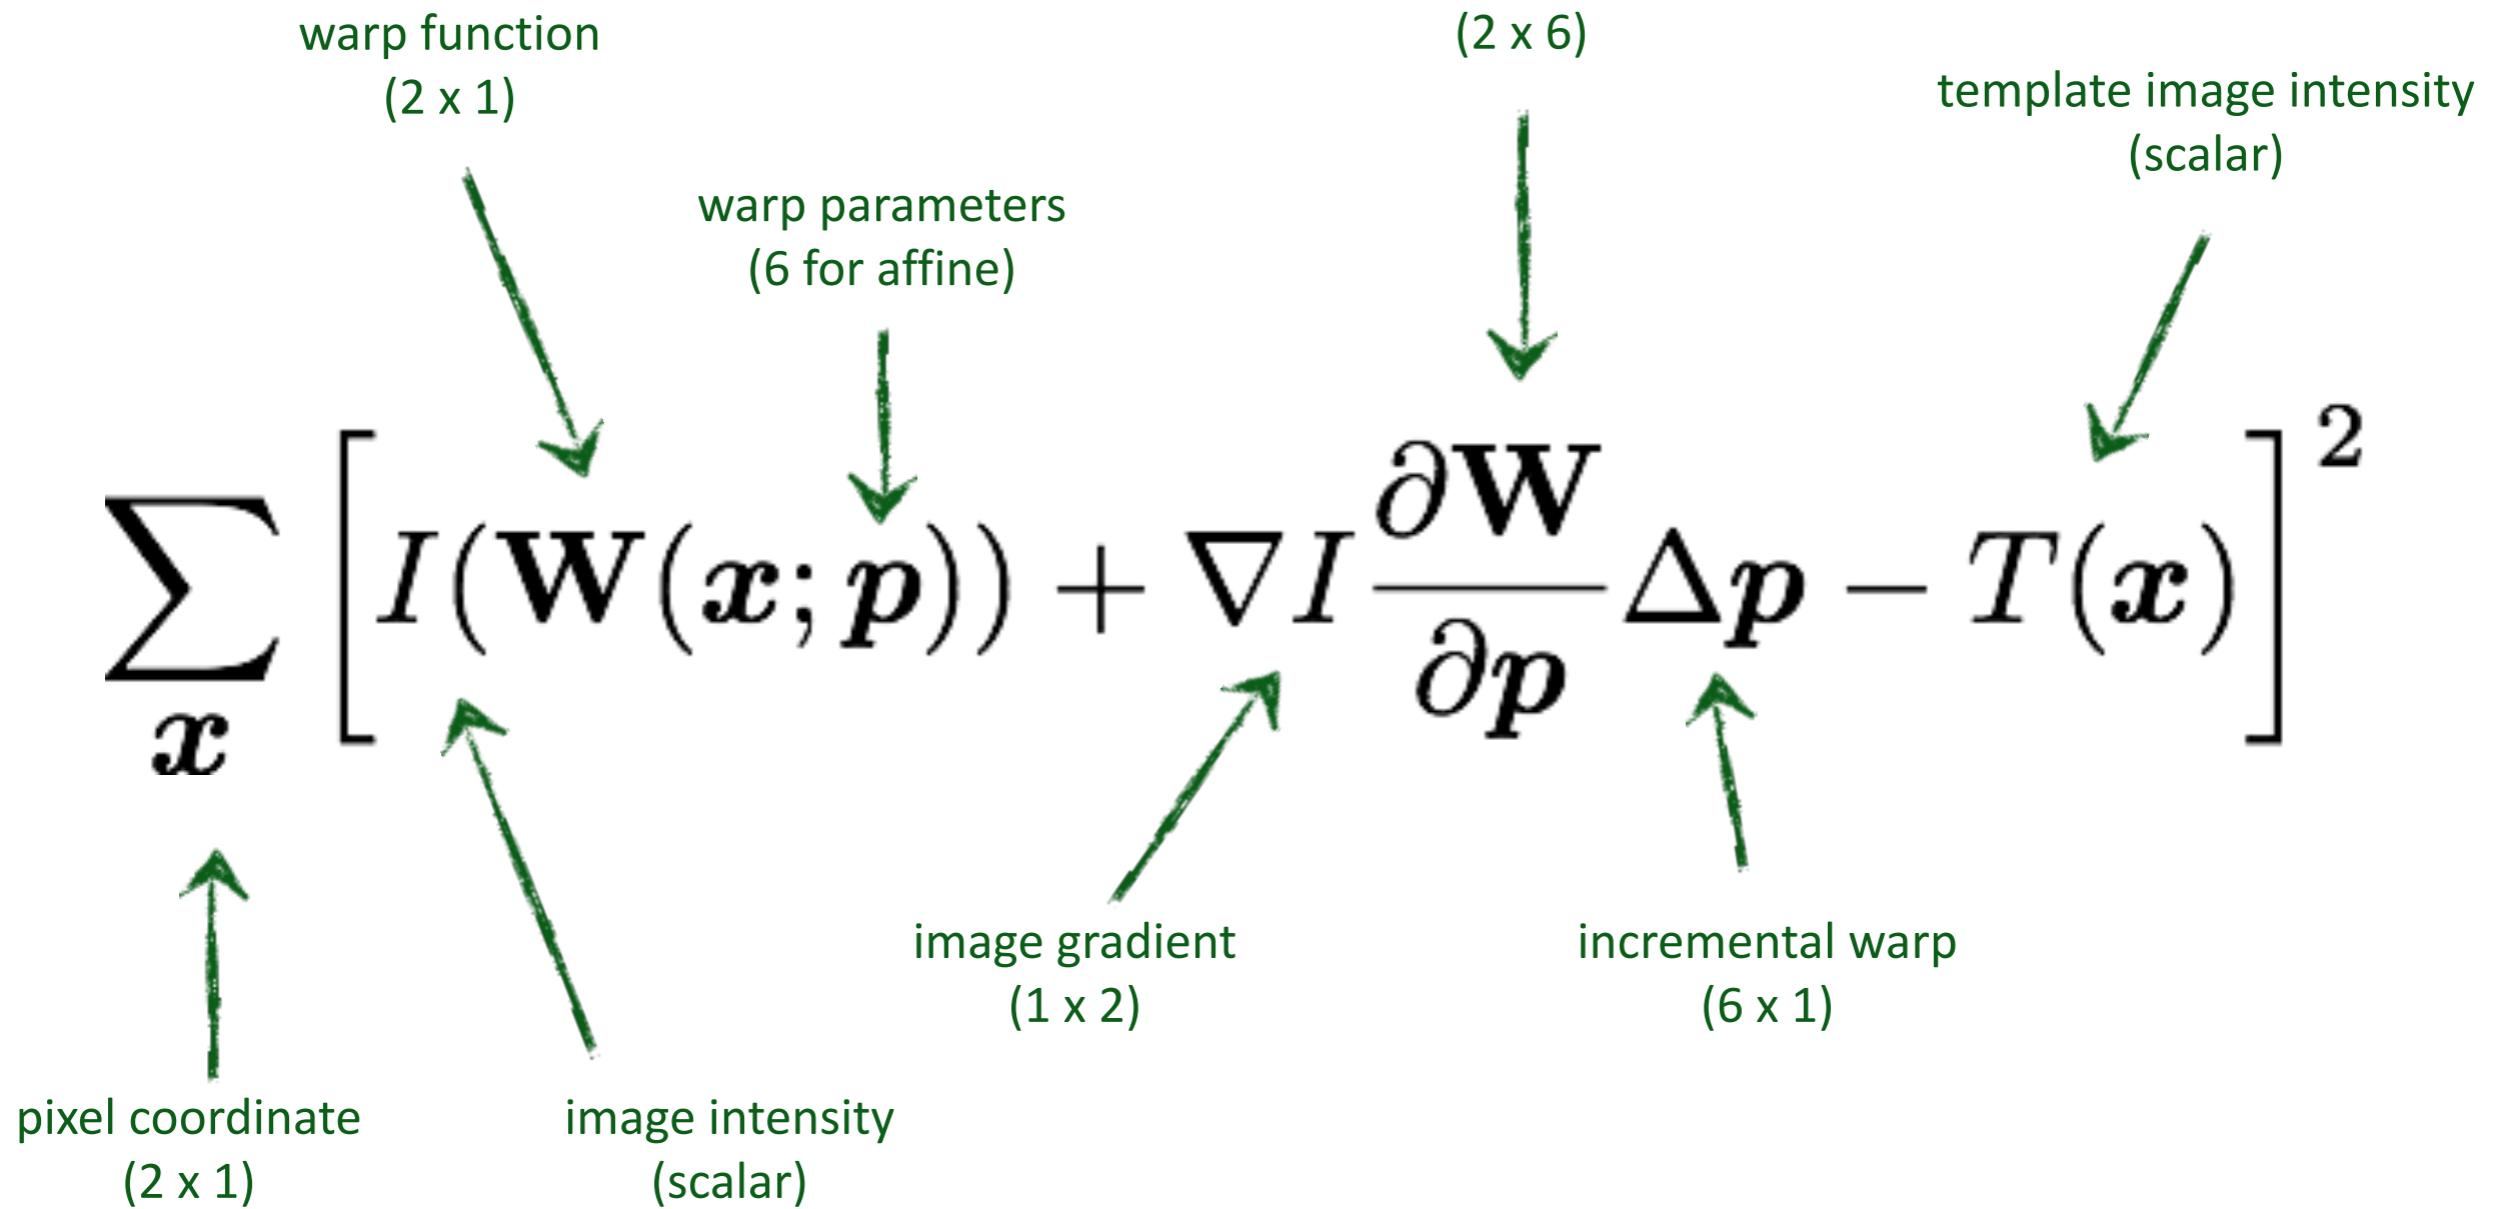
Partial derivatives of warp function
(2×6)



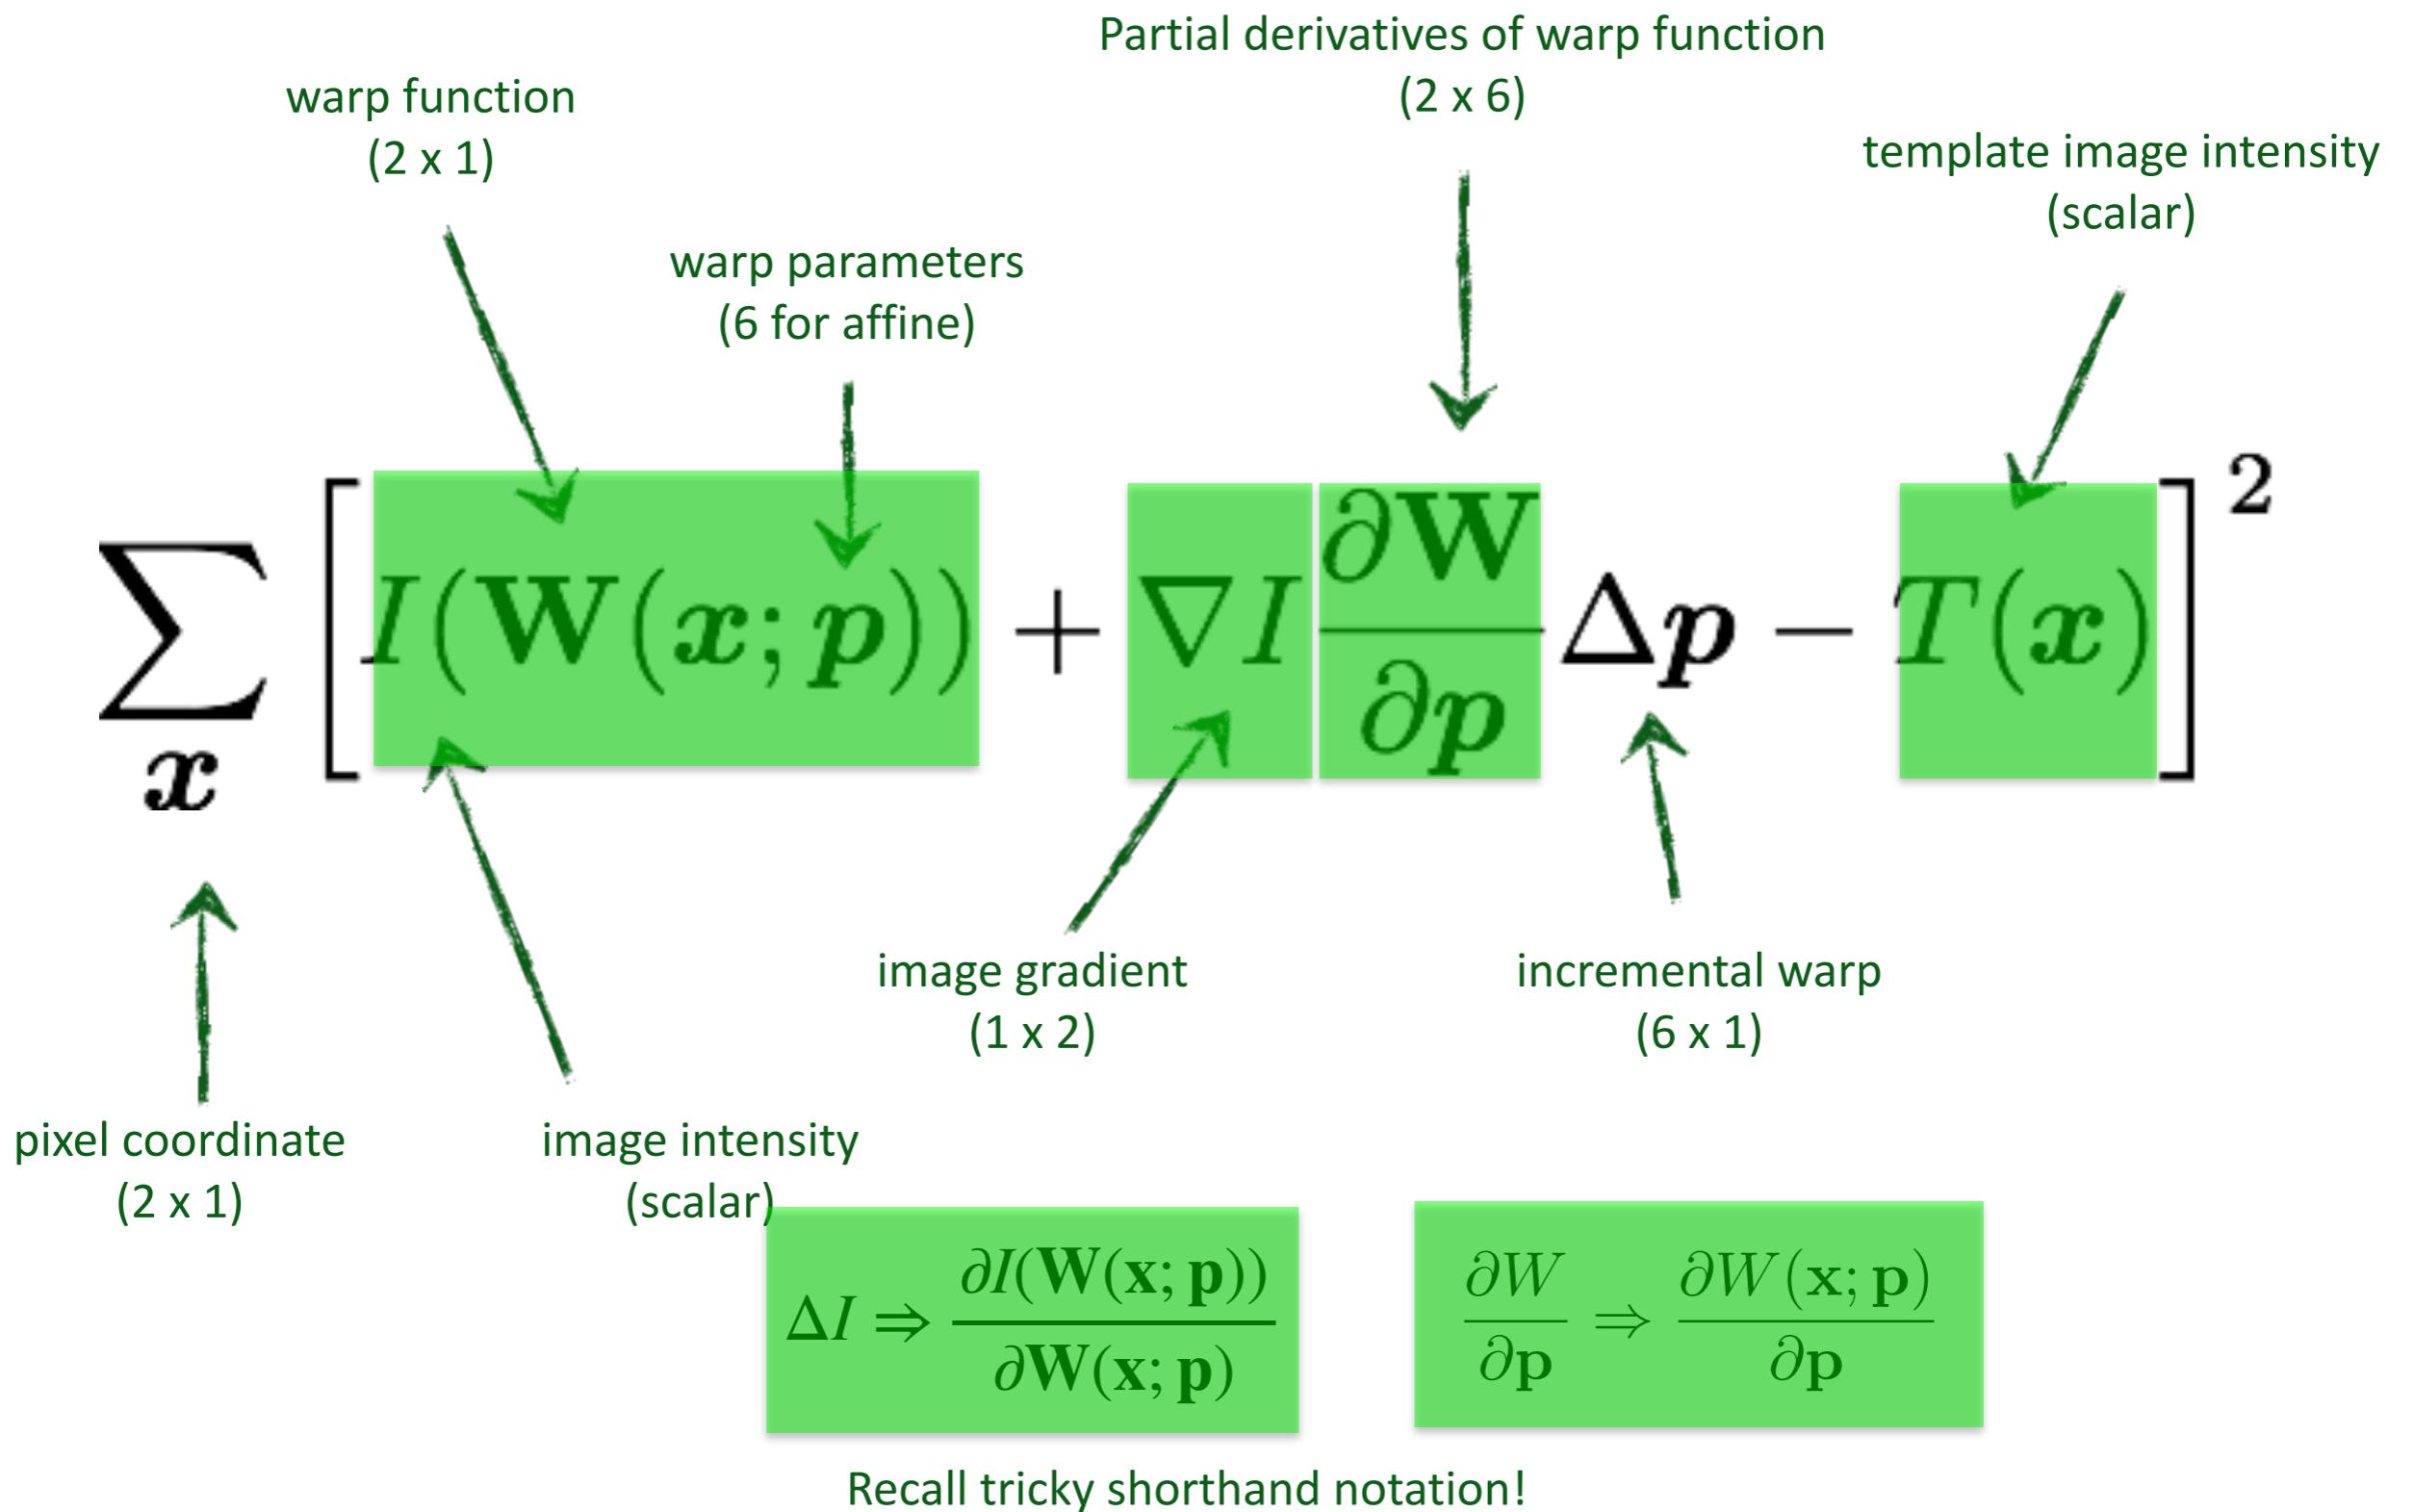
$$\begin{bmatrix} W_x(\mathbf{x}; \mathbf{p}) \\ W_y(\mathbf{x}; \mathbf{p}) \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p})) = I(p_1x + p_2y + p_3, p_4x + p_5y + p_6)$$

Partial derivatives of warp function
(2×6)



What variables depend on pixel coordinates \mathbf{x} ?



Summary

(of Lucas-Kanade Image Alignment)

Problem:

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

warped image template image

Difficult non-linear optimization problem

Strategy:

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

Solve for increment assuming a good initial guess

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Assume increment will be small and
linearize with Taylor series

... and then solve for $\Delta\mathbf{p}$

OK, so how to solve this?

$$\min_{\Delta p} \sum_x \left[I(\mathbf{W}(x; p)) + \nabla I \frac{\partial \mathbf{W}}{\partial p} \Delta p - T(x) \right]^2$$

(move terms around)

$$\min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - \{T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))\} \right]^2$$

vector of constants vector of variables constant

look familiar?

$$\min_{\mathbf{x}} \sum_i (A[i, :] \mathbf{x} - b[i])^2$$

$$\min_{\mathbf{x}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2$$

what are dimensions of A?
#pixels by 6 (for affine)

Hint: when implementing, explicitly compute matrix A and vector b in parallel and avoid looping over pixels
You may solve for Δp with `np.linalg.lstsq(A,b)` or `np.linalg.svd(A)` or ... [next slides]

Recall least squares soln for overdetermined (tall,skinny) A is given by

$$x = (A^\top A)^{-1} A^\top b$$

$$\min_{\Delta p} \sum_x \left[\nabla I \frac{\partial \mathbf{W}}{\partial p} \Delta p - \{T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; p))\} \right]^2$$

is optimized when

$$\Delta p = H^{-1} \sum_x \left[\nabla I \frac{\partial \mathbf{W}}{\partial p} \right]^\top [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; p))] \quad \begin{matrix} \text{after applying} \\ x = (A^\top A)^{-1} A^\top b \end{matrix}$$

approximate
Hessian matrix

$$\text{where } H = \sum_x \left[\nabla I \frac{\partial \mathbf{W}}{\partial p} \right]^\top \left[\nabla I \frac{\partial \mathbf{W}}{\partial p} \right] \quad A^\top A$$

what are the dimensions of H ? 6x6 (for affine)

Solve:

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

warped image template image

Difficult non-linear optimization problem

Strategy:

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) - T(\mathbf{x})]^2$$

Assume known approximate solution
Solve for increment

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Taylor series approximation
Linearize

Solution:

$$\Delta \mathbf{p} = H^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^\top [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

Solution to least squares
approximation

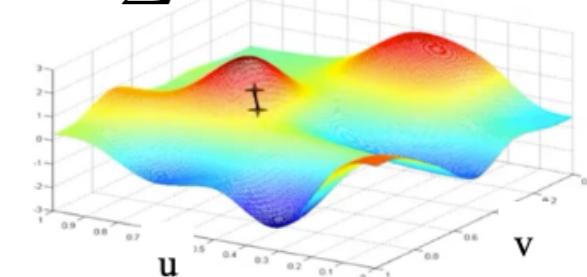
$$H = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^\top \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

Approximate Hessian

What's really going on under the hood:
nonlinear least squares optimization via Gauss-Newton descent

$$E(u, v) = \frac{1}{2} \left\| \begin{bmatrix} f(u, v) \\ g(u, v) \end{bmatrix} \right\|^2 = \frac{1}{2} \|\mathbf{f}(\mathbf{u})\|^2$$

$$E(u, v) = \sum [I(x + u, y + v) - T(x, y)]^2$$



1. First order method (gradient descent)

$$\begin{bmatrix} u \\ v \end{bmatrix} := \begin{bmatrix} u \\ v \end{bmatrix} - \alpha \begin{bmatrix} \frac{\partial f}{\partial u} & \frac{\partial f}{\partial v} \\ \frac{\partial g}{\partial u} & \frac{\partial g}{\partial v} \end{bmatrix}^T \begin{bmatrix} f(u, v) \\ g(u, v) \end{bmatrix}$$

jacobian $\mathbf{J}(\mathbf{u})$

1.5 Gauss-Newton descent

https://en.wikipedia.org/wiki/Gauss–Newton_algorithm

$$\mathbf{f}(\mathbf{u} + \Delta \mathbf{u}) \approx \mathbf{f}(\mathbf{u}) + \mathbf{J}(\mathbf{u}) \Delta \mathbf{u}$$

$$\mathbf{u} := \mathbf{u} + \arg \min_{\Delta \mathbf{u}} \frac{1}{2} \|\mathbf{f}(\mathbf{u}) + \mathbf{J}(\mathbf{u}) \Delta \mathbf{u}\|^2$$

Solve for update with least squares $\min_x \|\mathbf{-b} + \mathbf{Ax}\|^2$

Get the performance of second-order optimizers with first-order compute
(without computing second derivatives)!

Lucas Kanade (Additive alignment)

1. Warp image

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$$

2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$

3. Compute gradient

$$\nabla I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$$

dx,dy gradients of warped image

4. Evaluate Jacobian

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$$

5. Compute Hessian

$$H$$

$$H = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^{\top} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

6. Compute

$$\Delta \mathbf{p}$$

$$\Delta \mathbf{p} = H^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^{\top} [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

7. Update parameters

$$\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$$

Just 8 lines of code!

HW: Tracking by iterative template alignment



Start with template from first frame and repeat:

1. Align template to new frame with Lucas Kanade
2. Update template with new frame (or not?)

Important extensions

Lucas-Kanade 20 Years On: A Unifying Framework

SIMON BAKER AND IAIN MATTHEWS

The Robotics Institute, Carnegie Mellon University, 5000 Forbes Avenue, Pittsburgh, PA 15213, USA

simonb@cs.cmu.edu

iainm@cs.cmu.edu

image



template



Key insight: exploit flexibility of warping the image *or* warping the template

Image Alignment

(start with an initial solution, match the image and template)

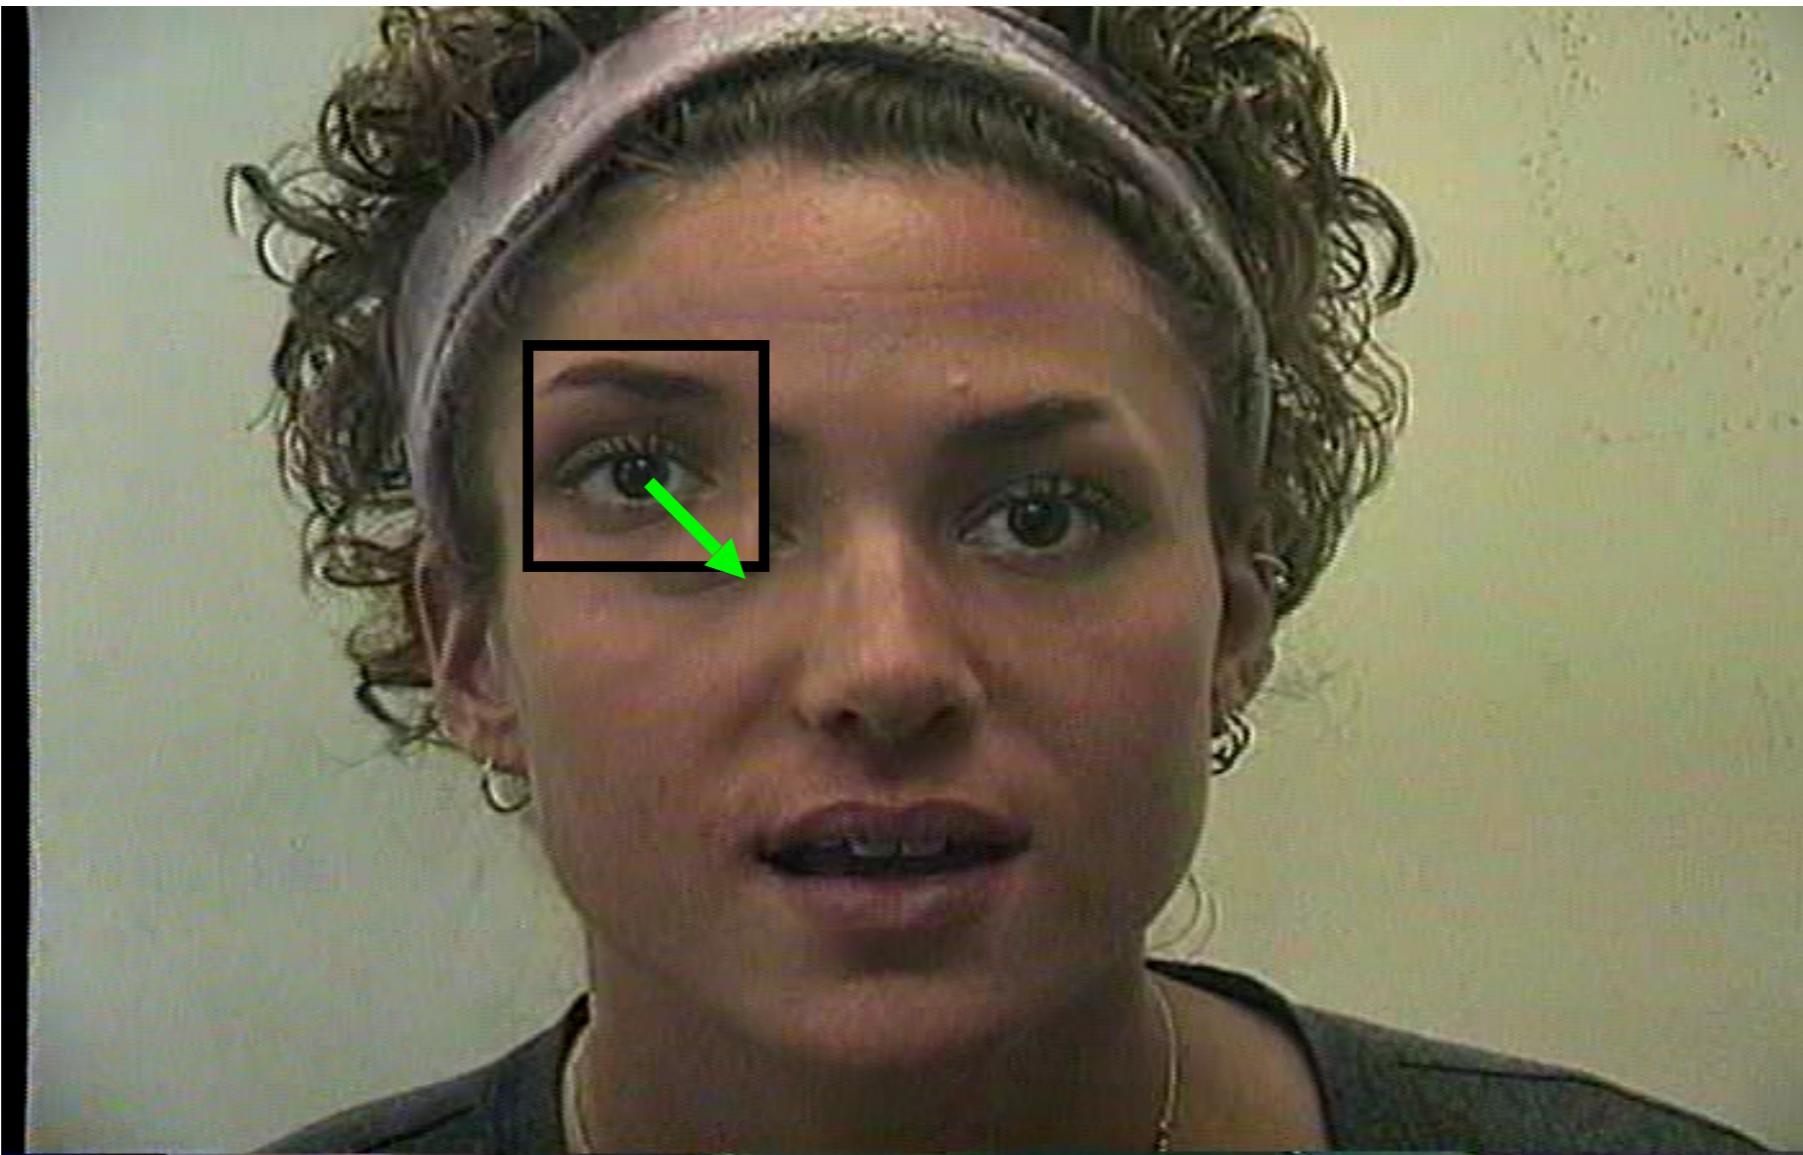


Image Alignment Objective Function

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Given an initial solution...several possible formulations

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) - T(\mathbf{x})]^2$$

incremental perturbation of parameters

Image Alignment Objective Function

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Given an initial solution...several possible formulations

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

incremental perturbation of parameters

Compositional Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$$

incremental warps of image

Additive strategy



[Loose analogy]

Compositional strategy



Additive



$$w(x;p)$$

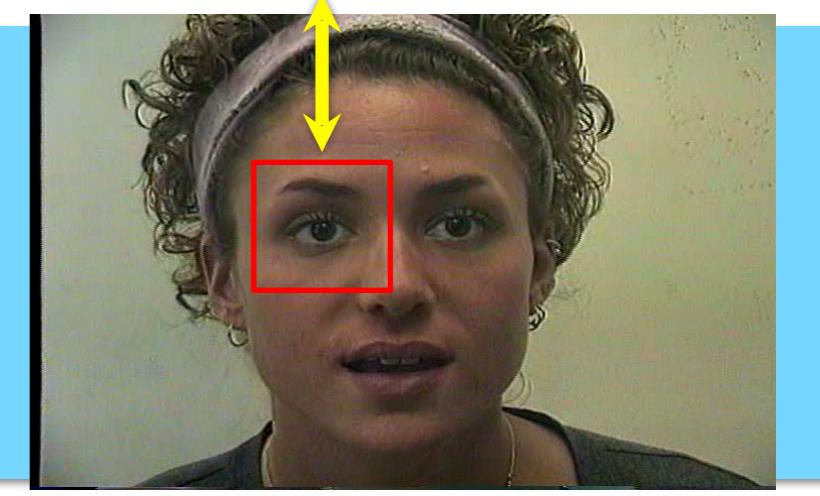
Additive



$W(x; p + \Delta p)$



$T(x)$



$W(x; p)$

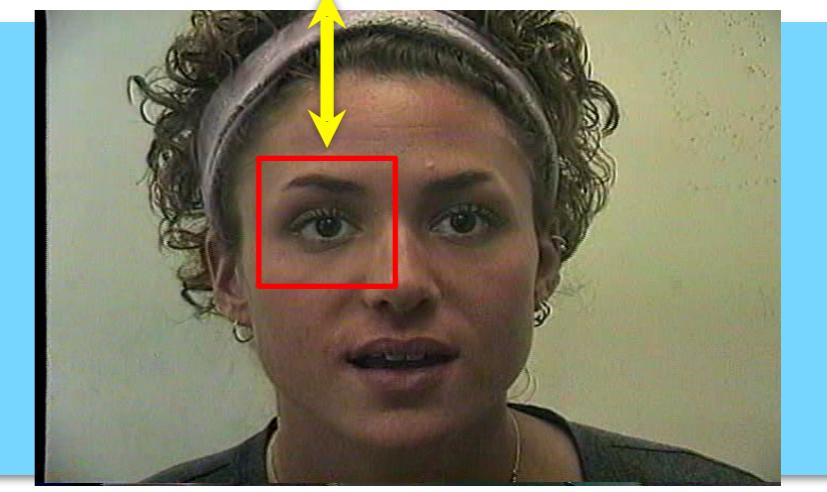
Additive



$$W(x; p + \Delta p)$$



$$T(x)$$



Compositional (with a twist; apply Δp warp *first* instead of *second*. We'll see why soon!)



$$W(x; p)$$



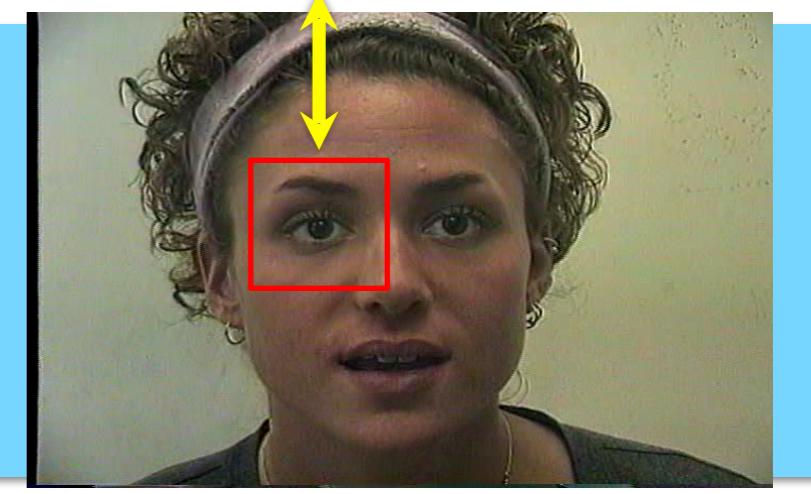
Additive



$W(x; p + \Delta p)$



$T(x)$



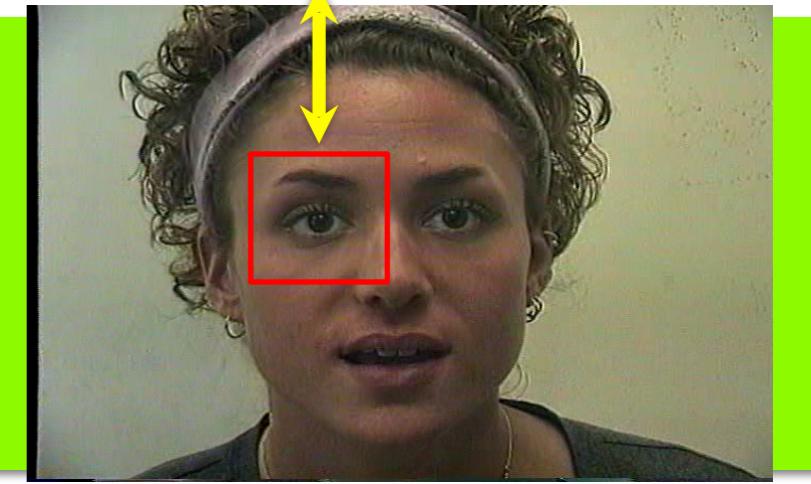
Compositional



$W(x; p) \circ W(x; \Delta p)$



$T(x)$



$W(x; 0 + \Delta p) = W(x; \Delta p)$

$W(x; p)$

Compositional Alignment

Original objective function (SSD)

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Assuming an initial solution \mathbf{p} and a compositional warp increment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p}) - T(\mathbf{x})]^2$$

Compositional Alignment

Original objective function (SSD)

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Assuming an initial solution \mathbf{p} and a compositional warp increment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p}) - T(\mathbf{x})]^2$$

Another way to write the composition

Identity warp

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p}) \equiv \mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})$$

$$\mathbf{W}(\mathbf{x}; \mathbf{0})$$

Compositional Alignment

Original objective function (SSD)

$$\min_{\mathbf{p}} \sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]^2$$

Assuming an initial solution \mathbf{p} and a compositional warp increment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p}) - T(\mathbf{x})]^2$$

Another way to write the composition

Identity warp

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p}) \equiv \mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})$$

$$\mathbf{W}(\mathbf{x}; \mathbf{0})$$

Skipping over the derivation...the new update rule is

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta\mathbf{p})$$

So what's so great about this compositional form?

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

linearized form

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I(\mathbf{x}') \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Compositional Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$$

linearized form

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I(\mathbf{x}') \frac{\partial \mathbf{W}(\mathbf{x}; \mathbf{0})}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

linearized form

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I(\mathbf{x}') \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Jacobian of $\mathbf{W}(\mathbf{x}; \mathbf{p})$

Compositional Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$$

linearized form

$$\sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I(\mathbf{x}') \frac{\partial \mathbf{W}(\mathbf{x}; \mathbf{0})}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$

Jacobian of
 $\mathbf{W}(\mathbf{x}; \mathbf{0})$

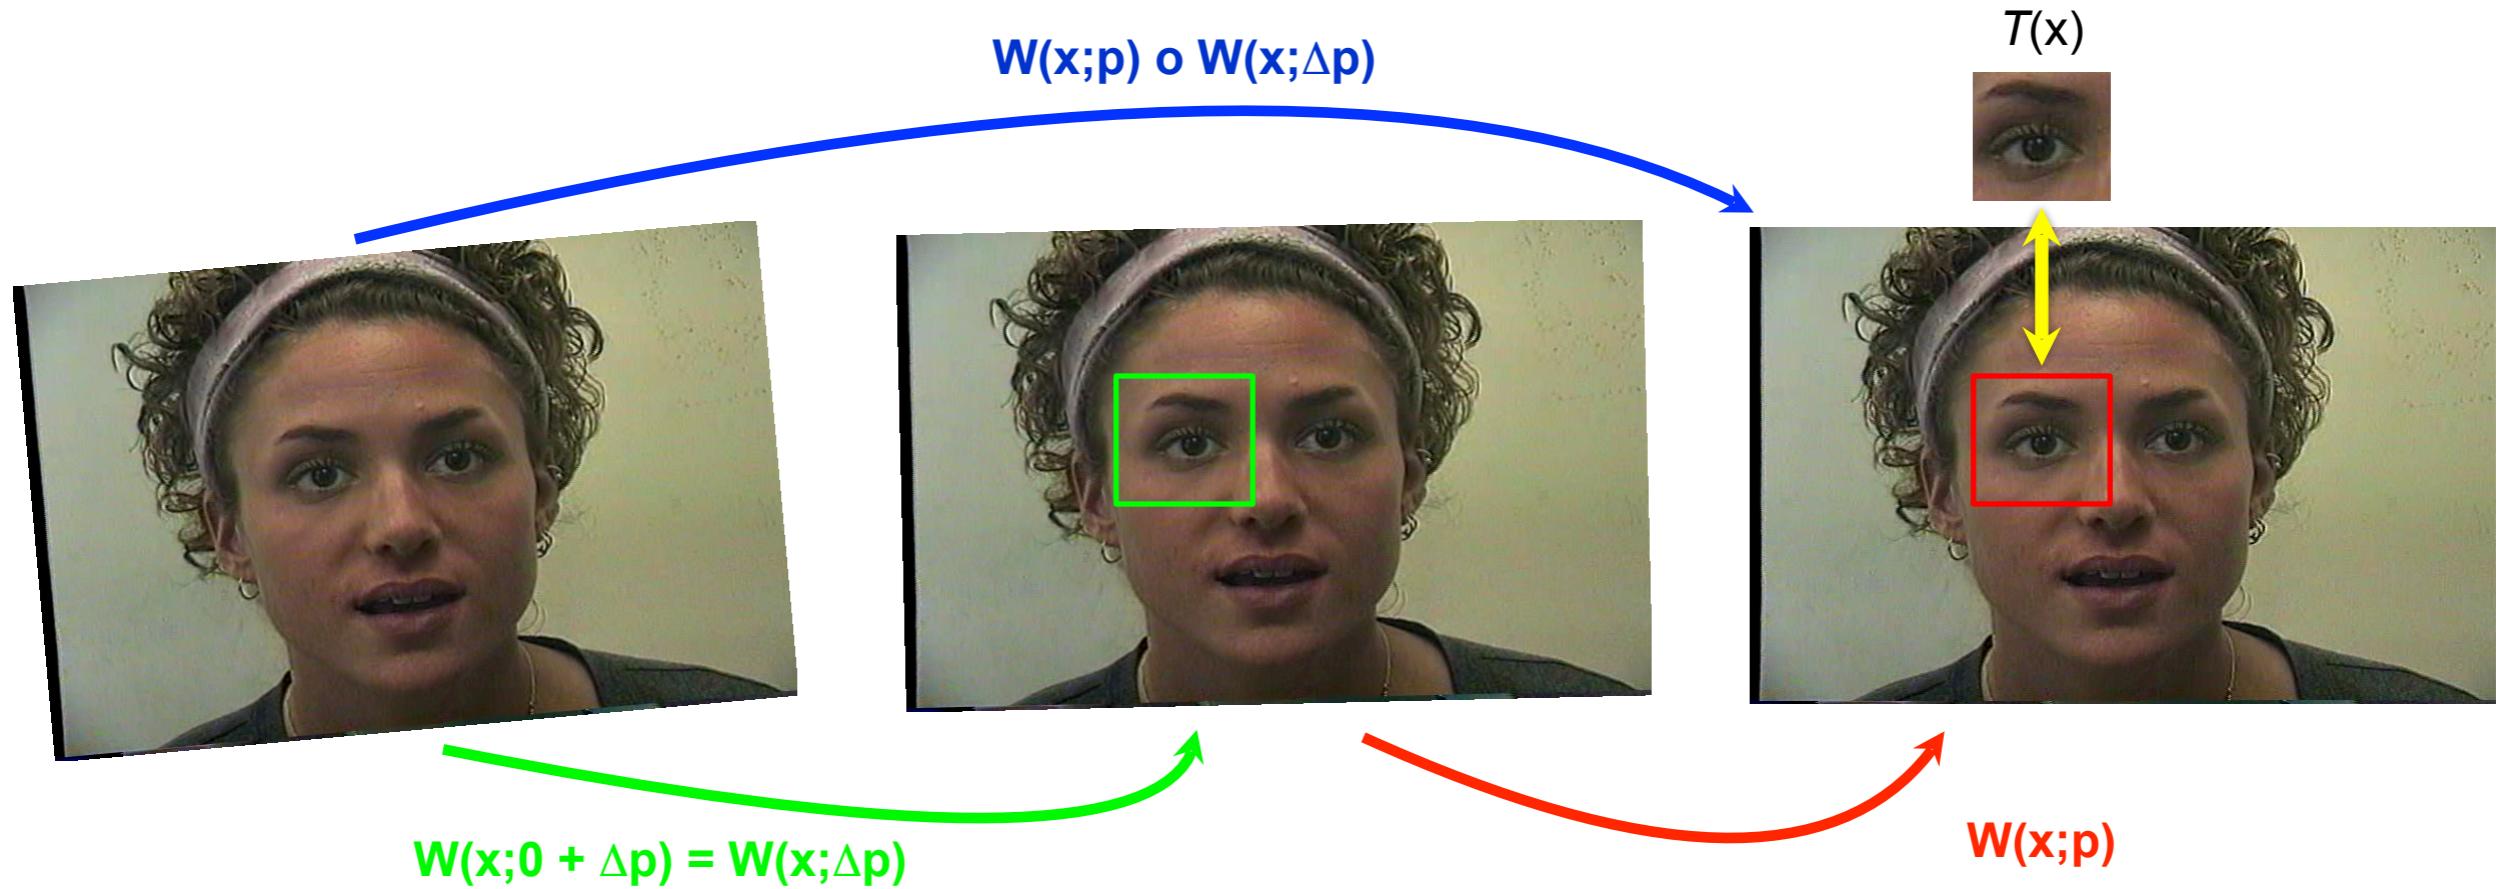
The Jacobian is constant and can be precomputed!

(benefit of applying incremental update *first*, since it requires linearizing about $\mathbf{p} = \mathbf{0}$)

Compositional Image Alignment

Minimize

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta \mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2 \approx \sum_{\mathbf{x}} \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I(\mathbf{W}) \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]^2$$



Jacobian is simple and can be precomputed

Lucas Kanade (Additive alignment)

1. Warp image $I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]^2$
3. Compute gradient $\nabla I(\mathbf{x}')$
4. Evaluate Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$
5. Compute Approx Hessian H
6. Compute $\Delta \mathbf{p}$
7. Update parameters $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$

Shum-Szeliski (Compositional alignment)

1. Warp image $I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]^2$
3. Compute gradient $\nabla I(\mathbf{x}')$
4. Evaluate Jacobian $\frac{\partial \mathbf{W}(\mathbf{x}; \mathbf{0})}{\partial \mathbf{p}}$
5. Compute Approx Hessian H
6. Compute $\Delta \mathbf{p}$
7. Update parameters $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta \mathbf{p})$

Any other speed up techniques?

Inverse alignment

Why not compute warp updates on the template?

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

Compositional Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$$

Why not compute warp updates on the template?

Additive Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x})]^2$$

Compositional Alignment

$$\sum_{\mathbf{x}} [I(\mathbf{W}(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p}); \mathbf{p})) - T(\mathbf{x})]^2$$

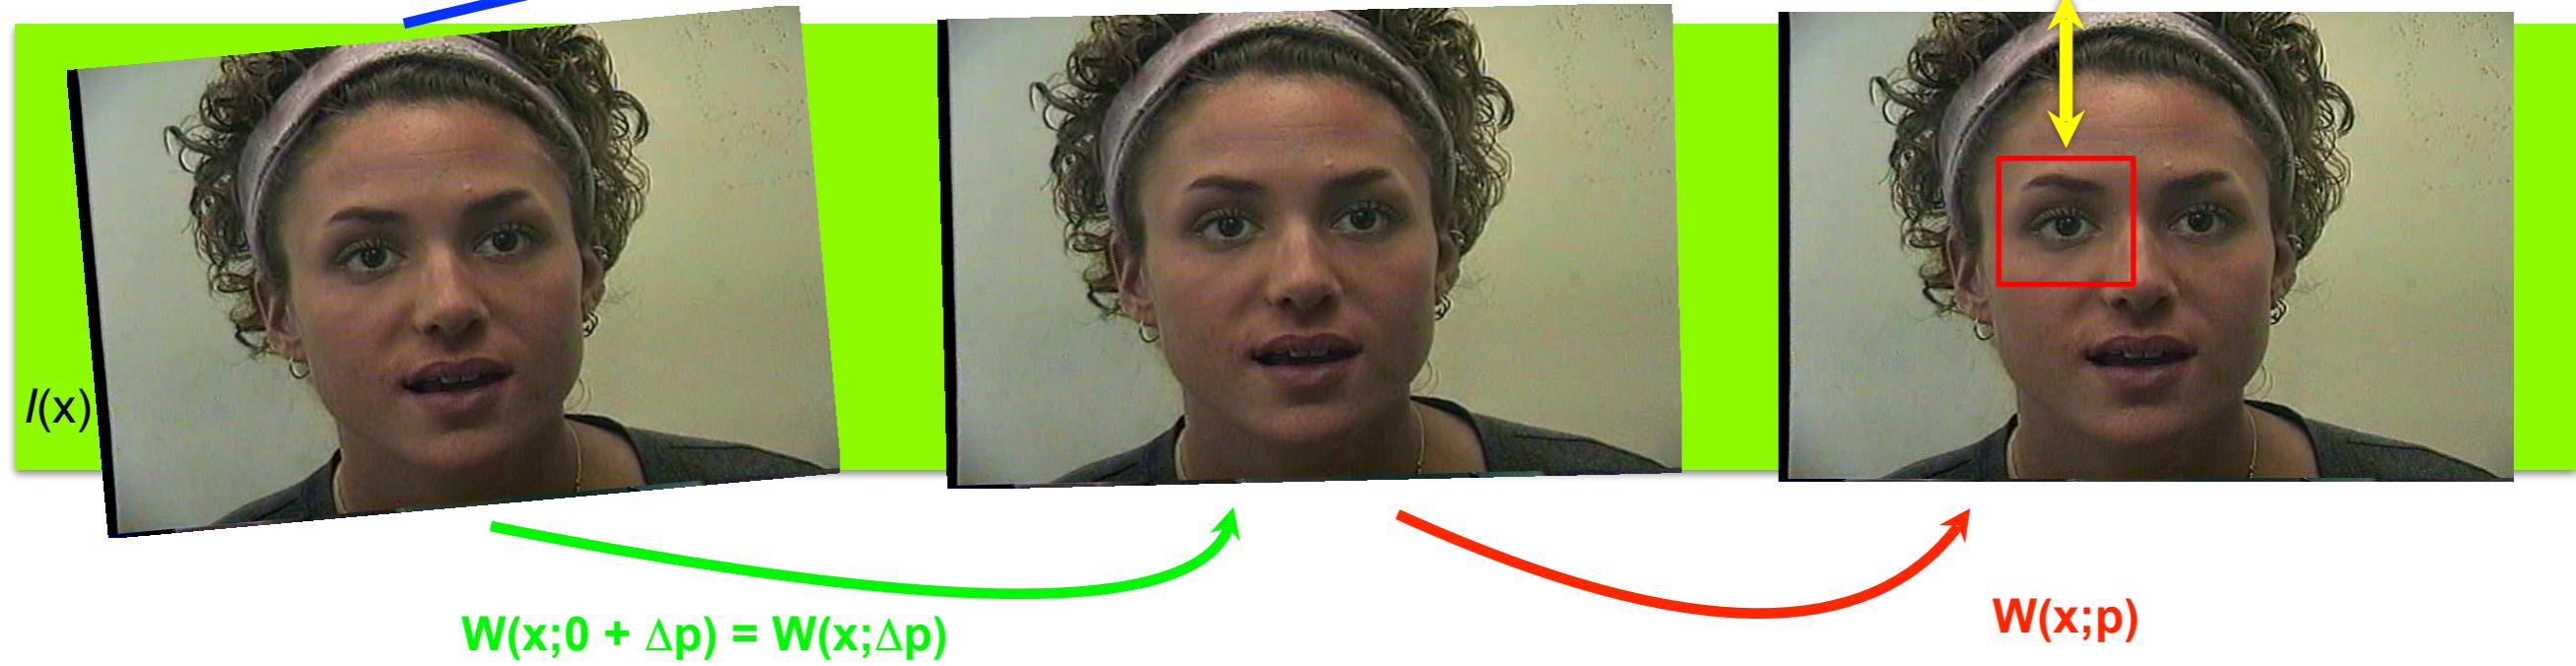
What happens if you let the template
be warped too?

Inverse Compositional Alignment

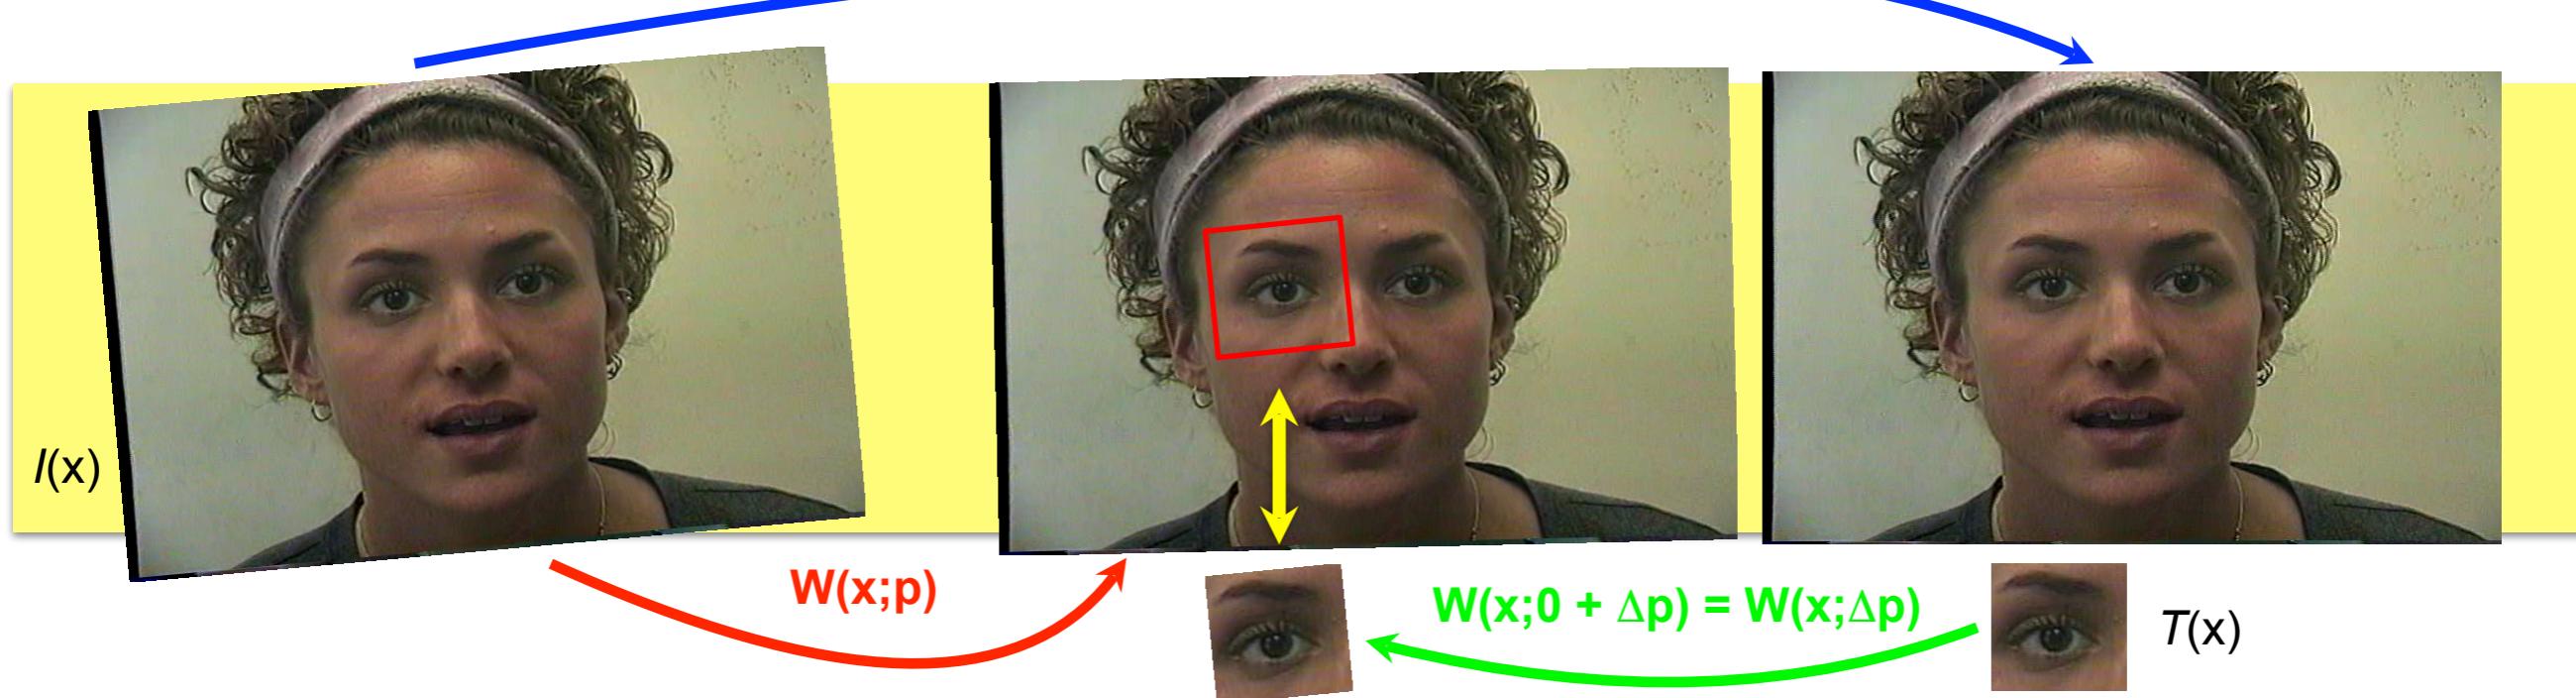
$$\sum_{\mathbf{x}} [T(\mathbf{W}(\mathbf{x}; \Delta\mathbf{p})) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]^2$$

1. Compute $\Delta\mathbf{p}$ warp on template.
2. But instead of applying it to the template, apply its *inverse* warp to the image. *This implies the linearized template (ie, jacobians, hessians) never need be recomputed since the template never changes!*

Compositional



Inverse compositional



Compositional strategy



Inverse Compositional strategy



So what's so great about this inverse compositional form?

Inverse Compositional Alignment

Minimize

$$\sum_{\mathbf{x}} [T(\mathbf{W}(\mathbf{x}; \Delta p)) - I(\mathbf{W}(\mathbf{x}; p))]^2 \approx \sum_{\mathbf{x}} \left[T(\mathbf{W}(\mathbf{x}; \mathbf{0})) + \nabla T \frac{\partial \mathbf{W}}{\partial p} \Delta p - I(\mathbf{W}(\mathbf{x}; p)) \right]^2$$

Solution

$$H = \sum_{\mathbf{x}} \left[\nabla T \frac{\partial \mathbf{W}}{\partial p} \right]^\top \left[\nabla T \frac{\partial \mathbf{W}}{\partial p} \right]$$

can be precomputed from template!

$$\Delta p = \sum_{\mathbf{x}} H^{-1} \left[\nabla T \frac{\partial \mathbf{W}}{\partial p} \right]^\top [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; p))]$$

Update

$$\mathbf{W}(\mathbf{x}; p) \leftarrow \mathbf{W}(\mathbf{x}; p) \circ \mathbf{W}(\mathbf{x}; \Delta p)^{-1}$$

Properties of inverse compositional alignment

Jacobian can be precomputed

It is constant - evaluated at $W(x; 0)$

Gradient of template can be precomputed

It is constant

Approx Hessian can be precomputed

$$H = \sum_{\mathbf{x}} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^\top \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

$$\Delta \mathbf{p} = \sum_{\mathbf{x}} H^{-1} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^\top [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

(main term that needs to be computed)

Warp must be invertible

Lucas Kanade (Additive alignment)

1. Warp image $I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]^2$
3. Compute gradient $\nabla I(\mathbf{W})$
4. Evaluate Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$
5. Compute Hessian H
6. Compute $\Delta \mathbf{p}$
7. Update parameters $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$

Shum-Szeliski (Compositional alignment)

1. Warp image $I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$
3. Compute gradient $\nabla I(\mathbf{x}')$
4. Evaluate Jacobian $\frac{\partial \mathbf{W}(\mathbf{x}; \mathbf{0})}{\partial \mathbf{p}}$
5. Compute Approx Hessian H
6. Compute $\Delta \mathbf{p}$
7. Update parameters $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta \mathbf{p})$

Baker-Matthews (Inverse Compositional alignment)

1. Warp image $I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $[T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$

3. Compute gradient $\nabla T(\mathbf{W})$

4. Evaluate Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$

5. Compute Approx Hessian H

$$H = \sum_{\mathbf{x}} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^{\top} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$$

6. Compute $\Delta \mathbf{p}$

$$\Delta \mathbf{p} = \sum_{\mathbf{x}} H^{-1} \left[\nabla T \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^{\top} [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

7. Update parameters $\mathbf{W}(\mathbf{x}; \mathbf{p}) \leftarrow \mathbf{W}(\mathbf{x}; \mathbf{p}) \circ \mathbf{W}(\mathbf{x}; \Delta \mathbf{p})^{-1}$

Algorithm	Efficient	Authors
Forwards Additive	No	Lucas, Kanade
Forwards compositional	No	Shum, Szeliski
Inverse Additive	Yes	Hager, Belhumeur
Inverse Compositional	Yes	Baker, Matthews

A look back

- Geometric warps
 - Foreshortening-vs-projective
 - Piecewise affine
 - Fwd-vs-inverse warps
- Lucas-Kanade Alignment
 - Gauss Newton optimization of nonlinear least squares
 - Inverse composition

